

Distributed Databases

And their relevance for NOSQL



NOSQL Goals

Recall the NOSQL Goals:

- Schemaless: Allow flexible (even runtime) schema definition [data structure]
- Reliability / availability: Keep delivering service even if its software or hardware components fail [recovery]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

NOSQL Goals

Recall the NOSQL Goals:

- Schemaless: Allow flexible (even runtime) schema definition [data structure]
- Reliability / availability: Keep delivering service even if its software or hardware components fail [distribution]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

NOSQL Goals

- Reliability / availability: Keep delivering service even if its software or hardware components fail [distribution]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

NOSQL Goals

- Reliability / availability: Keep delivering service even if its software or hardware components fail [distribution]
- Scalability: Continuously evolve to support a growing amount of tasks [distribution]
- Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

Most NOSQL goals can be achieved by means of distribution!

Activity: Why Distribution?




- Objective: Recognize the benefits of distributing data
- Tasks:

1. (15') By pairs, answer the following questions:

- How long would it take to read 1TB with sequential access (fig. a)? (in secs)
 - Can you identify any additional drawback to be considered?
- How long would it take to read 1TB with parallel access (fig. b)? Assume 100 disks on the same machine with shared-memory and infinite CPU capacity.
 - Can you identify any additional drawback to be considered?
- How long would it take to read 1TB with distributed access (fig. c)? Assume 100 shared-nothing machines in a star-shape LAN in a single rack where all data is sent to the centre.
 - Can you identify any additional drawback to be considered?
- Now, repeat the exercise considering a single random access. What changes?

2. (5') Discussion

Type	Latency	Bandwidth
Disk	$\approx 5 \times 10^{-3}$ s (5 millisec.);	At best 100 MB/s
LAN	$\approx 1 - 2 \times 10^{-3}$ s (1-2 millisec.);	$\approx 1\text{GB/s}$ (single rack); $\approx 100\text{MB/s}$ (switched);
Internet	Highly variable. Typ. 10-100 ms.;	Highly variable. Typ. a few MB/s.;
Bottom line (1): it is approx. one order of magnitude faster to exchange main memory data between 2 machines in a data center, than to read on the disk.		
Bottom line (2): exchanging through the Internet is slow and unreliable with respect to LANs.		



Activity: Why Distribution?

- Objective: Recognize the benefits of distributing data
- Tasks:
 1. (15') By pairs, answer the following questions:
 - a) How long would it take to read 1TB with sequential access (fig. a)? (in secs)
 - Can you identify any additional drawback to be considered?
 - b) How long would it take to read 1TB with parallel access (fig. b)? Assume 100 disks on the same machine with shared-memory and infinite CPU capacity.
 - Can you identify any additional drawback to be considered?
 - c) How long would it take to read 1TB with distributed access (fig. c)? Assume 100 shared-nothing machines in a star-shape LAN in a single rack where all data is sent to the centre.
 - Can you identify any additional drawback to be considered?
 - d) Now, repeat the exercise considering a single random access. What changes?
 2. (5') Discussion


Conclusions:

- Disk transfer rate is a bottleneck for batch processing of large scale datasets; parallelization and distribution can help to eliminate this bottleneck
- Disk seek time is a bottleneck for transactional applications that submit a high rate of random accesses; replication, distribution of writes and distribution of reads make such applications scalable
- When possible, data should be accessed where they are stored (or near) to avoid costly data exchange over the network

Definition and Distributed Architectures

Distributed Database

- A distributed database (DDB) is a database where data management is distributed over several nodes in a network.
 - Each node is a database itself
 - Potential heterogeneity
 - Nodes communicate through the network




Distributed Architectures

- Main objective: hide implementation (i.e., physical) details to the users who should not know they are dealing with a distributed system
 - Network transparency
 - Data access must be independent regardless where data is stored
 - Each data object must have a unique name
 - Replication transparency
 - The user must not be aware of the existing replicas
 - Fragmentation transparency
 - The user must not be aware of partitioning



[**Desiderata:** Data independency at the logical and physical level must be guaranteed to avoid accessing directly to physical structures]

- As seen in centralized DBs (ANSI SPARC)

From a Centralized to a Distributed DBMS Architecture



From a Centralized to a Distributed DBMS Architecture



Challenges

Principles of Distributed Databases




Challenges in Distributed Databases

- I. Distributed DB design
 - Node distribution
 - Data fragments
 - Data allocation (replication)
- II. Distributed DB catalog
 - Fragmentation trade-off: Where to place the DB catalog
 - Global or local for each node
 - Centralized in a single node or distributed
 - Single-copy vs. Multi-copy
- III. Distributed query processing
 - Data distribution / replication
 - Communication overhead
- IV. Distributed transaction management
 - How to enforce the ACID properties
 - Replication trade-off: Queries vs. Data consistency between replicas (updates)
 - Distributed recovery system
 - Distributed concurrency control system
- V. Security issues
 - Network security

Challenges in Data Distribution

A	B	C	D	E

Conceptual View




Physical View

Challenges in Data Distribution

Conceptual View

A	B	C	D	E



Physical View

Challenge I

Database Design: Fragmentation and Allocation (Replication)




Challenge I: DDB Design

- Given a DB and its workload, how should the DB be split and allocated to sites as to optimize certain objective functions
 - For example, minimize resource consumption for query processing
- Two main issues:
 - Data fragmentation
 - Data allocation (data replication)

Challenge I: Data Fragments

Conceptual View

A	B	C	D	E




Physical View

Challenge I: Data Fragments

Conceptual View


A	B	C	D	E
Apply a fragmentation strategy				



Physical View

Challenge I: Data Fragments

Conceptual View



Secondary servers


Physical View

Challenge I: Data Allocation

- Given a set of fragments, a set of sites on which a number of applications are running, **allocate** each fragment such that some optimization criterion is met (subject to certain constraints)
- It is known to be a NP-hard problem
 - The optimal solution depends on many factors
 - Location in which the query originates
 - The query processing strategies (e.g., join methods)
 - Furthermore, in a dynamic environment the workload and access pattern may change
- Most advanced approaches build *cost models* and any optimization algorithm can be adapted to solve it
 - Simplify the problem with assumptions (e.g., only communication cost considered)
 - Heuristics are also available: (e.g., best-fit for non-replicated fragments)
 - Sub-optimal solutions (i.e., applied per fragment individually)

Challenge I: Data Allocation

Conceptual View




Physical View

Challenge I: Data Allocation

Conceptual View



Apply an
allocation
strategy



Physical View


Challenge I: Data Allocation

Conceptual View

F2

F3

Apply an
allocation
strategy




Physical View

Challenge I: Data Allocation

Conceptual View

F3 Apply an
allocation
strategy




Physical View

Challenge I: Data Allocation


Conceptual View

Apply an
allocation
strategy



Challenge I: Data Allocation

Conceptual View




Physical View

Challenge I: Data Allocation

Conceptual View


If a fragment is placed in more than one server, then, we are **replicating** it




Physical View

Challenge I: Data Replication Management


- Replicating fragments improves the system throughput but raises some other issues:
 - Consistency
 - Update performance
- Most used replication protocols
 - Eager – Lazy replication
 - Primary – Secondary versioning




a) Eager replication with primary copy



b) Lazy replication with primary copy
(a.k.a Master-Slave replication)




c) Eager replication, distributed



d) Lazy replication, distributed
(a.k.a. Master-Master replication)

Challenge I: Data Replication Management

- Replicating fragments improves the system throughput but raises some other issues:
 - Consistency
 - Update performance
- Most used replication protocols
 - Eager – Lazy replication
 - Primary – Secondary versioning



Activity: Data Replication Issues


- Objective: Understand the consequences behind each data replication strategy
- Tasks:
 1. (10') By pairs, answer the following questions:
 - I. Discuss the questions below with your peer
 - II. What is the most important feature for each scenario?
 2. (5') Discussion
- You are a customer using an e-commerce based on heavy replication (e.g., Amazon):
 - a) Show a database replication strategy (e.g., sketch it) where:
 1. You buy an item, but this item does not appear in your basket.
 2. You reload the page: the item appears.What happened?
 - b) Show a database replication strategy (e.g., sketch it) where:
 1. You delete an item from your command, and add another one: the basket shows both items.What happened?
Will the situation change if you reload the page?

Challenge II



The Global Catalog



Challenge II: Global Catalog




Challenge II: Global Catalog




Physical View

Challenge II: Global Catalog




How to know this table is
fragmented and where the
fragments are?




Challenge II: Global Catalog

- Centralized version (@master)
 - Accessing it is a bottleneck
 - Single-point failure
 - May add a mirror
 - Poorer performance
- Distributed version (several *masters*)
 - Replica synchronization
 - Potential inconsistencies

Challenge II: Global Catalog




Conceptual View



Physical View

Challenge II: Global Catalog



Conceptual View


Catalog:

T <<fragmentation strategy>>

F1: @S1, @S2, @S4


F2: @S2, @S3, @S4

F3: @S1, @S3, @S4, @Sn



Physical View

Challenge II: Global Catalog



Conceptual View

Catalog:


T <<fragmentation strategy>>

F1: @S1, @S2, @S4

F2: @S2, @S3, @S4

F3: @S1, @S3, @S4, @Sn

This information is typically stored in a distributed index




Physical View

Challenge II: Global Catalog

Conceptual View

Centralized Version




Physical View


Challenge II: Global Catalog

Conceptual View

Primary server



Centralized Version




Secondary servers

Physical View





Challenge II: Global Catalog

Conceptual View


Primary server



Centralized Version




...



Secondary servers

Physical View


Challenge II: Global Catalog







Challenge II: Global Catalog

Conceptual View


Primary server



Distributed Version



...




Secondary servers

Physical View





Challenge II: Global Catalog

Conceptual View


Primary server



Distributed Version




...




Secondary servers

Physical View


Challenge II: Global Catalog



Challenge II: Global Catalog



Challenge II: Global Catalog



Summary

- Benefits of distributed systems and their relevance for NOSQL
- What is a distributed DBMS
- System architecture of a DDBMS
 - Distribution transparency
 - Replication transparency
 - Fragmentation transparency
- Distributed Database design
 - Data fragmentation
 - Data allocation
 - Data replication
- Global catalog management

Bibliography

- M.T. Özsü and P. Valduriez. *Principles of Distributed Database Systems*. Second edition. Prentice Hall, 1999
- Serge Abiteboul, Ioana Manolescu, Philippe Rigaux, Marie-Christine Rousset, Pierre Senellart. *Web Data Management*. Cambridge Press, 2011.
- L. Liu, M.T. Özsü (Eds.). *Encyclopedia of Database Systems*. Springer, 2009