

# Column-Oriented Databases

Exploiting Vertical Fragmentation to the Extreme

# Column-Oriented Databases

- Meant to accelerate read-only workloads
  - Use, to the extreme, of vertical partitioning
  - Do not allow variable record sizes and apply efficient compression techniques
  - Specific query processing techniques assuming the items above
    - In-memory query processing
- Recall the limitations yielded by vertical partitioning when reconstructing the original tuples
  - As such, these DBs are meant for **read-only databases**
  - Extremely inefficient in front of write-intensive database workloads

# Column-Oriented Databases: Types

- Column-oriented DBs are inherently aligned with decisional systems
  - A must for Data Warehousing! ~ Terabytes
  - A must for read-only Big Data Systems ~ Petabytes
- Two main types:
  - Relational Column-Oriented DBs (aka NewSQL)
    - First system: C-Store
    - Industrial examples: MonetDB, HP Vertica, SAP Hana, Oracle in-memory column store, MariaDB ColumnStore, PostgreSQL Zedstore... in general, ANY relational database provides, in one way or another, a columnar engine
  - Column-Oriented NOSQL databases
    - Apache Druid (first OLAP-like NOSQL engine) - first non-relational column-oriented DB
    - Hadoop Ecosystem
      - Apache Parquet and Apache Arrow file-formats for HDFS
      - Apache Kudu
    - Google BigQuery (former Dremel)
    - Amazon Redshift (former DynamoDB)
    - Snowflake

In general, most Cloud Providers provide a column-oriented PaaS

*NOTE: Many classify Apache HBase as a column-oriented DB. IT IS NOT. HBase applies a hybrid partitioning strategy with horizontal partiotining as its primary partitioning strategy. Therefore, it cannot apply to its whole the optimizations for query processing*

# Column-Oriented Databases: Features

- Column-Oriented Specific Features:
  - Data model
    - Pure Vertical Partitioning
    - Tuples are identified by their position (no PK needed to be replicated in each fragment)
      - Multiple sorting of data (if needed, different in each replica)
    - Remove variable size records and work with fix-sized records (dictionaries or bitmaps needed)
    - Column-specific compression techniques
  - Specific query processing
    - Late materialization: apply as many processing operators to the vertical fragments before joining them
    - Block iteration: exploit the fix-sized records to process data per blocks
    - *Vectorized query processing*: when late materialization and block iteration are combined
    - Specific join algorithms (e.g., invisible join) exploiting the previous items

# *Activity: Column-Oriented Specific Features*

- Objective: Understand the column-oriented specificities that tuned row-oriented databases cannot meet
- Tasks:
  1. (15') In group of two, each of you must read one of these specific features. Namely:
    - I. Compression
    - II. Late materialization and block iteration
  2. (5') Think tank

# Data Model: Vertical Partitioning

- Most column-oriented DBs create a fragment per column

**Table T**

	A	B	C
	Bravo	Lleida	A
	Bravo	BCN	A
	null	Girona	A
	null	Lleida	E
	null	Vic	C
	Charlie	Salt	E
	Charlie	BCN	E
	Charlie	BCN	A

# Data Model: Vertical Partitioning

- Most column-oriented DBs create a fragment PER column

Table T (partitioned)	A	B	C
	Bravo	Lleida	A
	Bravo	BCN	A
	null	Girona	A
	null	Lleida	E
	null	Vic	C
	Charlie	Salt	E
	Charlie	BCN	E
	Charlie	BCN	A

- Relevantly, some DBs allow to define groups of columns

# Data Model: Compression

- Each column is not stored as a regular file, but as a compressed file
- In these DBs, the compression main objective is not reducing data space but reducing I/Os
  - Yet, data in columnar format is more compressible than data stored in rows
    - High data value locality (less value entropy)
    - Benefits from sorting
- Two main trends
  - Heavy weight compression algorithms (e.g., Lempel-Ziv)
    - In general, not that useful but it might be if there is a (huge) gap between memory bandwidth and CPU performance
  - Lightweight compression (e.g., Run-Length Encoding) may allow the query optimizer work directly on compressed data
    - Improves performance by reducing CPU cost
    - Decompression is needed in front of bitwise AND / OR

# Data Model: Lightweight Compression


- Mainly based on dictionaries or bitmaps
- Example of Run-length Encoding (with dictionary)

Bravo
Bravo
null
null
null
Charlie
Charlie
Charlie

Column A

# Data Model: Lightweight Compression


- Mainly based on dictionaries or bitmaps
- Example of Run-length Encoding (with dictionary)




Column A

# Data Model: Lightweight Compression

- Mainly based on dictionaries or bitmaps
- Example of Run-length Encoding (with dictionary)



Column A




1
1
0
0
0
2
2
2

Dictionary  
positions


# Data Model: Lightweight Compression

- Mainly based on dictionaries or bitmaps
- Example of Run-length Encoding (with dictionary)




# Data Model: Lightweight Compression

- Mainly based on dictionaries or bitmaps
- Example of Run-length Encoding (with dictionary)



# Data Model: Lightweight Compression

- Mainly based on dictionaries or bitmaps
- Example of Run-length Encoding (with dictionary)



Realise that the ERI and Dictionary positions are always fixed-size

# Data Model: Reconstructing Records

- From each table,  $3 \times N$  vectors are generated ( $N = \# \text{attributes}$ )
- To reconstruct the original tuple, the original vertical fragmentation strategy we study used the PK (to be included in each fragment)
- **Alternatively**, in column-oriented DBs each record has the same position in all the columns (vectors) created from that table
  - Thus, column-oriented DBs do not join fragments to reconstruct the record
- Relevantly, a certain order might benefit a column and harm another
  - If necessary, in the presence of replicas, each replica might use a different order  
*(the order must be the same for all fragments of a table in a given replica)*

# *Activity: Data Model (I)*

- Objective: *Understand Run-Length Encoding*
- Tasks:
  1. (15') *Apply the Run-Length Encoding for the table in the next slide (dictionary-based + Ending-Row Index)*  
*For this exercise, do not play with the order of the rows. Use the one given. At the end, identify those columns where a different order might have generated a more compact representation*
  1. (5') *Think tank*

# *Activity: Data Model (I)*

BookID	Date	Price	#ItemsBought
1	1/01/2012	19,99	1
99	1/01/2012	9,99	1
301	2/01/2012	19,99	1
44	2/01/2012	9,99	1
56	2/01/2012	9,99	1
1	2/01/2012	19,99	1
77	3/01/2012	9,99	2
8	3/01/2012	19,99	1
78	3/01/2012	9,99	1
10	3/01/2012	19,99	1

# Data Model: Create Vertical Partitioning

- Creating a fragment per column, in general, is suboptimal. Thus, most advanced DBs allow the user to define partitions (each, potentially, as a set of columns)
- In these cases, finding the optimal vertical partitioning is a problem that must take into account the query workload
  - Identify correlations in the query workload and group those columns with high *affinity* (i.e., frequently queried together)
  - For  $m$  non-key attributes the search space is  $B(m)$ , i.e., the  $m$ th Bell number, which counts the possible partitions of a set of  $m$  elements
  - For large numbers,  $B(m) \approx m^m$
- Two main approaches to compute the optimal partitioning
  - Grouping
    - Attribute affinity matrix
    - Clustering algorithms
    - ...
  - Splitting
- [RECAP] Partitioning must guarantee
  - Completeness
  - Disjointness
  - Reconstruction

# Data Model: Create Vertical Partitioning

- Creating a fragment per column, in general, is suboptimal. Thus, most advanced DBs allow the user to define partitions (each, potentially, as a set of columns)
- In these cases, finding the optimal vertical partitioning is a problem that must take into account the query workload
  - Identify correlations in the query workload and group those columns with high *affinity* (i.e., frequently queried together)
  - For  $m$  non-key attributes the search space is  $B(m)$ , i.e., the  $m$ th Bell number, which counts the possible partitions of a set of  $m$  elements
  - For large numbers,  $B(m) \approx m^m$
- Two main approaches to compute the optimal partitioning
  - Grouping
    - Attribute affinity matrix
    - Clustering algorithms
    - ...
  - Splitting
- [RECAP] Partitioning must guarantee
  - Completeness
  - ~~Disjointness~~ ((for the sake of performance, some DBMS might sacrifice it))
  - Reconstruction

# Example

## Schema:

```
Compres(llibreIdFK, date, preu, numUnitats)
Llibre(llibreId, autor, any, editorial, ISBN)
```

## Queries:

```
SELECT llibreId, SUM(numUnitats) FROM compres c, llibre l
WHERE c.llibreId = l.llibreId AND editorial = 'RBA'
GROUP BY llibreId
```

```
SELECT editorial, AVG(preu) FROM compres c, llibres l
WHERE c.llibreId = l.llibreId
GROUP BY editorial
```

```
SELECT AVG(numUnitats) FROM compres c
WHERE date BETWEEN '01/01/yyyy' AND '31/12/yyyy'
GROUP BY llibreID
```

```
SELECT autor, any, COUNT(*) FROM llibre l
GROUP BY autor, any
```



# Attribute Affinity Matrix

- Algorithm:
  1. For each relation, generate the **attribute usage matrix**

	<i>llibreId</i>	<i>date</i>	<i>preu</i>	<i>numUnitats</i>
<i>Q1</i>	1	0	0	1
<i>Q2</i>	1	0	1	0
<i>Q3</i>	1	1	0	1
<i>Q4</i>	0	0	0	0

Now, create the attribute usage matrix for *Llibre*

# Attribute Affinity Matrix

- Algorithm:
  2. For each relation, generate the **attribute affinity matrix**
    - Consider Q1 frequency is 50%, Q2 10%, Q3 30%, Q4 10%
    - For a pair of attributes  $A_i, A_j$ , compute its affinity by adding up all the frequencies in which they appear together

	<i>llibreId</i>	<i>date</i>	<i>preu</i>	<i>numUnitats</i>
<i>Q1</i>	1	0	0	1
<i>Q2</i>	1	0	1	0
<i>Q3</i>	1	1	0	1
<i>Q4</i>	0	0	0	0

*Step 1 Output: Query-Attribute Usage Matrix*

Now, create the attribute affinity matrix for *Llibre*

# Attribute Affinity Matrix

- Algorithm:
  2. For each relation, generate the **attribute affinity matrix**
    - Consider Q1 frequency is 50%, Q2 10%, Q3 30%, Q4 10%
    - For a pair of attributes  $A_i, A_j$ , compute its affinity by adding up all the frequencies in which they appear together

	<i>llibreId</i>	<i>date</i>	<i>preu</i>	<i>numUnitats</i>
<i>Q1</i>	1	0	0	1
<i>Q2</i>	1	0	1	0
<i>Q3</i>	1	1	0	1
<i>Q4</i>	0	0	0	0

Step 1 Output: Query-Attribute Usage Matrix

	<i>llibreId</i>	<i>date</i>	<i>preu</i>	<i>numUnitats</i>
<i>llibreId</i>	90	30	10	80
<i>date</i>	30	30	0	30
<i>preu</i>	10	0	10	0
<i>numUnitats</i>	80	30	0	80

Step 2 Output: Query-Attribute Affinity Matrix

Now, create the attribute affinity matrix for *Llibre*

# Attribute Affinity Matrix

- Algorithm:
  3. For each matrix, reorganize the attribute orders to form clusters where the attributes in each cluster show high affinity to one another

	<i>llibreId</i>	<i>numUnitats</i>	<i>preu</i>	<i>date</i>
<i>llibreId</i>	90	80	10	30
<i>numUnitats</i>	80	80	0	30
<i>preu</i>	10	0	10	0
<i>date</i>	30	30	0	30

Now, do the same for *Llibres*

# Attribute Affinity Matrix

- Algorithm:
  3. For each matrix, reorganize the attribute orders to form clusters where the attributes in each cluster show high affinity to one another

	<i>llibreId</i>	<i>numUnitats</i>	<i>preu</i>	<i>date</i>
<i>llibreId</i>	90	80	10	30
<i>numUnitats</i>	80	80	0	30
<i>preu</i>	10	0	10	0
<i>date</i>	30	30	0	30

Now, do the same for *Llibres*

# Attribute Affinity Matrix

- Output:
  - P1: LlibreID, numUnitats
  - P2: Preu
  - P3: Date
- Assess the result: compute the **effective read ratio**
  - For each query, for each partition it must read, compute the following ratio:

$$\#Attr(Q_j, P_i) / \#Attr(P_i)$$

Where  $\#Attr(Q_j, P_i)$  means the number of attributes  $Q_j$  must read from  $P_i$  and  $\#Attr(P_i)$  is the total number of attributes in  $P_i$

- Example:  $RR(Q_1, P_1) : 2/2 = 1$


# Data Processing: Late Materialization

- *Tuple reconstruction* can be done at the beginning or at the end of the query

# Data Processing: Late Materialization

- *Tuple reconstruction* can be done at the beginning or at the end of the query


Beginning:




# Data Processing: Late Materialization

- *Tuple reconstruction* can be done at the beginning or at the end of the query

Beginning:




End:



# Data Processing: Late Materialization

- *Tuple reconstruction* can be done at the beginning or at the end of the query



- Advantages of reconstructing the tuple at the end:
  - Some tuples do not need to be constructed (because of selections and projections)
  - Some columns remain compressed more time
  - Cache performance is improved (kept at column level)
  - Helps block iteration for values of fixed length columns

# Data Processing: Block Iteration

- Blocks of values of the same column are passed to the next operation in a single function call
- Values inside the block can be:
  - Iterated as in an array (fixed-width)
  - Remain compressed together
    - Not necessarily using multiples of 8 bits
    - Counting or even identifying the tuples for which the predicate is true
  - Exploits parallelism / pipelining

1	1
4	0
7	2

Ending Row Index      Dictionary positions

# Summary

- Advantages of column-oriented databases
  - Bring into memory only relevant data
  - Provide fewer and simpler internal functions
  - Easier to recognize all execution strategies
  - Simpler tuning required by users
- Pioneers: Daniel Abadi & Michael Stonebraker (M.I.T.)
  - C-Store (blueprint)
  - MonetDB (academic edition, open-source)
    - <http://www.monetdb.org/Home>
  - Vertica (commercial version, bought by HP)
    - <http://www.vertica.com/>