


Key-Value Stores

Very Large Databases

The Complexity of Big Data Management




HDFS, HBase and MapReduce

An Example of Key-Value Ecosystem


Key-Values: A Piece of History

- Key-values were born as a desperate answer to the RDBMS limitations
 - It is widely assumed that Google is the father of Key-value stores
- High-performance is mainly achieved by means of parallelism
 - To achieve large-scale parallelism, data is distributed across a cluster
 - Google File System (GFS)
 - GFS + Bigtable
 - MapReduce
 - It is a “query language” that provides parallelism in a transparent manner



Key-Values: A Piece of History

- Hadoop Distributed File System (HDFS)
 - Based on *The Google File System* (2003)
- Hadoop MapReduce
 - Based on *MapReduce: Simplified Data Processing on Large Clusters* (2004)
- HBase
 - Based on *Bigtable: A Distributed Storage System for Structured Data* (2006)



The File System: HDFS

A Distributed File System

What is a file system?

The image shows two windows side-by-side. On the left is a Windows File Explorer window titled 'Teaching > BDM > Slides'. It displays a list of files and folders in a table format. The columns are 'Name', 'Date modified', 'Type', and 'Size'. The 'Name' column lists 'resources', '00-Introduction_withFonts', '01-BigDataDesign_withFonts', '02-DistributedDataManagement_withFo...', '03-DistributedDataProcessing_withFo...', and '04-HDFS'. The 'Date modified' column shows dates from 07/02/2020 to 22/02/2020. The 'Type' column indicates all are 'File folder' except for the first one which is 'File'. The 'Size' column shows sizes ranging from 6.078 KB to 12.108 KB. The '04-HDFS' file is selected. At the bottom of the File Explorer window, there is a status bar with the following information: '04-HDFS', 'Title: BDM-HDFS', 'Size: 6.15 MB', 'Tags: Add a tag', 'Microsoft PowerPoint Presentation', 'Authors: Sergi Nadal', and 'Date modified: 22/02/2020 10:02'. On the right is a terminal window titled 'cmd.exe' with the path 'C:\Windows\system32\cmd.exe'. The terminal output shows the following command and its results:

```
C:\Windows\system32\cmd.exe
Microsoft Windows [Version 6.1.7601]
Copyright (c) 2009 Microsoft Corporation. All rights reserved.

C:\Users\sergi>cd Desktop\Teaching\BDM\Slides
C:\Users\sergi\Desktop\Teaching\BDM\Slides>dir /p
Volume in drive C has no label.
Volume Serial Number is 5C20-684F


Directory of C:\Users\sergi\Desktop\Teaching\BDM\Slides

22/02/2020  10:02    <DIR>      .
22/02/2020  10:02    <DIR>      ..
07/02/2020  12:38    12.397.595 00-Introduction_withFonts.pptx
17/02/2020  07:45      6.223.187 01-BigDataDesign_withFonts.pptx
15/02/2020  15:21      7.876.848 02-DistributedDataManagement_withFonts.pptx
20/02/2020  09:51      5.059.894 03-DistributedDataProcessing_withFonts.pptx
22/02/2020  10:02      6.452.040 04-HDFS.pptx
22/02/2020  10:02    <DIR>      resources
                           5 File(s)   38.009.564 bytes
                           3 Dir(s)  20.948.156.416 bytes free

C:\Users\sergi\Desktop\Teaching\BDM\Slides>
```

Functionalities provided by a File System

- Creates a hierarchical structure of data
 - Splits and stores data into files and blocks
- Provides interfaces to read/write/delete
- Maintains directories/files metadata
 - Size, date of creation, permissions, ...



Hadoop File System (HDFS)


- Apache Hadoop (<http://hadoop.apache.org>)
 - Based on Google File System (GFS)
- Designed to meet the following requirements
 - Handle many and very large files
 - Data collections are written once and read many times (sequential access pattern)
 - Infrastructure consisting of thousands of inexpensive commodity machines with high failure probability
- Traditional network file systems do partially fulfil these requirements
 - File System Vs. Database Management System
 - Balancing *query* load (e.g., by means of *fragmentation and replication*) boosts availability and reliability
 - HDFS: Equal-sized file chunks evenly distributed

Chunks are
“equivalent” to blocks
in a centralized system

HDFS in a Nutshell (I)

- A single master (coordinator)
 - Receives client connections
 - Maintains the description of the global file system namespace
 - Keeps track of file chunks (default: 128Mb)
- Many servers
 - Receive file chunks and store them
- A single master design forfeits availability and scalability
 - Availability and reliability: Recovery system
 - Replication (a chunk always in 3 servers, by default)
 - Monitors the system with heartbeat messages to detect failures as soon as possible
 - Specific recovery system to protect the master
 - Scalability: Client cache

HDFS in a Nutshell (II)



HDFS File Formats

HDFS provides different file formats to store data:

Text / CSV / JSON formats

- Do not support block compression
- Fix-sized fragments with no metadata
- Not advisable for advanced processing
- Typically used to store *raw* data

Avro

- Native row-oriented storage
- Rich header metadata (including schema information)
- Supports block compression techniques
- Recognised by most popular processing engines, such as Spark
 - It allows SQL-like querying (e.g., Spark SQL)

Parquet

- Native column-oriented storage
- Rich header metadata (including schema information and statistics)
- Supports advanced block compression techniques
- Recognised by most popular processing engines, such as Spark
 - It allows SQL-like querying (e.g., Spark SQL)
 - Selections and Projections can be pushed down
 - Statistics used to skip whole fragments

... and many others: Arrow, ORC Files, Sequence Files, etc.

The Database: HBase

Richer Data Structure, Recovery, Concurrency and Indexing Capabilities


HBase

- Apache HBase (<http://hbase.apache.org>)
 - Based on Google's BigTable
- Designed to meet the following requirements
 - Access specific data out of petabytes
 - It must support key search, range search and high throughput for scans
- With regard to HDFS, it provides:
 - Richer key-value structure (hybrid fragmentation strategy: horizontal + vertical)
 - Indexing capabilities (to enable efficient key and range searches)
 - Concurrency
 - Recovery
- From an architecture point of view, it is a distributed index cluster on top of HDFS
 - Distributed clustered B+
 - Tuples are lexicographically sorted according to the key

Provides random
access capabilities on
top of HDFS

Schema Elements

- Stores tables (collections) and rows (instances)
 - A cell is identified by row key, column family, and column (arbitrary strings)
- Treats data as uninterpreted strings (without data types)
- Each cell of a BigTable can contain multiple versions of the same data
 - Stores different versions of the same values in the rows
 - Each version is identified by a timestamp
 - Timestamps can be explicitly or automatically assigned




Key-Value

- Key-value stores
 - Entries in form of key-values
 - One key maps only to one value
 - Query on key only
 - Schemaless



- HBase (Wide-column key-value store)
 - Entries in form of key-values
 - But now values are split in wide-columns (column families)
 - Typically query on key, but some support for values (columns)
 - Schemaless within a family



Columns are defined at insertion time

HBase Shell

- ALTER <tablename>, <columnfamilyparam>
- COUNT <tablename>
- CREATE TABLE <tablename>
- DESCRIBE <tablename>
- DELETE <tablename>, <rowkey>[, <columns>]
- DELETEALL <tablename>, <rowkey>
- DROP <tablename>
- EXIT
- EXISTS <tablename>
- GET <tablename>, <rowkey>[, <columns>]
- LIST
- PUT <tablename>, <rowkey>, <columnid>, <value>[, <timestamp>]
- SCAN <tablename>[, <startrow>, <stoprow>, <columns>]
- STATUS [{summary|simple|detailed}]
- SHUTDOWN


SCAN is parallel!

Activity: Key-Value Design


- Objective: Learn the basic design principles of key-value stores

- Tasks:

1. (20') By pairs, solve the following exercise
 - Model in HBase the Lineitem and Orders tables
2. (5') Discussion



No JOINS in HBase!



TPC-H Benchmark

Activity: Key-Value Design

- Objective: Learn the basic design principles of key-value stores

- Tasks:

1. (20') By pairs, solve the following exercise
 - Model in HBase the Lineitem and Orders tables
2. (5') Discussion


```
SELECT l_orderkey, sum(l_extendedprice*(1-l_discount)) as revenue,  
o_orderdate, o_shippriority
```

```
FROM customer, orders, lineitem
```

```
WHERE c_mktsegment = '[SEGMENT]' AND c_custkey = o_custkey AND  
l_orderkey = o_orderkey AND o_orderdate < '[DATE]' AND l_shipdate > '[DATE]'
```


```
GROUP BY l_orderkey, o_orderdate, o_shippriority
```

```
ORDER BY revenue desc, o_orderdate;
```




TPC-H Benchmark

Just Another Point of View



Physical Implementation




- It is a layer on top of HDFS (HDFS stores data)
- Each table is horizontally fragmented into *Regions* (*tablets in BigTable*)
 - Dynamic fragmentation
 - By default into few hundreds of MBs (system parameter)
 - Distributed on a cluster of machines or cloud
- At each Region rows are stored column-wise according to families (**hybrid fragmentation**)
 - In disk, as many files as families
 - Static fragmentation (determined by)
 - Block compression can be enabled

Implications on modeling!

- Massive usage of in-memory storage
 - Files in disk have an in-memory counterpart
 - Changes happen in main memory and are not flushed to disk until a **compaction** happens
 - A compaction merges the in-memory and disk files
- Metadata table (stores the **global catalog**)
 - It is the physical implementation of the catalog
 - Tuples are lexicographically sorted according to the key
 - Each row (entry) consists of <key, loc>
 - Key: it is the last key value in *that Region*
 - Loc: it is the physical address of a Region
 - LSM-tree index: combination of in-memory and disk-based structures

HBase: A Distributed Index Cluster




Serge Abiteboul et al.

Functional Components (I)

- Zookeeper
 - Open-source server for highly reliable distributed coordination of cloud applications
 - Stores the HBase cluster config info
- HMaster
 - Coordinates splitting of regions/rows across nodes
 - Controls distribution of HFile chunks
- Region Servers (HRegionServer)
 - Store regions (i.e., horizontal fragments)
 - Within the store region, the horizontal fragment is further fragmented vertically
 - Logs changes
 - Guarantees “atomic” updates
 - Holds (caches) chunks of HFile into Memstores, waiting to be written
- HDFS
 - Stores all data including columns and logs
 - DataNodes store chunks of a file
 - NameNode holds all metadata including namespace
 - HBase uses two HDFS file types
 - HFile: regular data files (holds column data)
 - HLog: region’s log file (allows flush/fsync for small append-style writes)
- Clients
 - Read and write chunks
 - Locality & load determine which copy to access

Functional Components (II)

Zookeeper for configuration and global coordination



- A primary copy must be stored in the same DataNode the HRegionServer runs on.
- Secondary copies can be stored in any DataNode different from the DataNode the HRegionServer runs on.
- All the stores of a given family correspond to the same table as this family.

$n \xrightarrow{m} A$ An element of class B is associated with n elements of A, and viceversa

$B \xrightarrow{\Delta} A$ Class B is a specialization (subtype) of class A


$B \xrightarrow{\blacklozenge} A$ Class A is composed by elements of class B

HBase Design Decisions (I)

- One master server
 - Responsible to maintain the table schemas and their fragments (i.e., catalog)
 - Assignment of regions to servers
 - Monitors the cluster (*heartbeating*)
- Many region servers
 - Each handling around 100-1000 regions (i.e., horizontal fragments)
 - Further apply vertical fragmentation (families) to regions
 - Apply concurrency and recovery techniques
 - Managing split of regions
 - A region server decides to split (e.g., >128MB)
 - Half of its regions are sent to another server
 - Managing merge of regions
- Client nodes



HBase Design Decisions (II)

- Split and merge affects the distributed tree, which must be updated
 - *Gossiping*
 - Lazy updates: discrepancies may cause out-of-range errors, which triggers a stabilization (**mistake compensation**) protocol
- Mistake compensation
 - The client keeps in cache the tree sent by the master and uses it to access data
 - If an out-of-range error is triggered, it is forwarded to the parent
 - In the worst case, 6 network round trips



Activity: Mistake Compensation

- Objective: Understand how the global catalog is handled in HBase
- Tasks:
 1. (10') By pairs, solve the following exercise
 - a) Number of round trips if Get(2602)
 - b) Number of round trips if Get(2602) again
 2. (5') Discussion



HBase in a Nutshell

- HBase was thought to provide a specific answer for a specific problem
 - Architecture: Distributed index cluster
 - In-memory sorted map for new rows (i.e., heavy in-memory processing)
 - Merge-sort for merging it with *old* data
 - LSM-tree: a distributed B+ able to deal with in-memory and disk data
 - Richer key-value data structure: Table / family / row
 - Hybrid fragmentation (first horizontal, then vertical)
 - Thus, it cannot fully benefit from column-oriented processing
 - Concurrency: Timestamping MVCC (Multiversion Concurrency Control)
 - Guarantees ACID semantics per-row (even across families)
 - Recovery: Logging is used at the Region Servers
 - WAL (write-ahead protocol)
 - REDO log (No Force / No Steal)

HBase Architecture

- Refreshing the NOSQL Challenges

- I. Distributed DB design

- Data fragments
 - Data replication
 - Node distribution

Horizontal fragmentation (fixed-size chunks)

Secondary vertical fragmentation (not allowing HBase to benefit from column-oriented processing)

Replication performed by HDFS: Eager / primary copy

Load balancing. Tuneable, by default depends on #RegionServers and size of the family file

- II. Distributed DB catalog

- Fragmentation trade-off: Where to place the DB catalog
 - Global or local for each node
 - Centralized in a single node or distributed
 - Single-copy vs. Multi-copy

Global catalog: distributed tree

Clustered data

Centralized and multi-copy catalog (if several masters)

Eager replication / secondary copy between the catalog copies

- III. Distributed query processing

- Data distribution / replication
 - Communication overhead

Not covered by HBase but done by MapReduce / Spark

HBase API only covers single-key or range key searches

- IV. Distributed transaction management

- How to enforce the ACID properties
 - Replication trade-off: Queries vs. Data consistency between replicas (updates)
 - Distributed recovery system
 - Distributed concurrency control system

Concurrency: per-row guarantees only

Recovery: checkpointing logging per Region Server

- V. Security issues

- Network security

Nothing!

HBase Architecture

- NOSQL Goals
 - Schemaless: No explicit schema [column-family key-value]
 - Reliability / availability: Keep delivering service even if its software or hardware components fail [recovery] / [distribution]
 - Scalability: Continuously evolve to support a growing amount of tasks [distribution]
 - Efficiency: How well the system performs, usually measured in terms of response *time* (latency) and *throughput* (bandwidth) [distribution]

Summary

- The Hadoop Ecosystem: Storage Layer
 - HDFS, HBase
- The File System (HDFS)
 - Architecture: functional components
 - File formats
- The Database (HBase)
 - Schema elements
 - Data model:
 - Data structure: Key-value
 - Operations: GET, SCAN
 - Constraints: NONE!
 - Architecture: functional components
 - Design decisions

Bibliography

- F. Chang et all. *Bigtable: A Distributed Storage System for Structured Data*. OSDI'06
- Sanjay Ghemawat et al. *The Google File System*. SOSP'03
- D. Jiang et al. *The performance of MapReduce: An In-depth Study*. VLDB'10