

# MapReduce

A Processing Framework for Massive Parallelism



# Motivation

Example: Card Counting

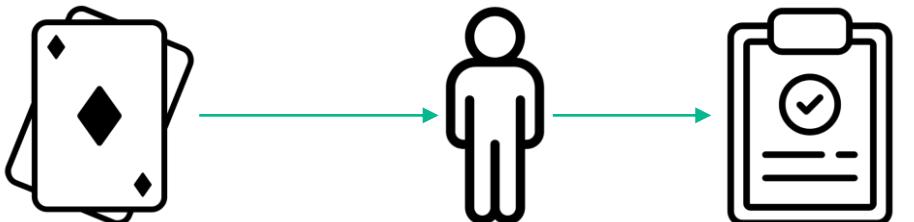
Input:

A bunch of cards

Result:

$A\spades \times 13, 2\spades \times 15, 3\spades \times 16, \dots, 3\hearts \times 12; 2\hearts \times 21; A\hearts \times 15$

How would you do it?



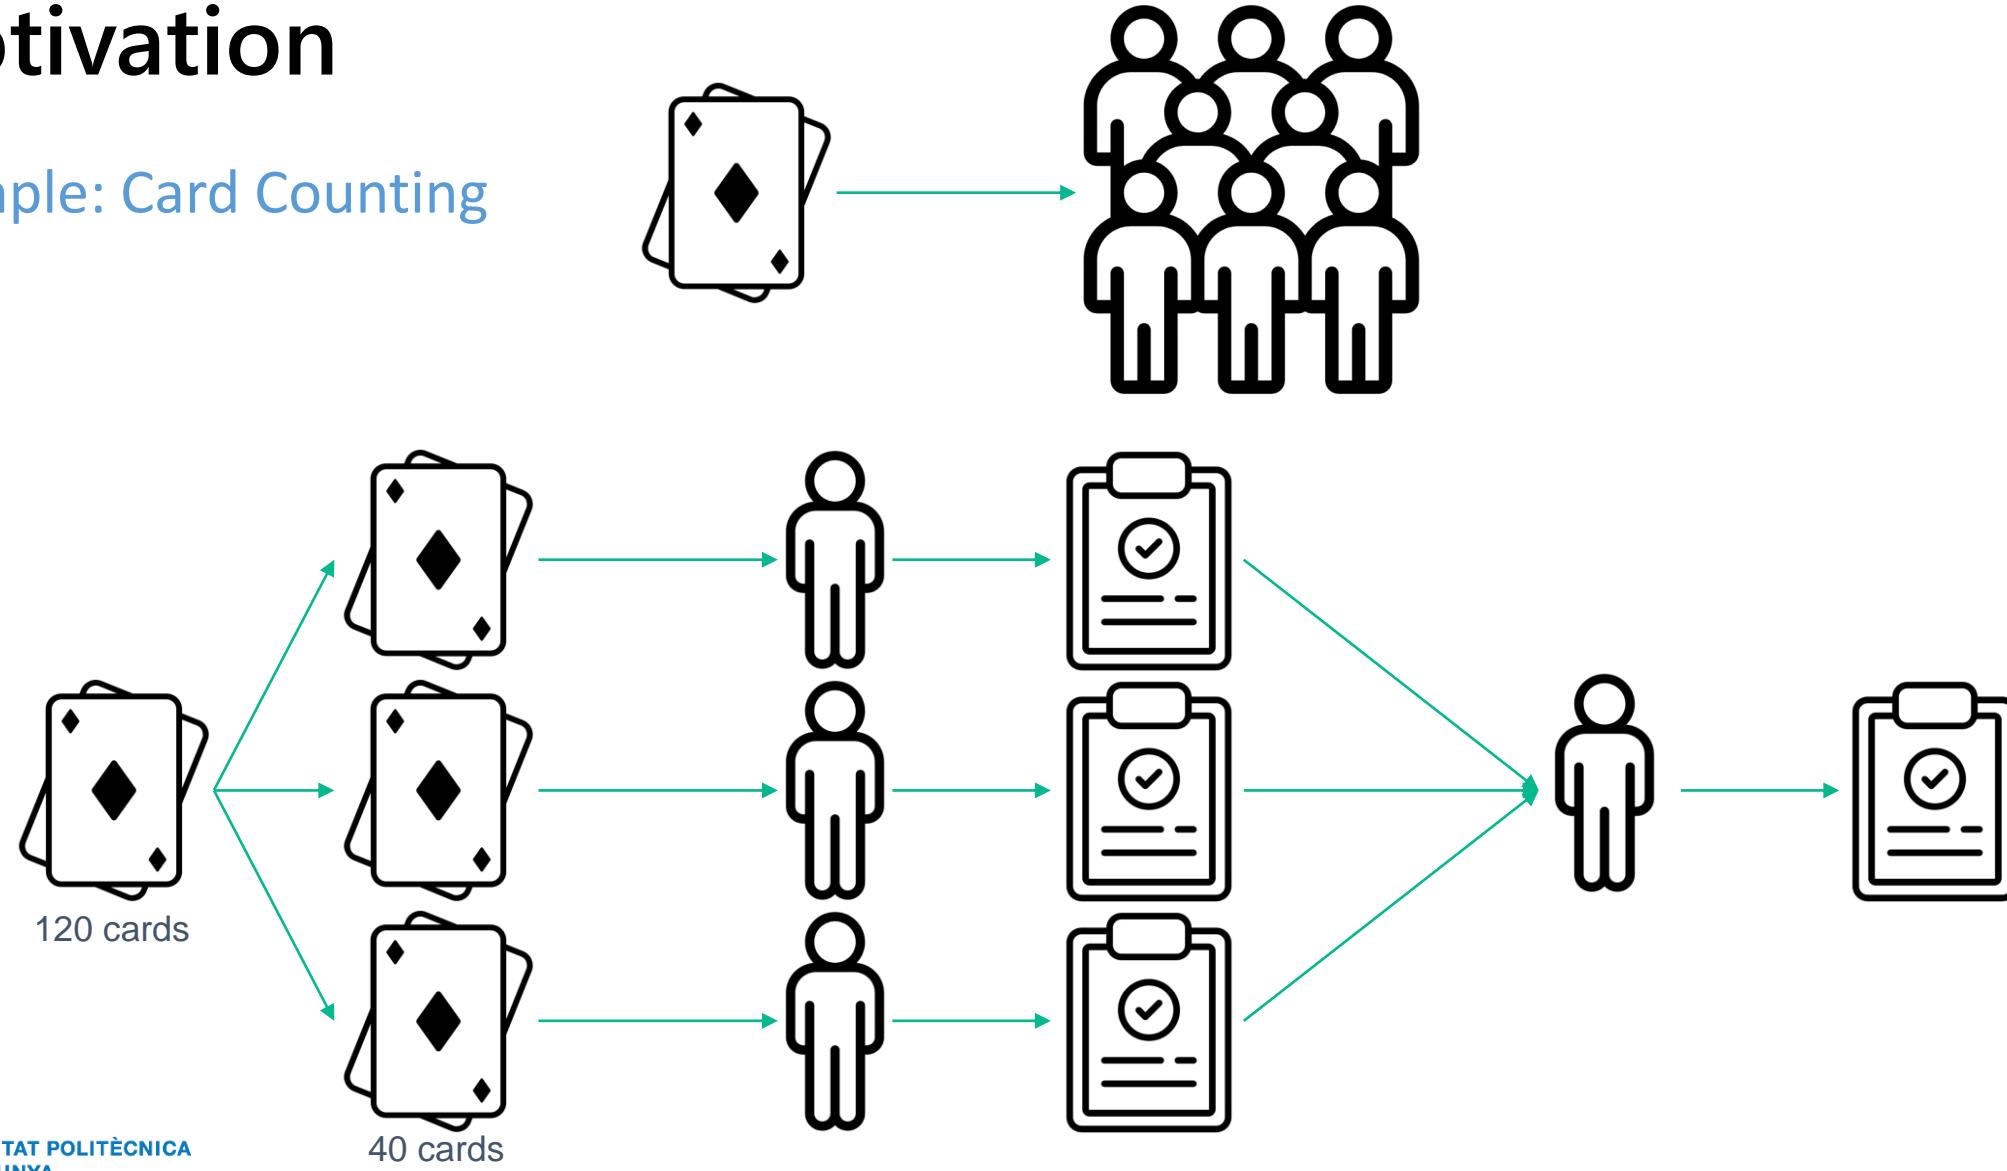
(sequential execution)



Photo by [Angshu Purkait](#) on Unsplash

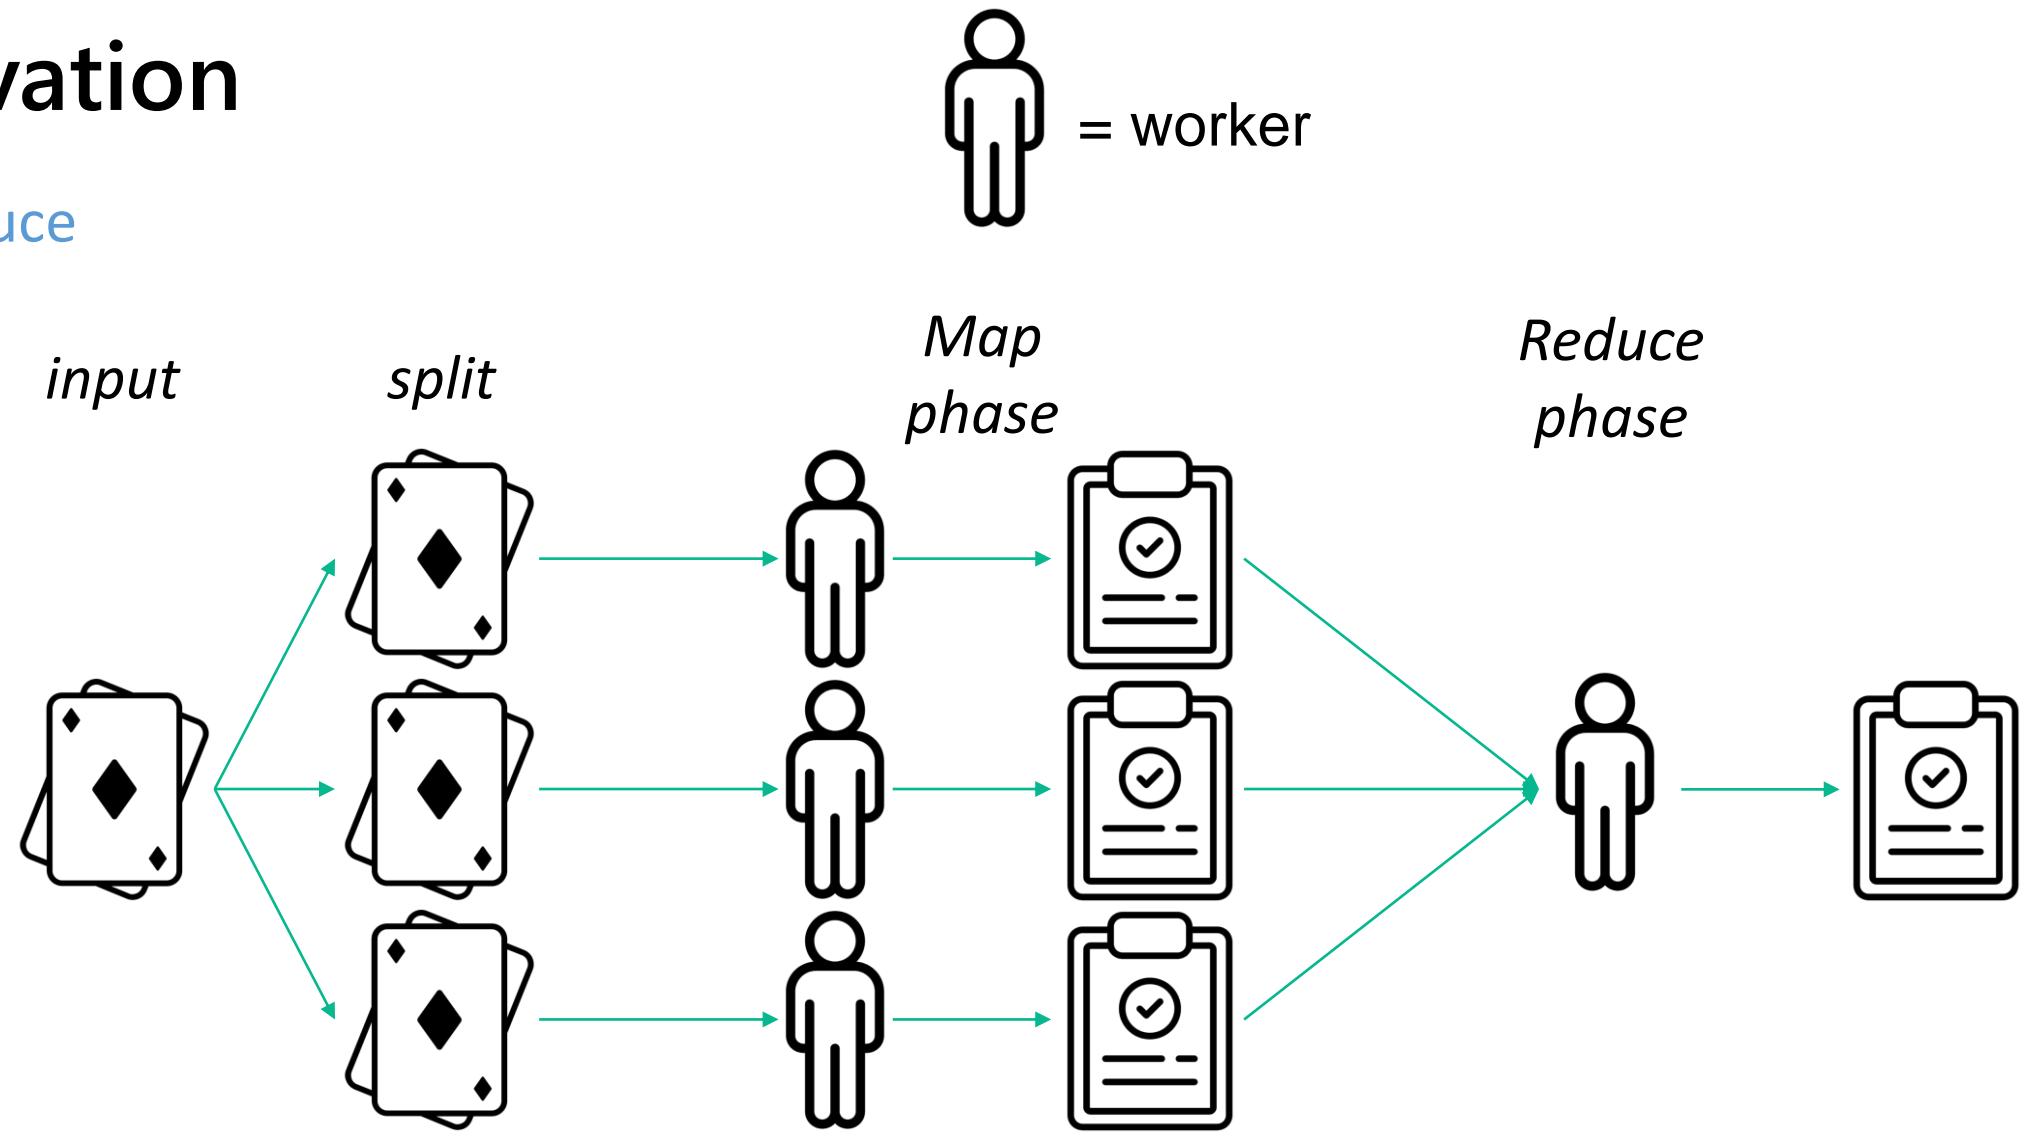
# Motivation

Example: Card Counting



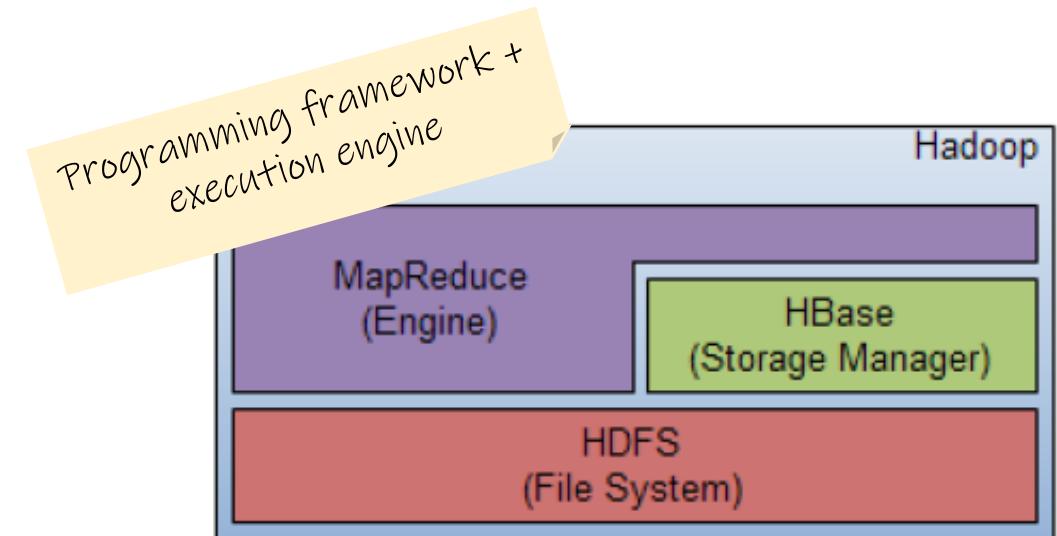
# Motivation

## MapReduce

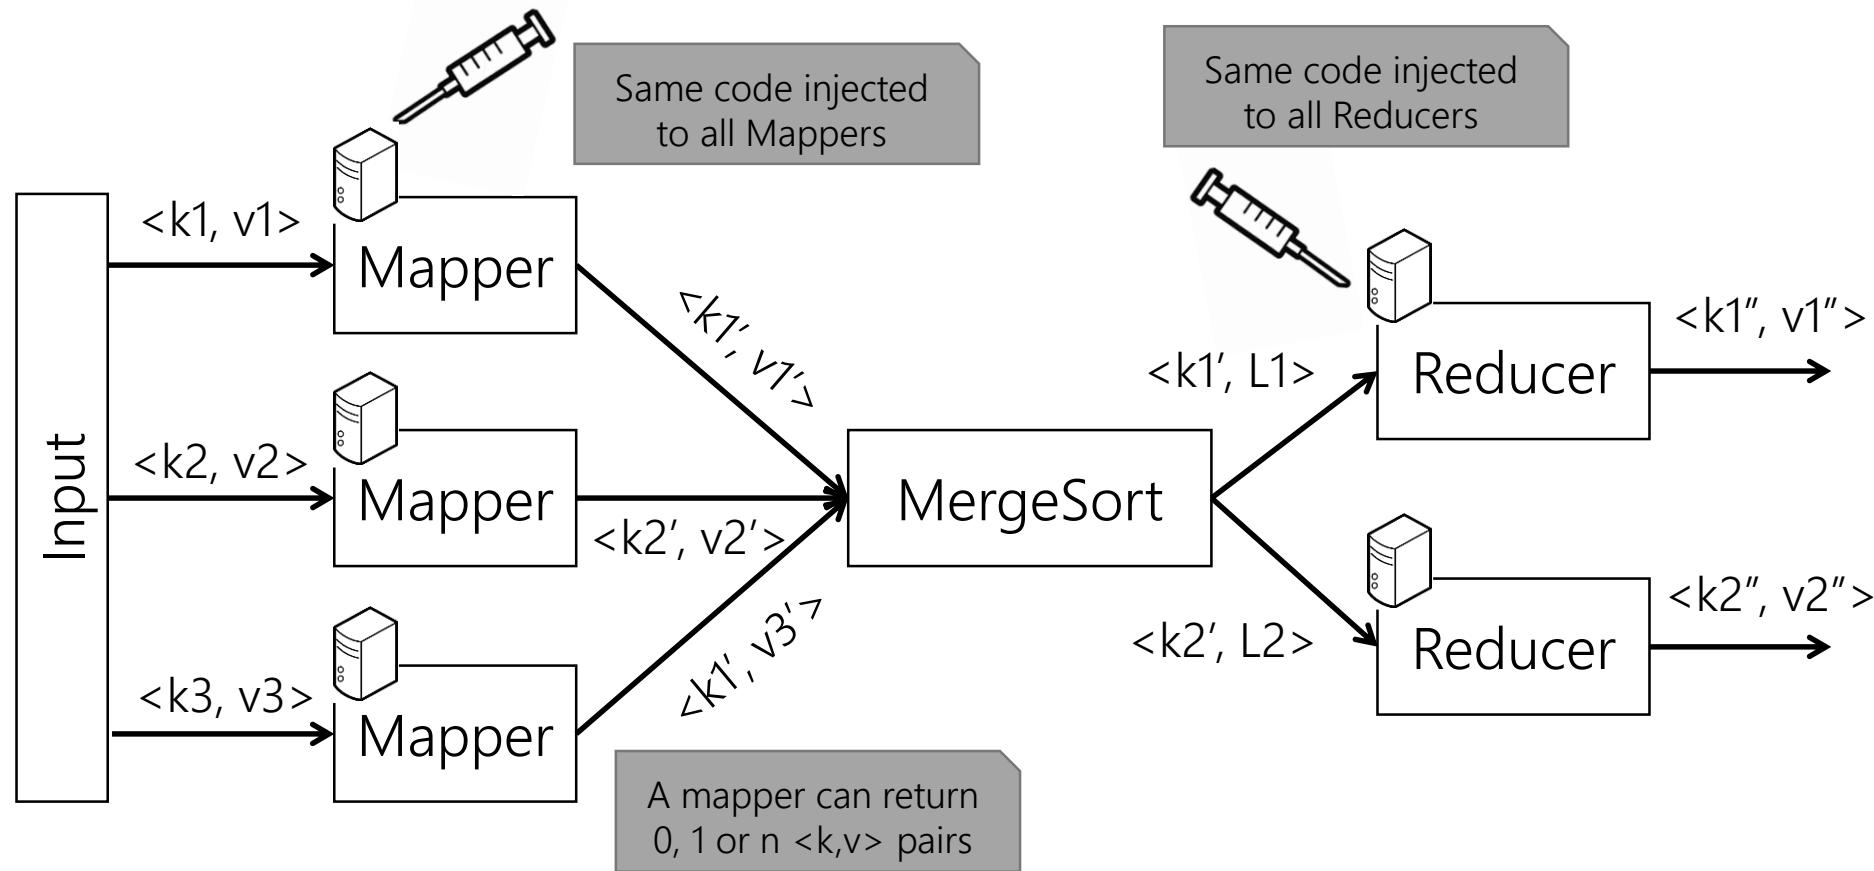


# MapReduce

- A piece of history...
  - Appeared together with Google File System (GFS) as a processing framework
  - Its main objective is to enable massive parallelism over distributed data
- Designed to meet the following requirements
  - Exploit distributed systems and provide full distributed-transparency for the end-user
    - Send the queries to data (i.e., query-shipping instead of data-shipping for exploiting the data locality principle)
    - Support parallelism and hide its complexity
      - Independent data (typically collected from the web)
        - Without references to other pieces of data
        - No joins
      - Exploit petabytes of data in batch mode
        - No transactions
    - Failure resilience
      - Cope with failures without aborting
  - Inspired in functional programming
    - Functions, immutability



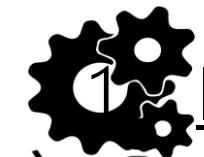
# The MapReduce framework



# The MapReduce framework in detail

1. Input: read input from distributed storage
2. Map: for each input  $\langle \text{key}_{\text{in}}, \text{value}_{\text{in}} \rangle$ , generate zero-to-many  $\langle \text{key}_{\text{map}}, \text{value}_{\text{map}} \rangle$
3. Partition: assign each  $\text{key}_{\text{map}}$  to a reducer machine
4. Shuffle: data are shipped to reducer machines using the distributed storage
5. Sort&Merge: reducers sort their input data by key
6. Reduce: for each  $\text{key}_{\text{map}}$ , the set  $\text{value}_{\text{map}}$  is processed to produce zero-to-many  $\langle \text{key}_{\text{red}}, \text{value}_{\text{red}} \rangle$
7. Output: writes the result of reducers to distributed storage

# The MapReduce framework in detail



Input: read input from distributed storage



Map: for each input  $\langle \text{key}_{\text{in}}, \text{value}_{\text{in}} \rangle$ , generate zero-to-many  $\langle \text{key}_{\text{map}}, \text{value}_{\text{map}} \rangle$



Partition: assign each  $\text{key}_{\text{map}}$  to a reducer machine



Shuffle: data are shipped to reducer machines using the distributed storage



Sort&Merge: reducers sort their input data by key



Reduce: for each  $\text{key}_{\text{map}}$ , the set  $\text{value}_{\text{map}}$  is processed to produce zero-to-many  $\langle \text{key}_{\text{red}}, \text{value}_{\text{red}} \rangle$



Output: writes the result of reducers to distributed storage

# MapReduce example

Word count

# Word count example

The Project Gutenberg eBook of The Outline of Science, Vol. 1 (of 4), by J. Arthur Thomson

This eBook is for the use of anyone anywhere at no cost and with almost no restrictions whatsoever. You may copy it, give it away or re-use it under the terms of the Project Gutenberg License included with this eBook or online at [www.gutenberg.org](http://www.gutenberg.org)

Title: The Outline of Science, Vol. 1 (of 4)  
A Plain Story Simply Told

Author: J. Arthur Thomson

Release Date: January 22, 2007 [EBook #20417]

Language: English

Character set encoding: ASCII

\*\*\* START OF THIS PROJECT GUTENBERG EBOOK OUTLINE OF SCIENCE \*\*\*

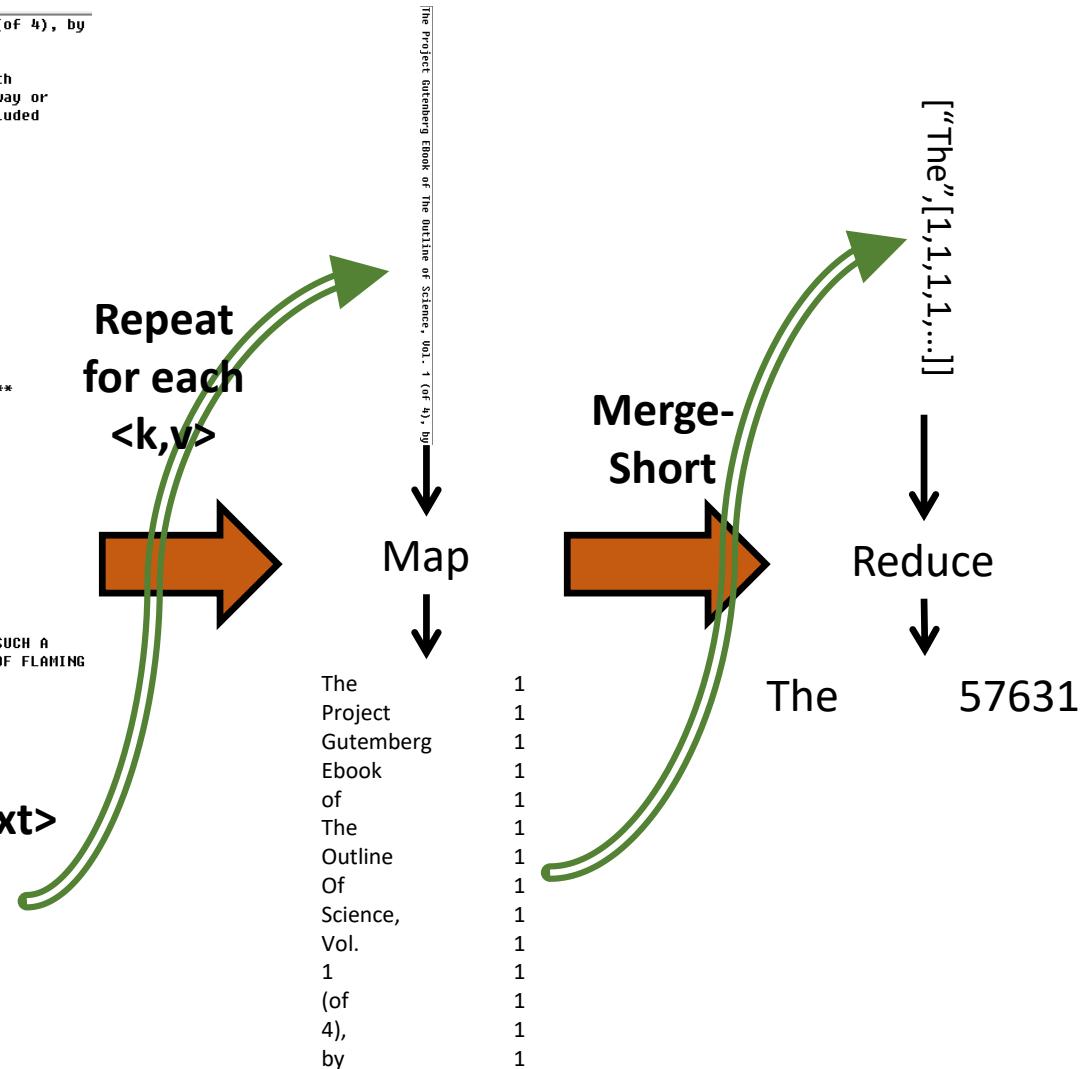
Produced by Brian Janes, Leonard Johnson and the Online Distributed Proofreading Team at <http://www.pgdp.net>

[ILLUSTRATION: THE GREAT SCARLET SOLAR PROMINENCES, WHICH ARE SUCH A NOTABLE FEATURE OF THE SOLAR PHENOMENA, ARE IMMENSE OUTBURSTS OF FLAMING HYDROGEN RISING SOMETIMES TO A HEIGHT OF 500,000 MILES]

THE  
OUTLINE OF SCIENCE  
A PLAIN STORY SIMPLY TOLD

EDITED BY  
J. ARTHUR THOMSON  
REGIUS PROFESSOR OF NATURAL HISTORY IN THE  
UNIVERSITY OF ABERDEEN

WITH OVER 800 ILLUSTRATIONS  
OF WHICH ABOUT 40 ARE IN COLOUR



# WordCount Code Example

```
public void map(LongWritable key, Text value) {  
    String line = value.toString();  
    StringTokenizer tokenizer = new StringTokenizer(line);  
    while (tokenizer.hasMoreTokens()) {  
        write(new Text(tokenizer.nextToken()), new IntWritable(1));  
    }  
}  
  
public void reduce(Text key, Iterable<IntWritable> values) {  
    int sum = 0;  
    for (IntWritable val : values) {  
        sum += val.get();  
    }  
    write(key, new IntWritable(sum));  
}
```

IntWritable, LongWritable, ...  
are variants of data types  
that are optimized for  
serialization in the Hadoop  
environment

# WordCount Code Example

```
public void map Key Value {
```

Blackbox

```
    write Key Value
```

```
}
```

```
}
```

```
public void reduce Key {Values} {
```

Blackbox

```
    write Key Value
```

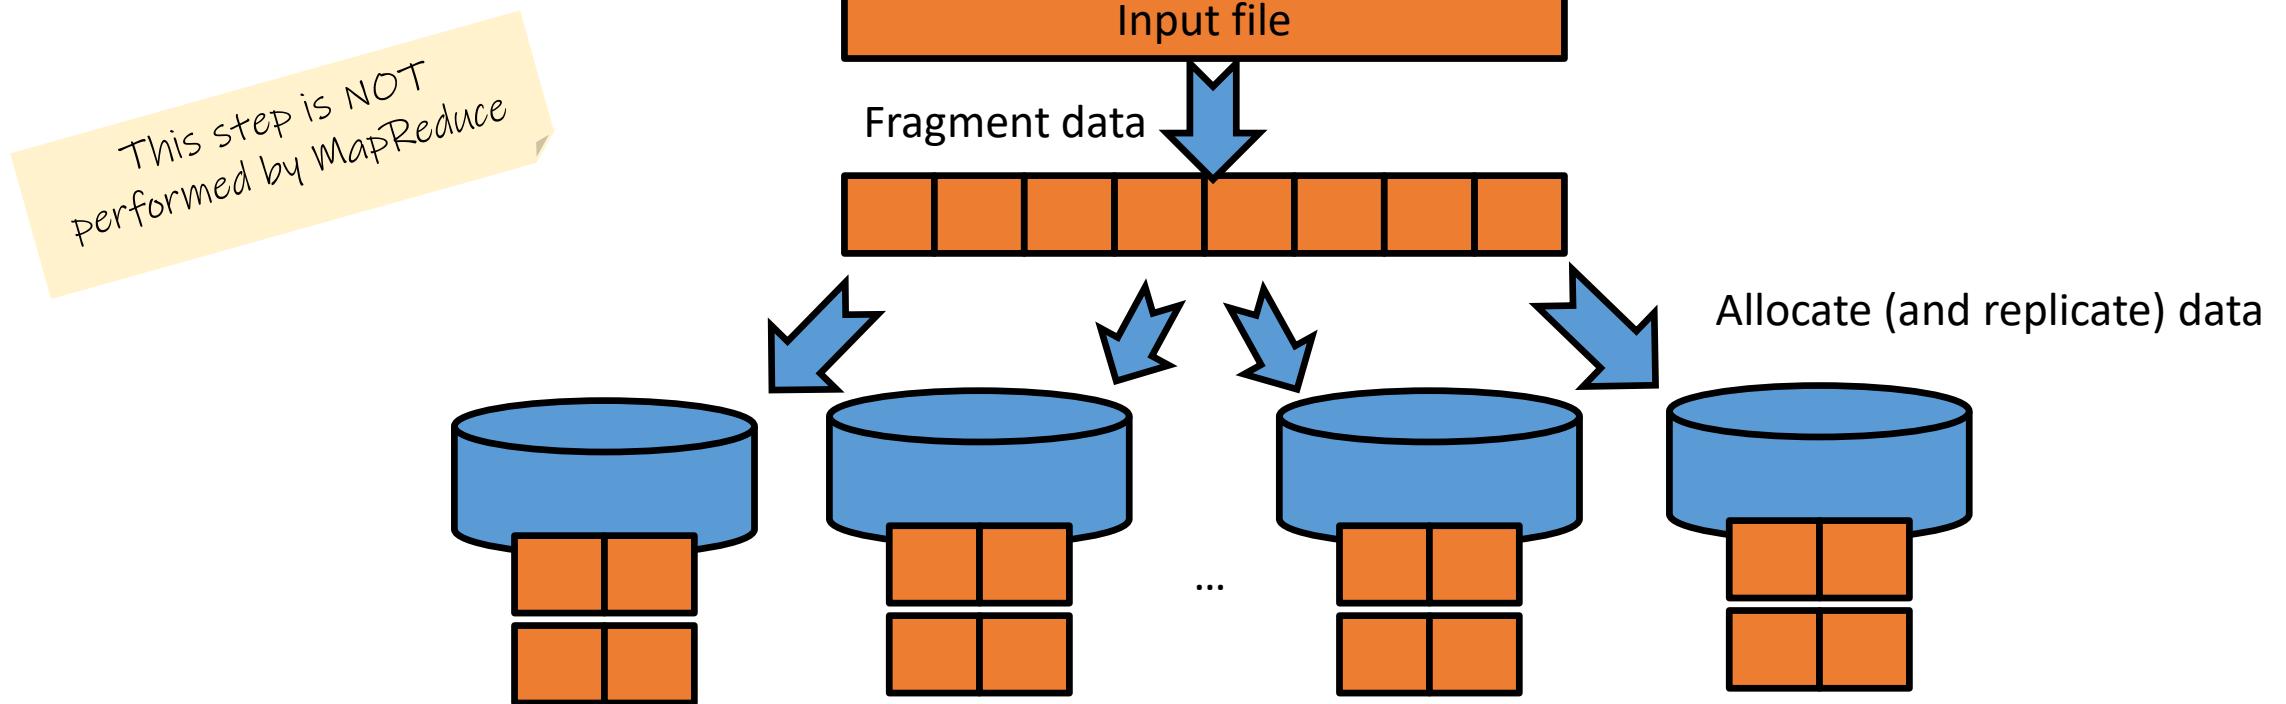
```
}
```

# Architecture

MapReduce implementation

# Pre-requisite: Data Loaded into a Distributed System

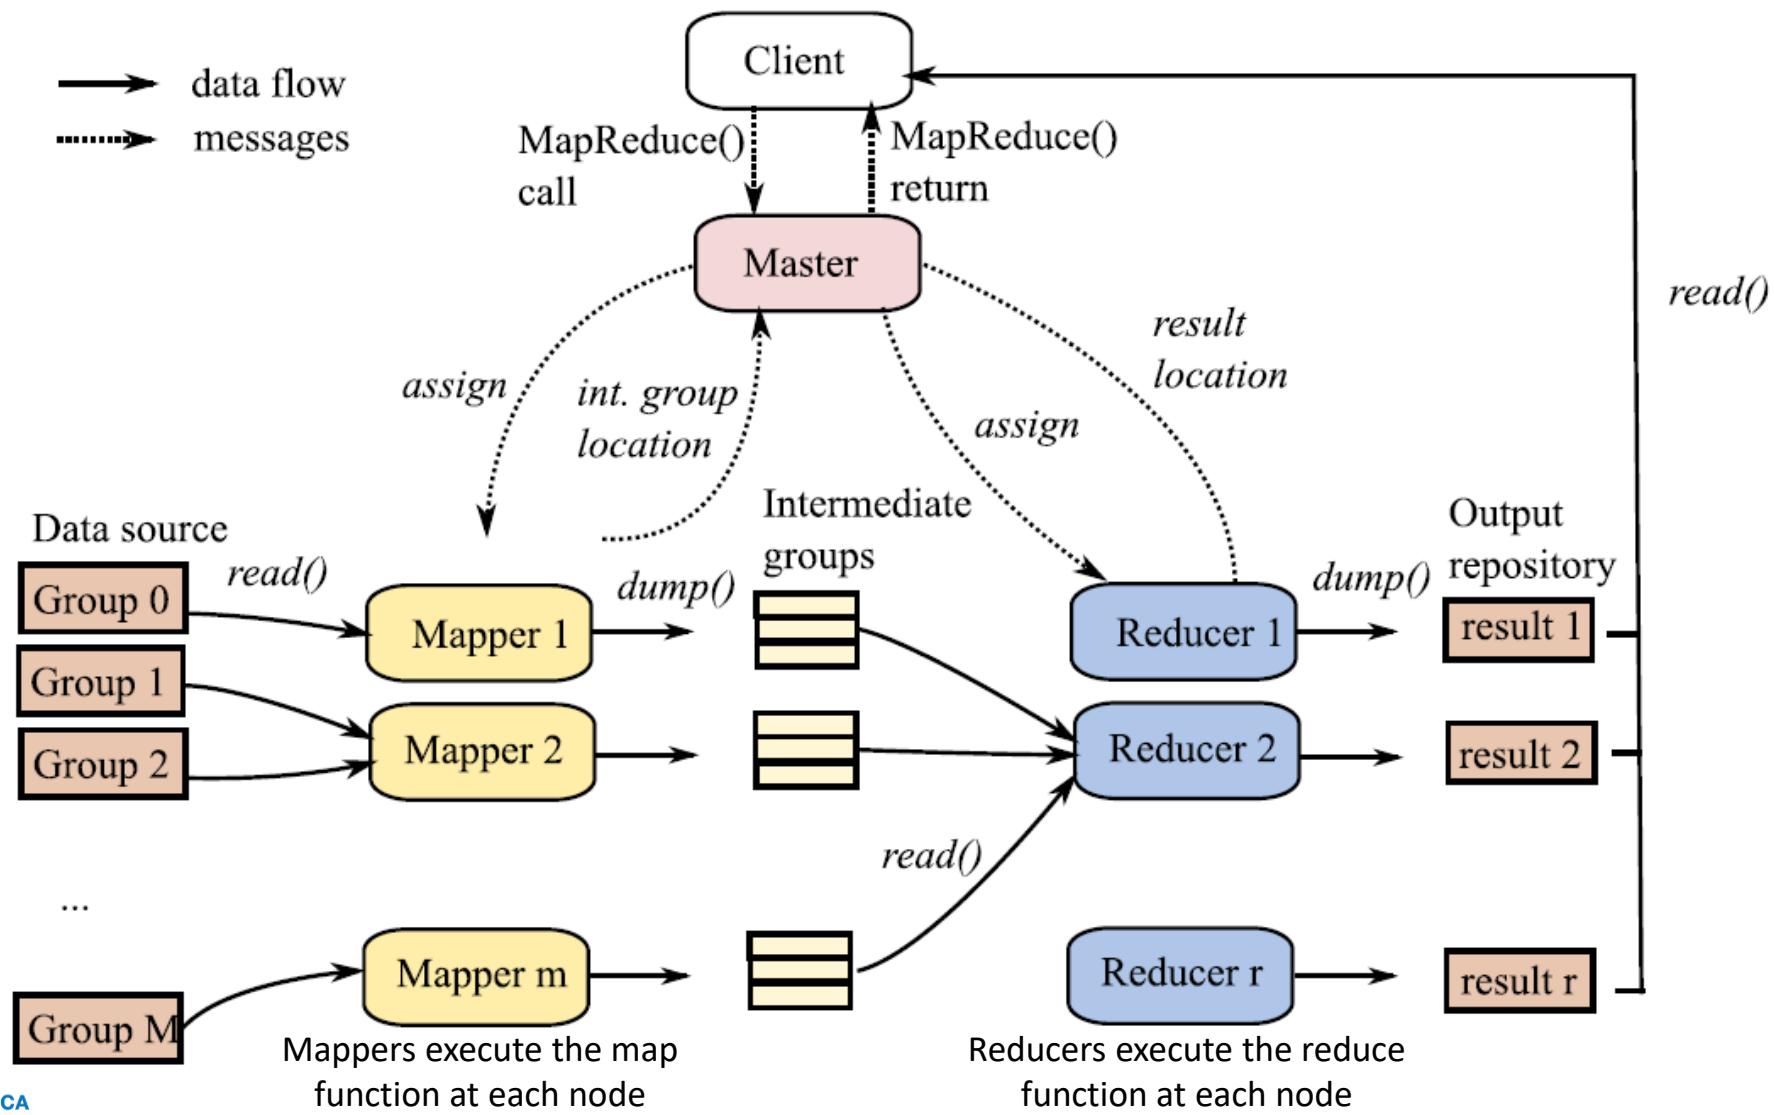
The **input data must be partitioned into fragments** (i.e., be already distributed) and allocated (**and** therefore **replicated**) in a distributed system (e.g., HDFS, HBase, Cassandra, MongoDB...)



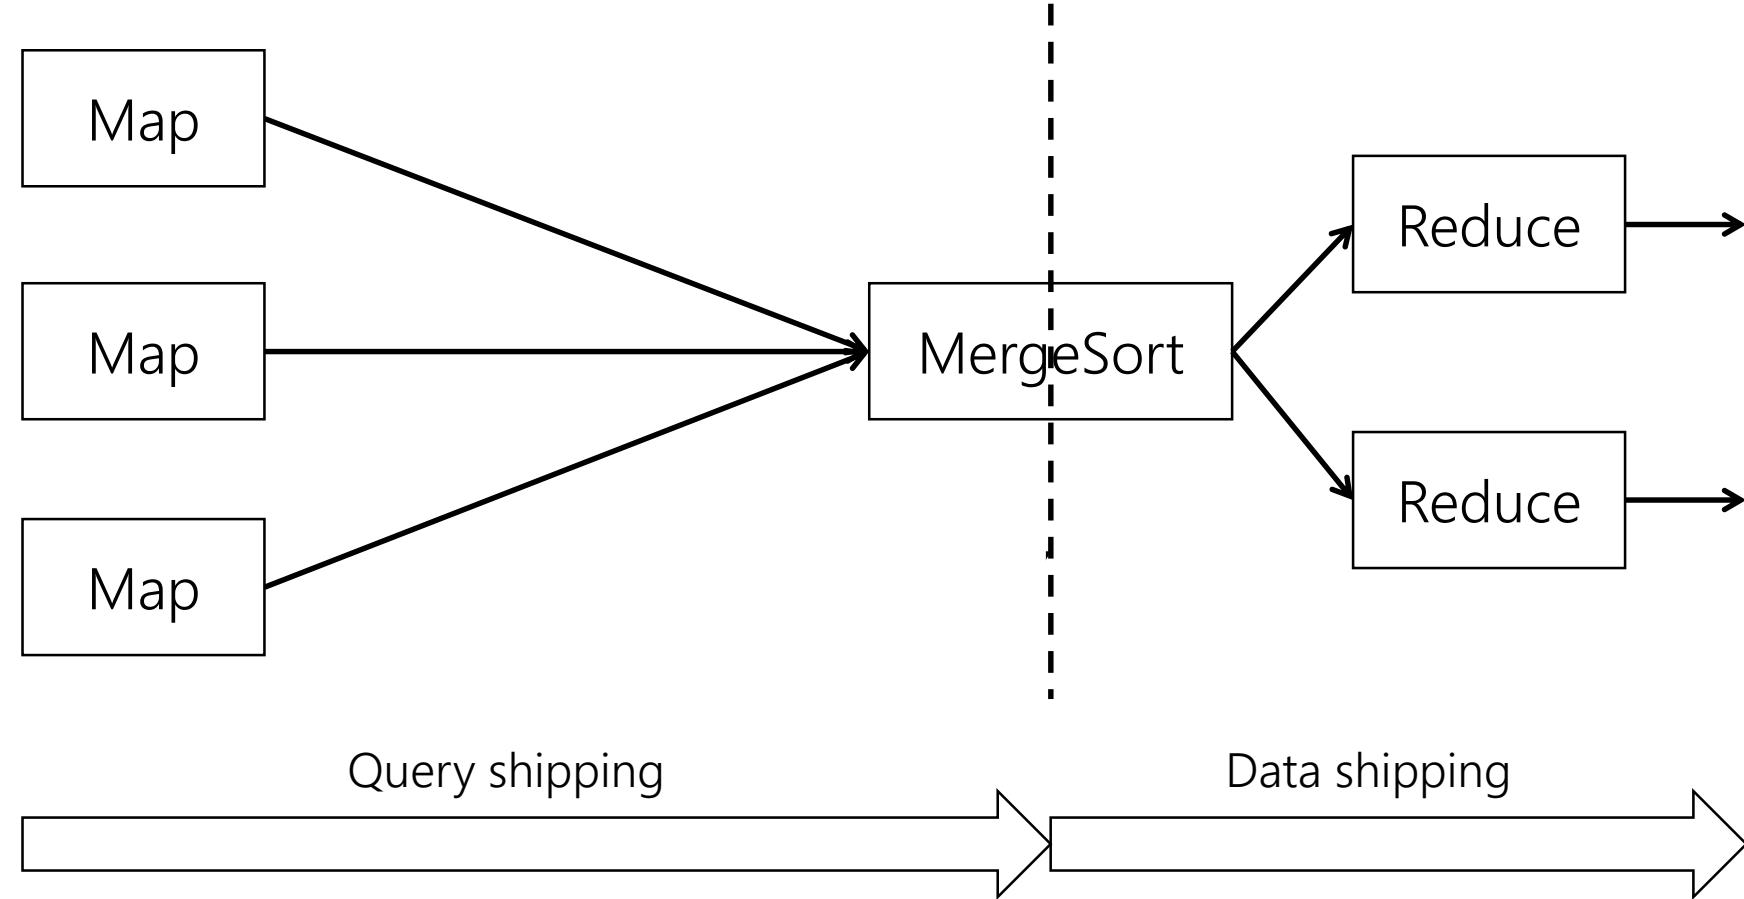
# MapReduce Main Components

- **Master:** receives the MapReduce (MR) job and schedules it
  - Send the map function to Mappers
  - Send the reduce function to Reducers
  - Monitors the overall execution and detects if a mapper or a reducer fails
  - Reassigns the tasks of a mapper or a reducer to another node upon failure
- **Mapper:** one per node assigned to execute the map function
  - The master assigns a set of blocks to each mapper
  - Executes the map function for each record in the set of blocks assigned
  - Stores the result of all the map executions in intermediate results
- **Reducer:** one per node assigned to execute the reduce function
  - The master assigns a set of intermediate results to each reducer
  - Executes the reduce function to each record in the intermediate results assigned
  - Stores the result of the reduce executions in the output storage specified in the MR job

# MapReduce Components: Tasks and Data Flows

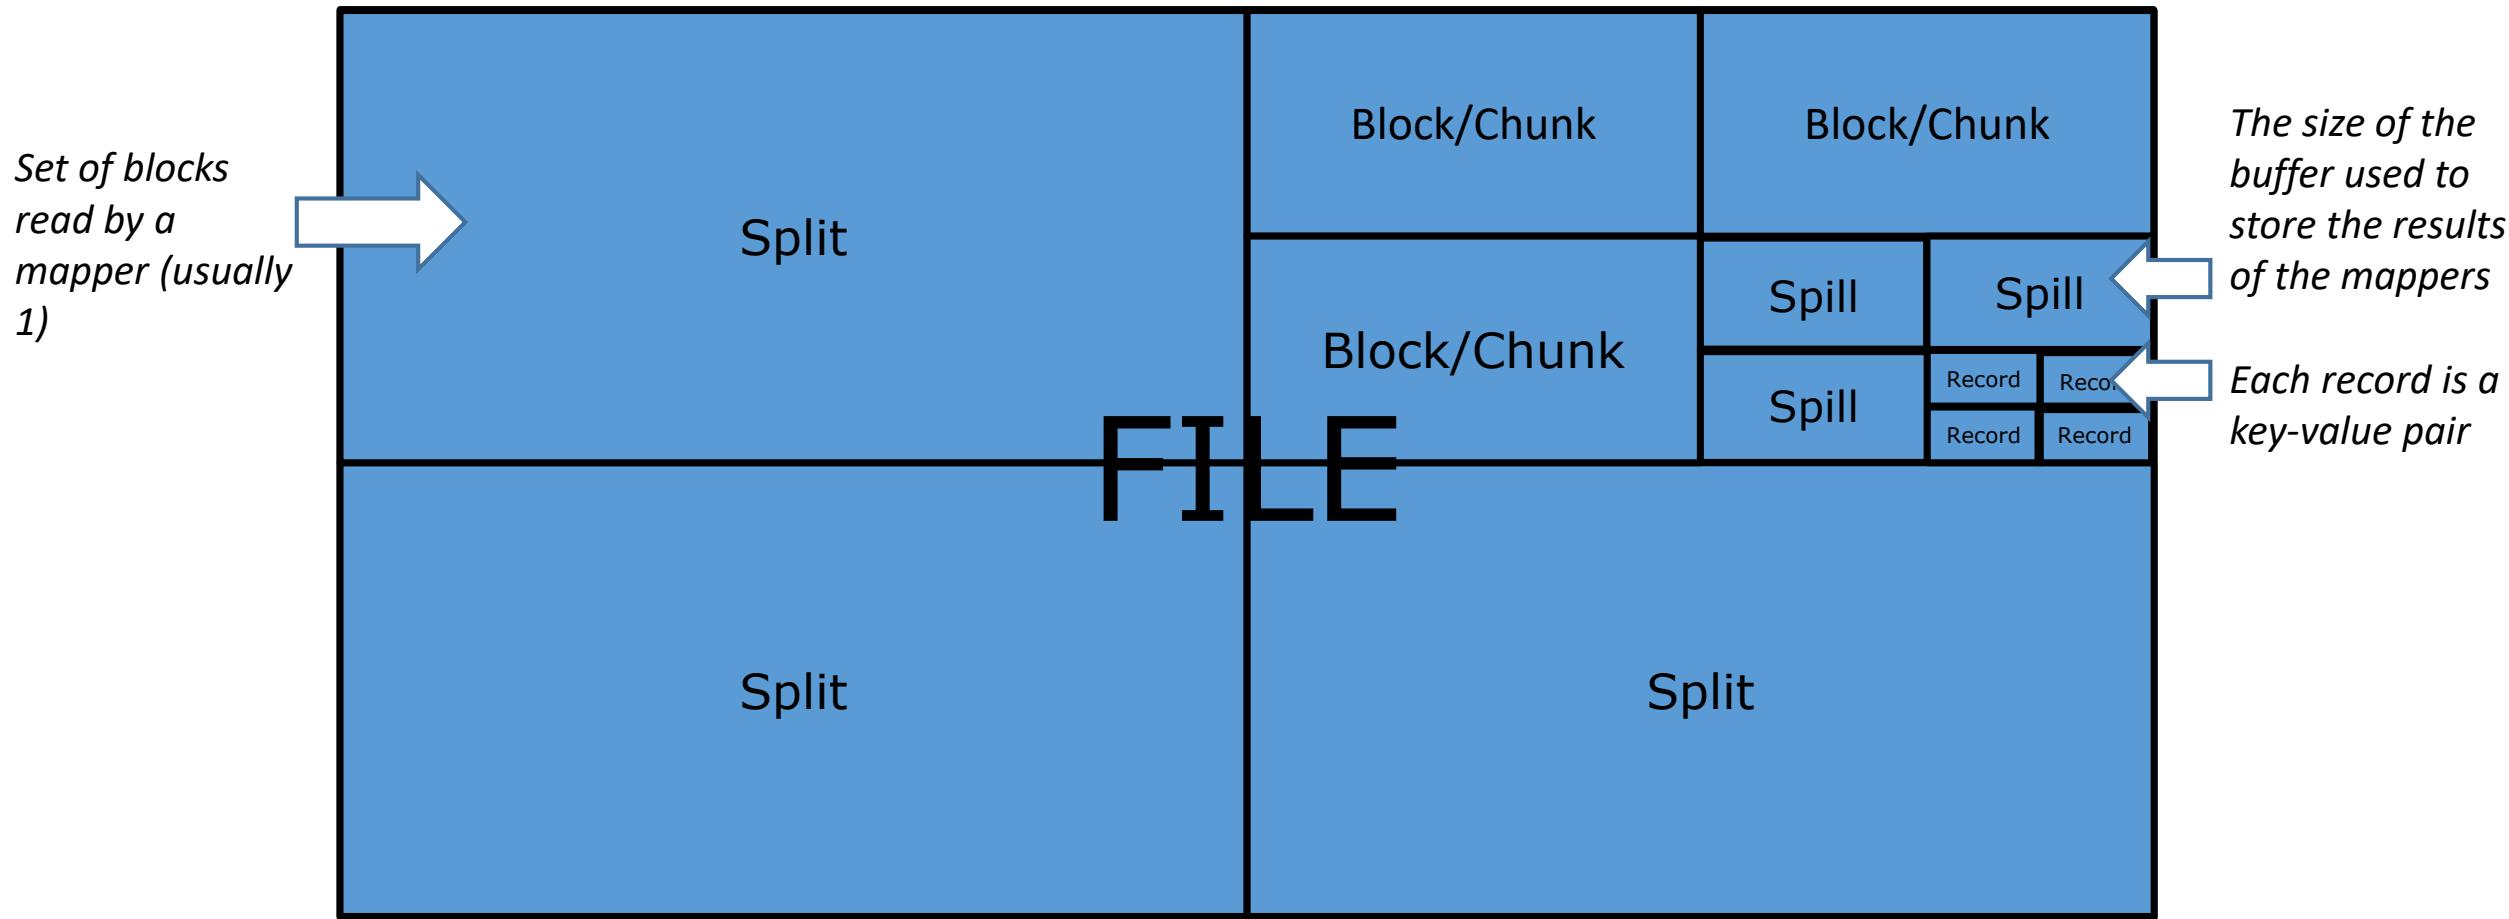


# Query shipping vs. data shipping



# MapReduce Objects

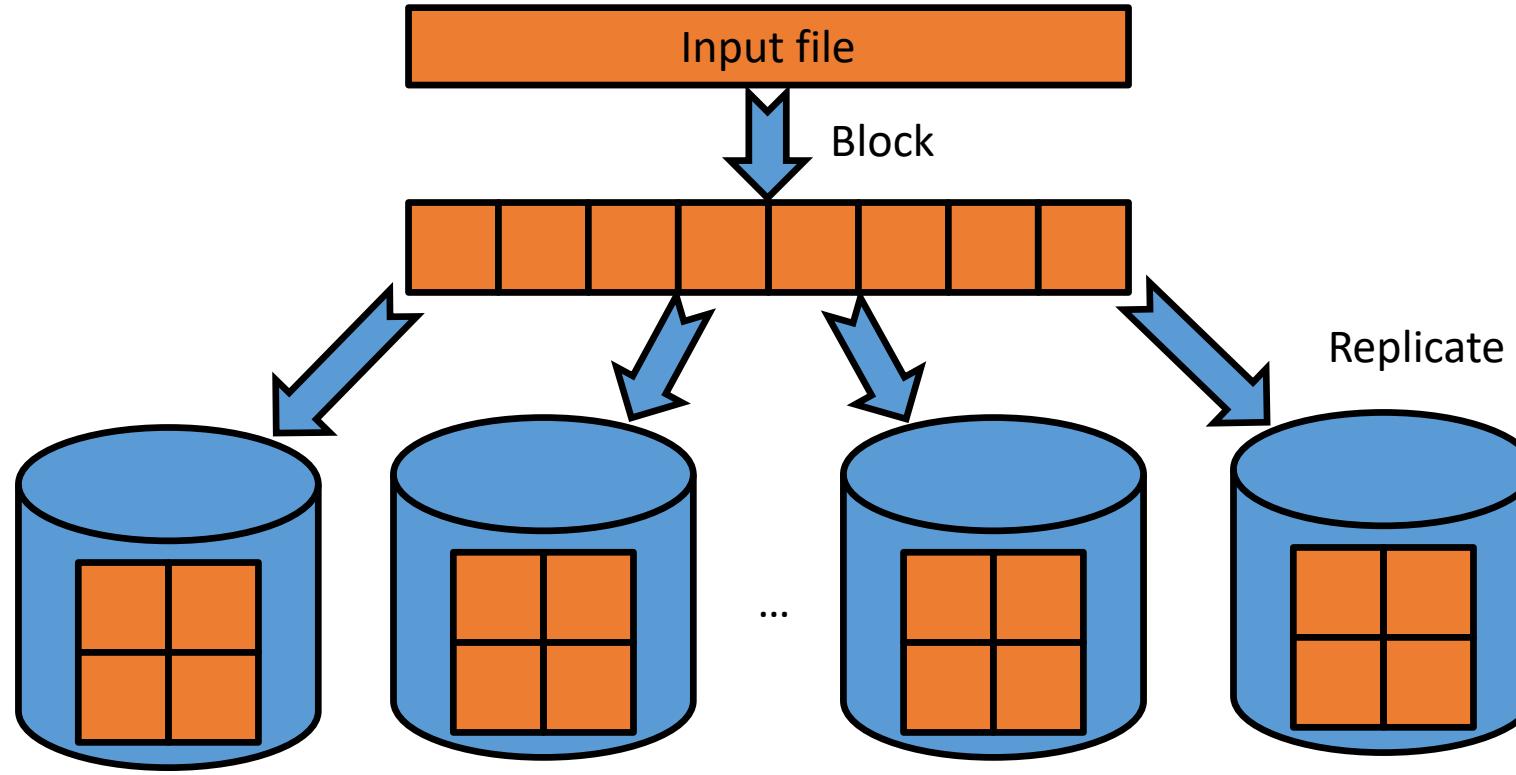
*Remember the block (sometimes also called chunk) is the atomic unit read from the distributed storage (see previous slide)*



# Internal algorithm

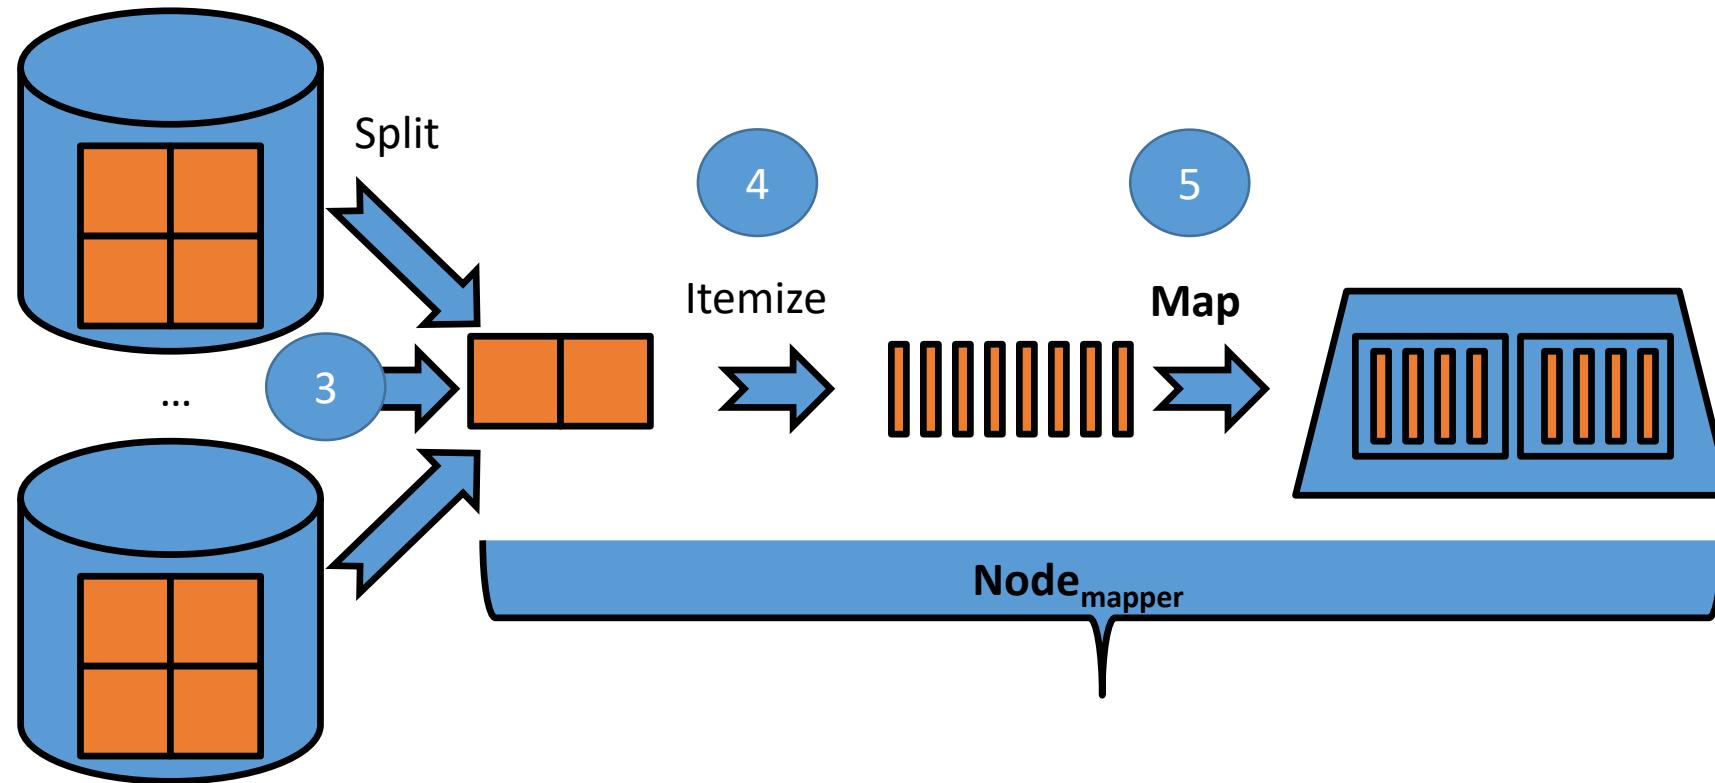
# Algorithm: Data Load

1. Upload the data to the distributed storage
  - o Partition them into blocks
  - o Using HDFS or any other storage (e.g., HBase, MongoDB, Cassandra, CouchDB, etc.)
2. Replicate them in different nodes



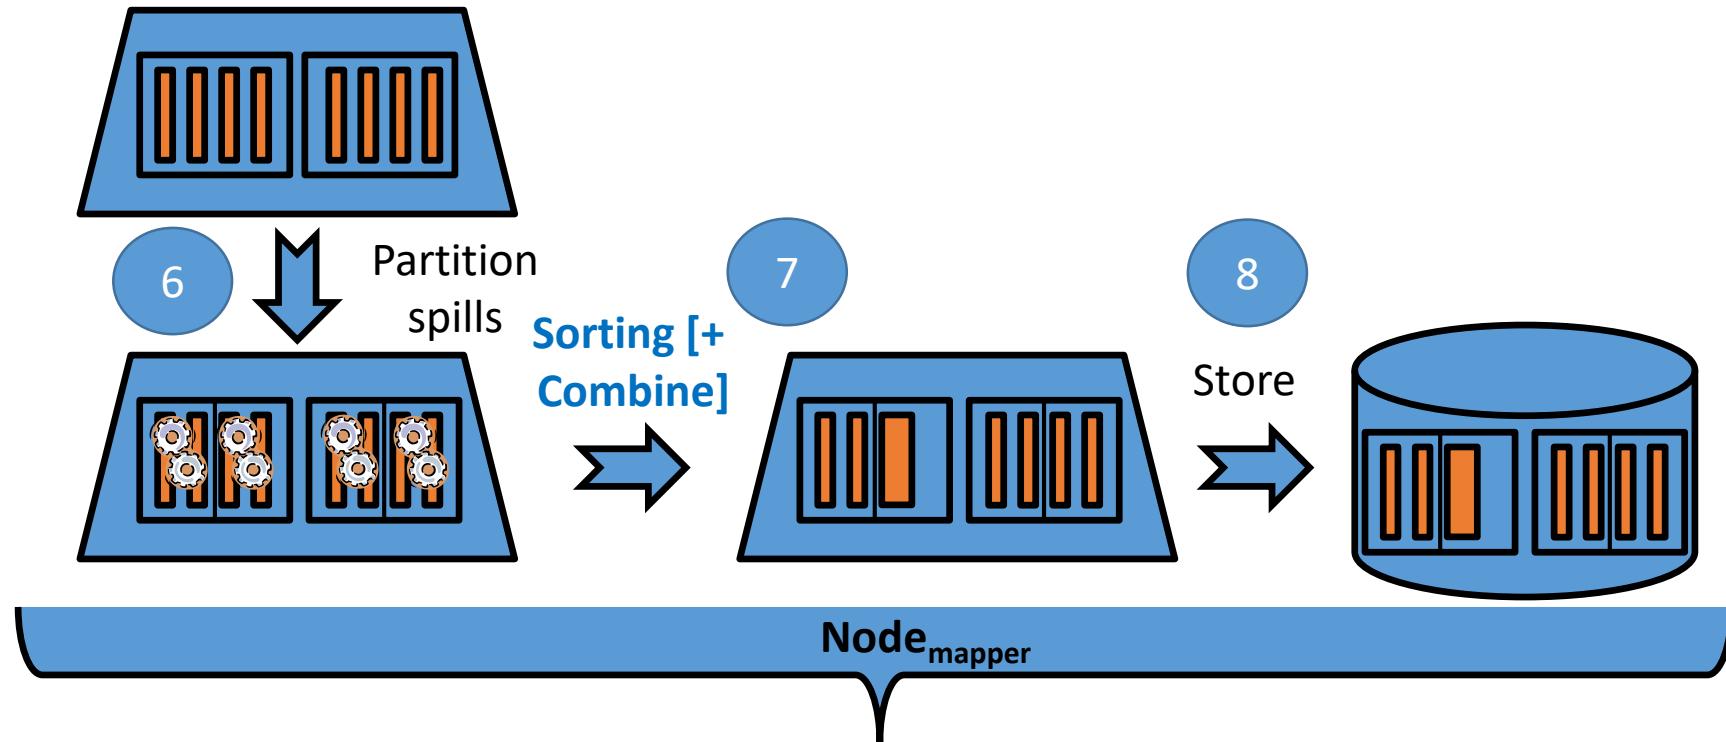
# MapReduce Algorithm: Map Phase (I)

3. Each mapper reads a split (logical chunk assigned to a map task, usually 1 block)
4. Divides the split into records, i.e. converts bytes onto key-value pairs (*itemization*)
5. Executes the map function for each record and leaves them in memory divided into spills



# MapReduce Algorithm: Map Phase (II)

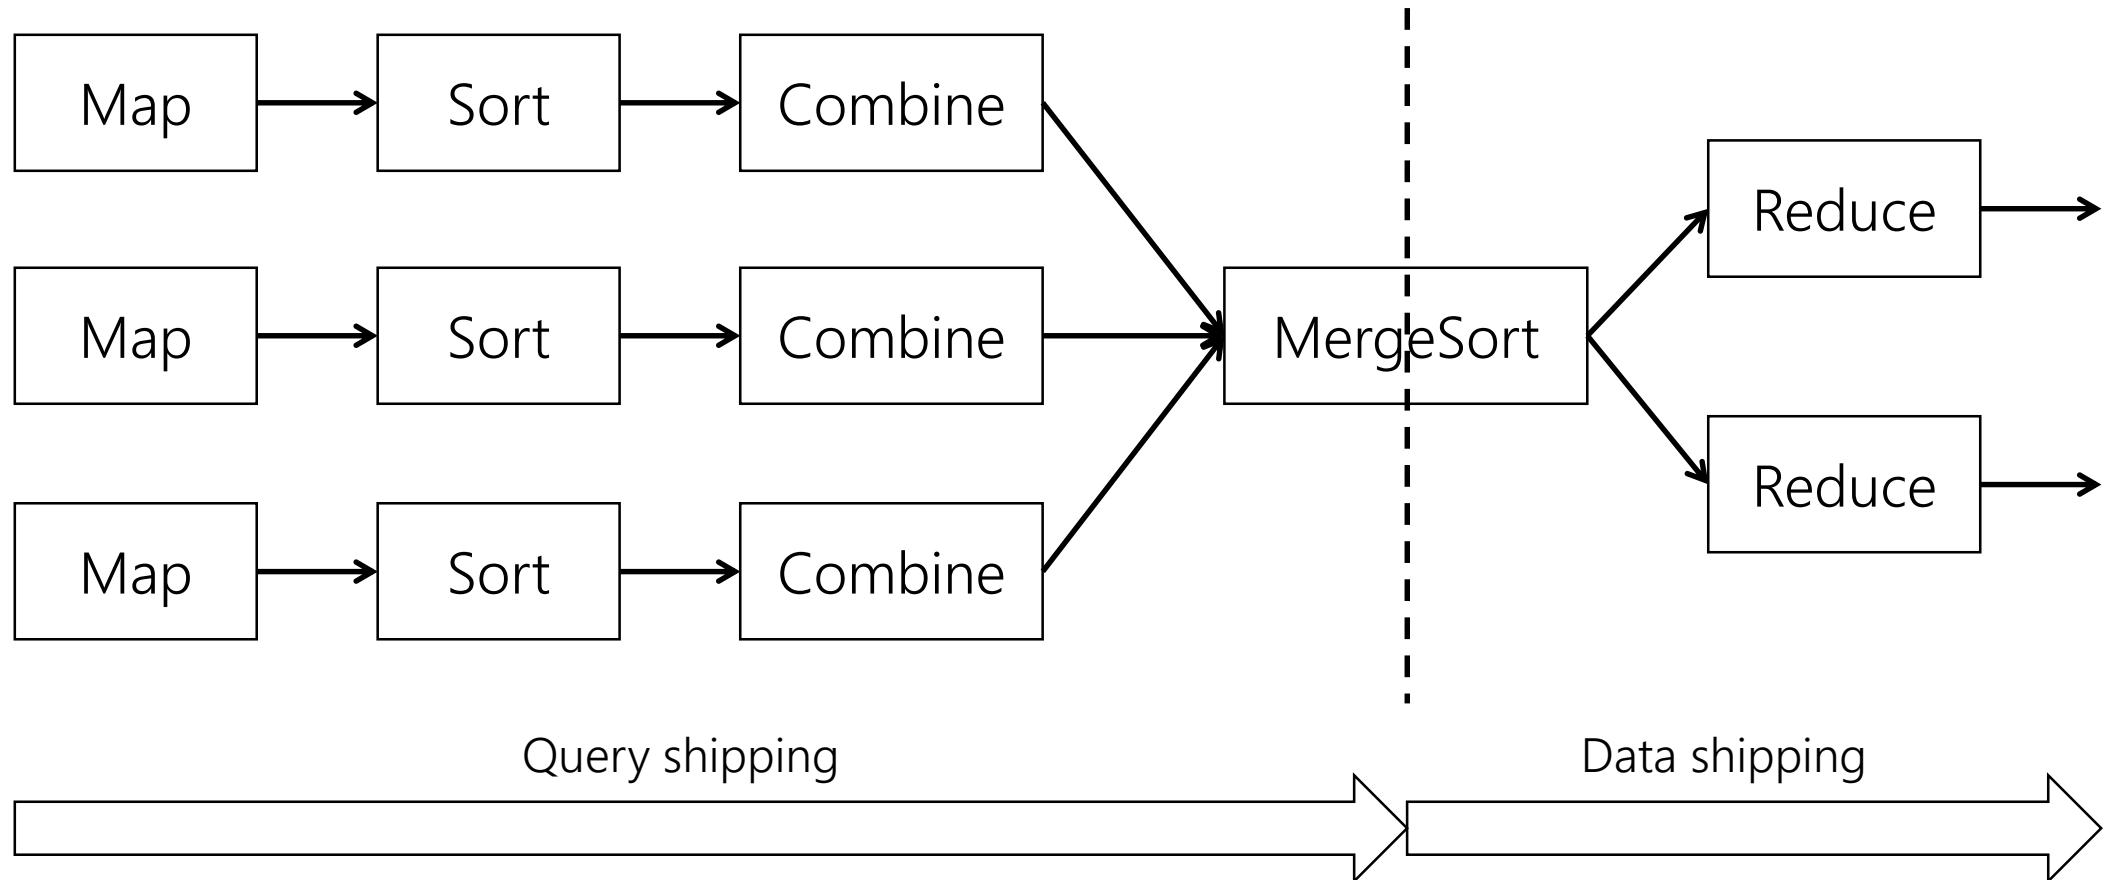
6. Each spill is then partitioned per reducers
  - o Using a hash function  $f$  over the new key, according to the number of reducers  $R$ ; i.e.,  $f(x) = x \% R$
7. Each spill partition is sorted independently
  - o If a **combine function** is defined, it is executed locally after sorting
8. Store the spill partitions into disk (massive writing)



# The Combine Function (I)

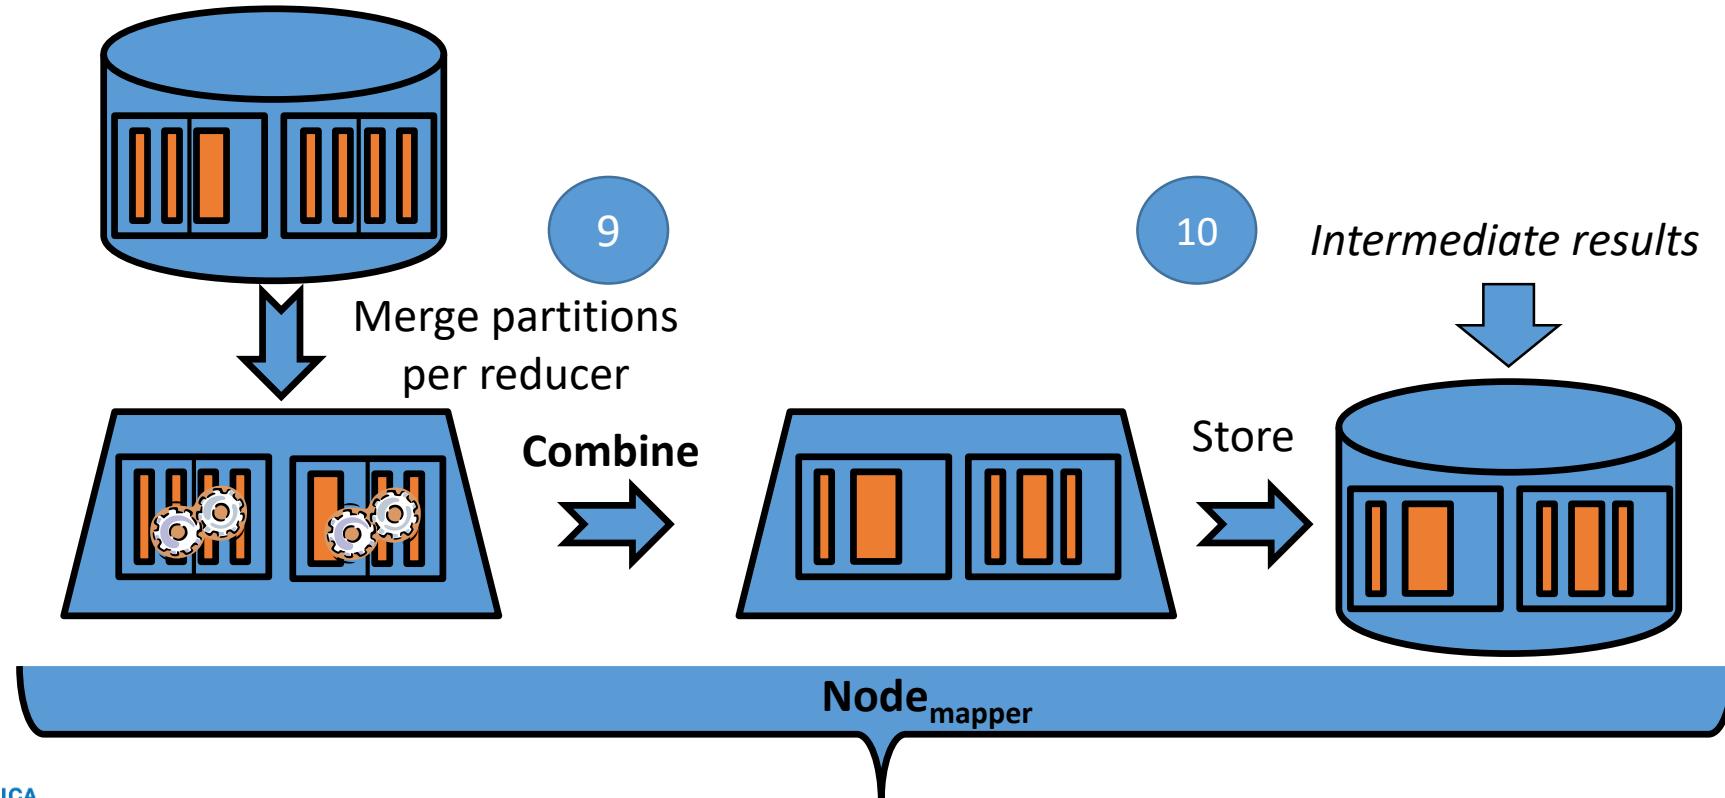
- Combine is executed locally in the Mapper nodes
  - Reduces the number of tuples sent to reducers
- Only possible when the combine function is
  - Commutative
  - Associative
- Typically, the combine and reduce functions are the same

# The Combine function (II)



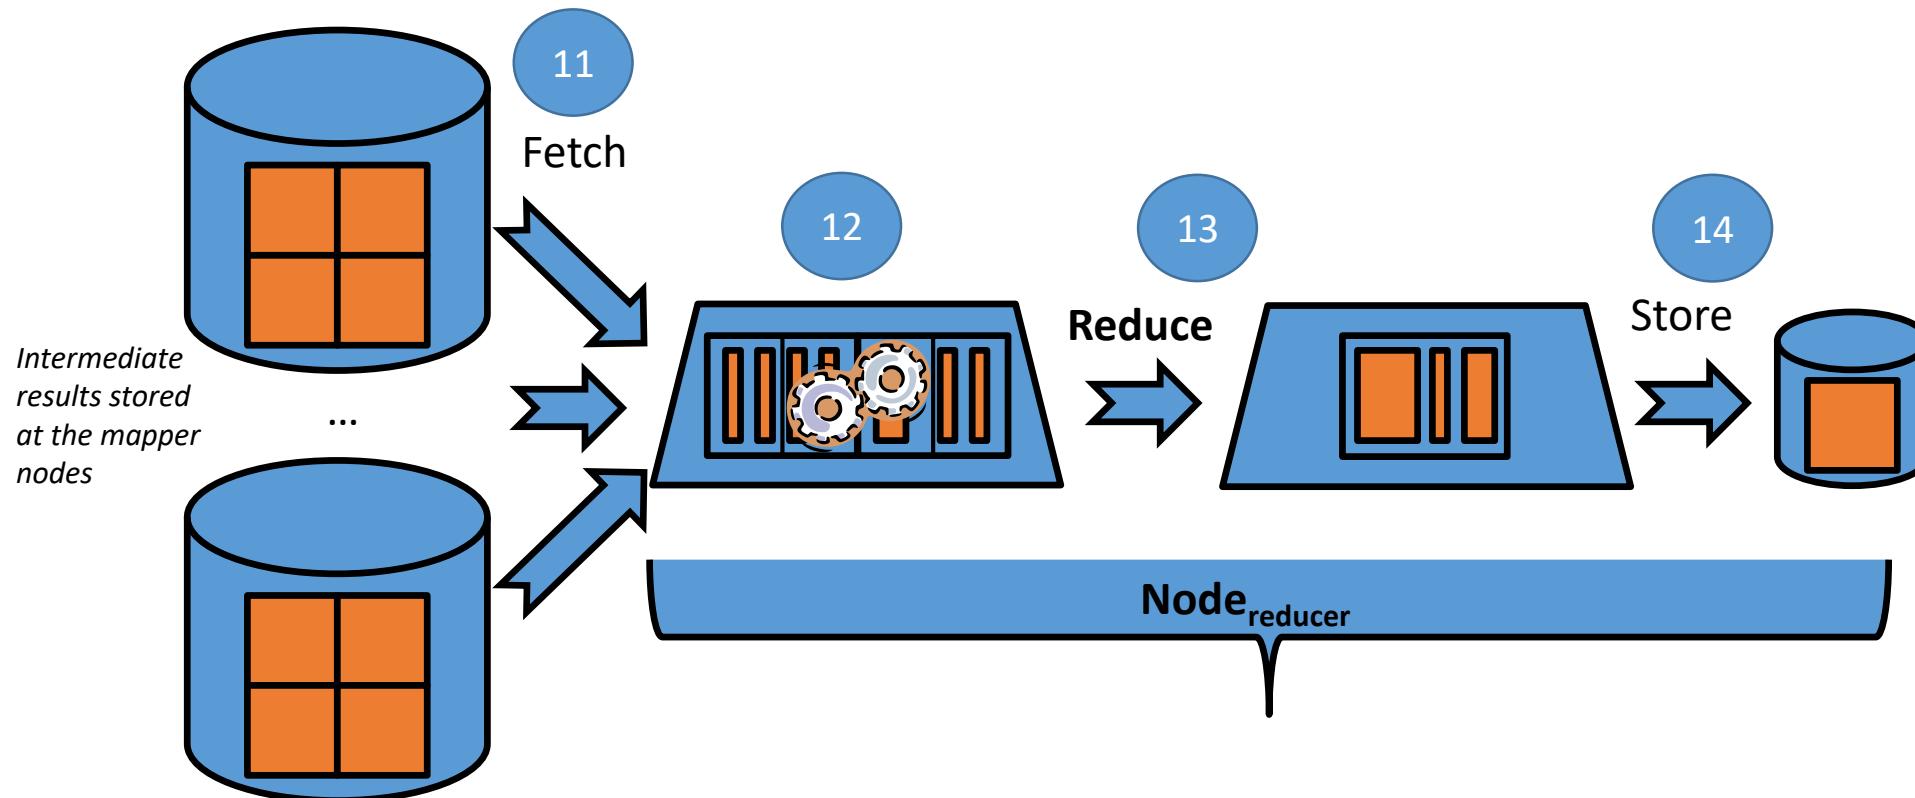
# MapReduce Algorithm: Map Phase (III)

9. Spill partitions are merged per reducer
  - Each merge is sorted independently
  - Combine is applied again
10. Store the result into disk into a single map output file



# MapReduce Algorithm: Shuffle and Reduce

11. Reducers fetch the intermediate results from mappers (massive data transfer)
12. Intermediate results fetched are merged and sorted
  - When merging files, for those  $\langle k, v \rangle$  with the same  $k$  a list of values is generated:  $k, \{v_1, \dots, v_n\}$
13. Reduce function is executed per key
14. Store the result into disk

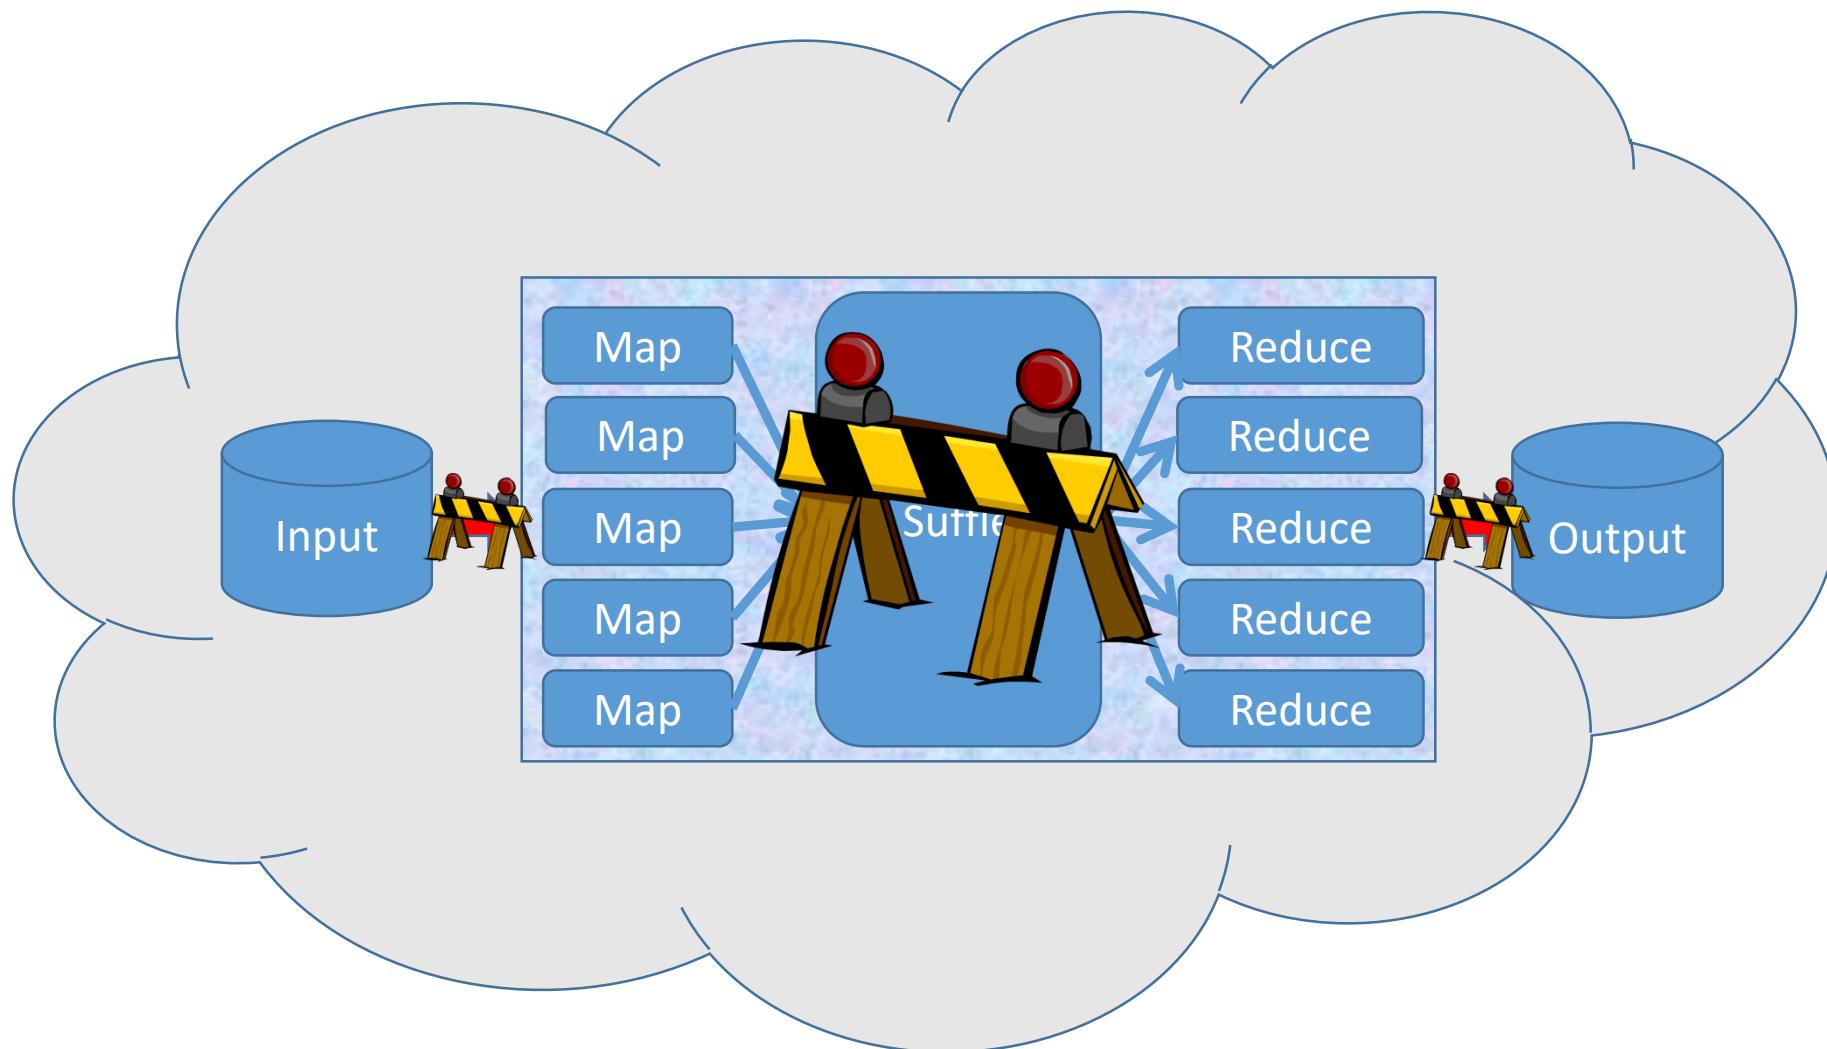


# Drawbacks

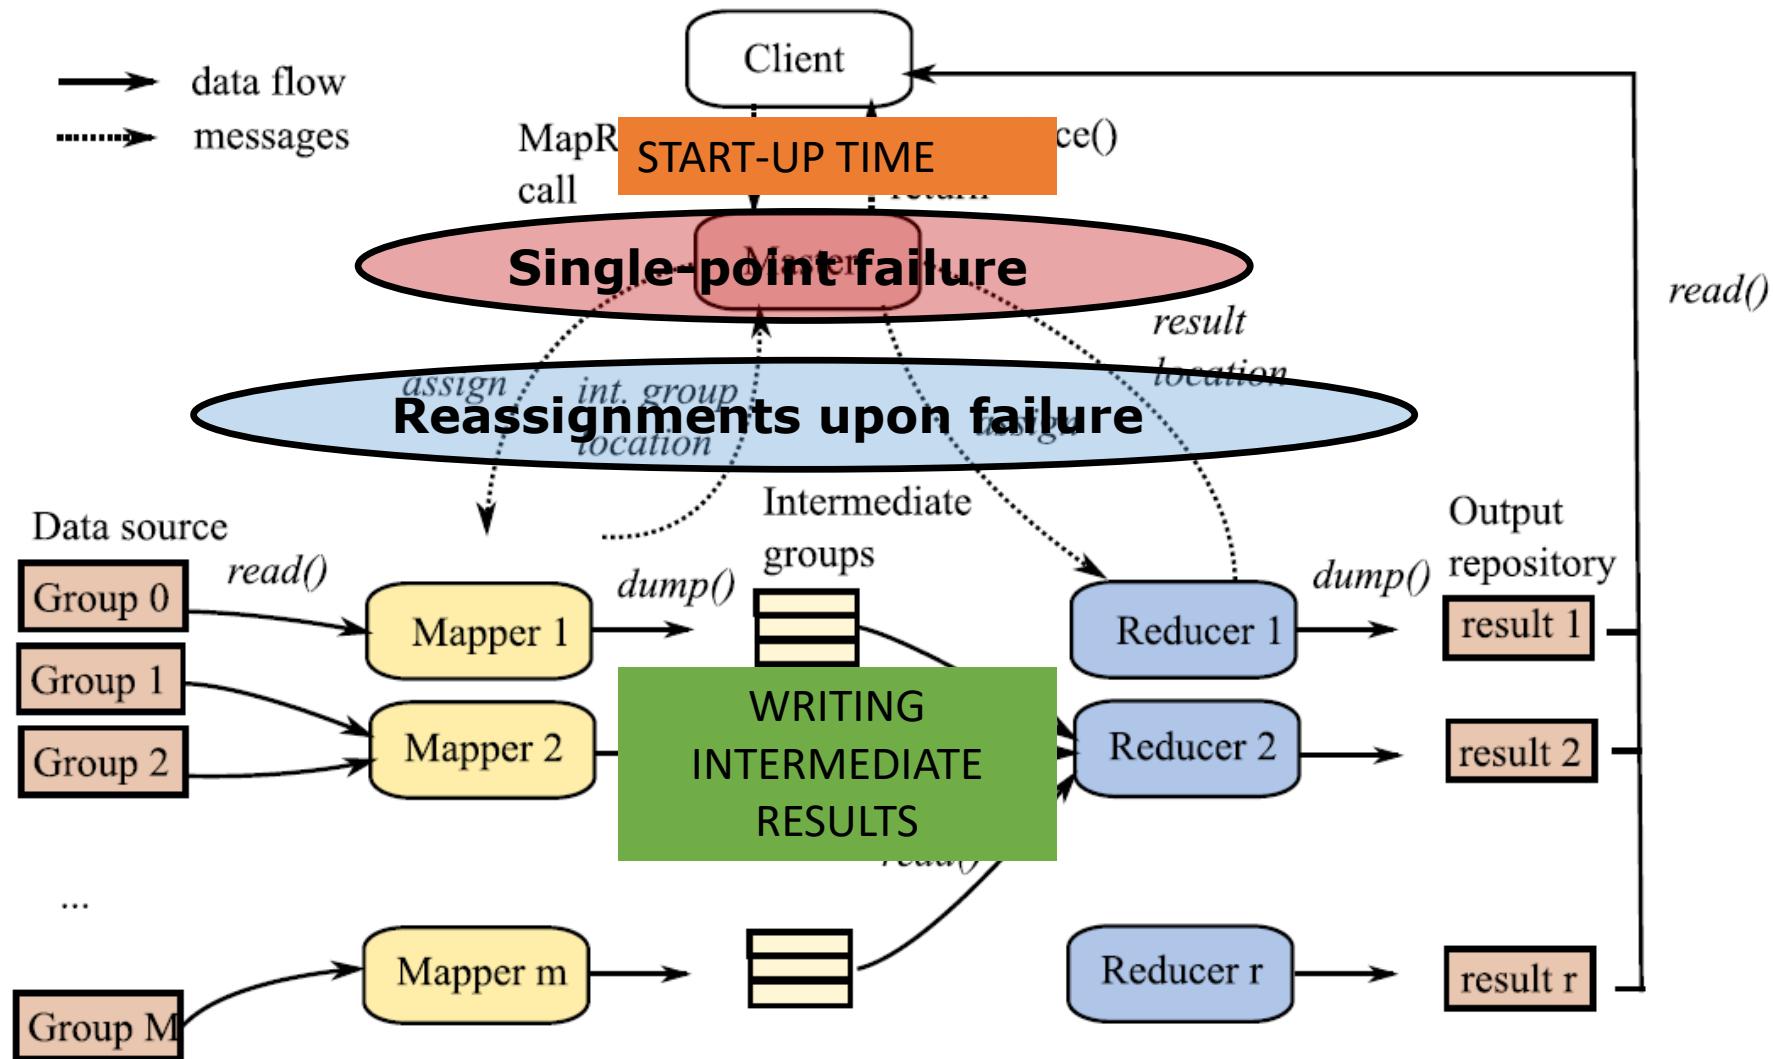
Bottlenecks of the algorithm



# Synchronization Barriers



# MapReduce: Tasks and Data Flows



# Implementing Operations with MapReduce

- The MapReduce paradigm program is computationally **complete** and ANY program can be adapted to it
- Furthermore, MapReduce's signature is closed
  - Thus, iterations can be achieved by nesting a MapReduce job
- However, some tasks better adapt to it than others
  - Easily adaptable:
    - Aggregations
    - Selections
    - Projections
    - Set operators
    - Sorting
  - Difficult for:
    - Joins
    - Any other operation needing to access data in another fragment to be executed (i.e., against the independence principle)

All these operators can be executed in each fragment independently  
(independence principle of their execution)



# Beyond MapReduce: Spark

*"Unified abstraction for cluster computing, consisting in a **read-only**, partitioned collection of records. Can only be created through deterministic operations on either (1) data in stable storage or (2) other RDDs."*

It is an evolution of MapReduce

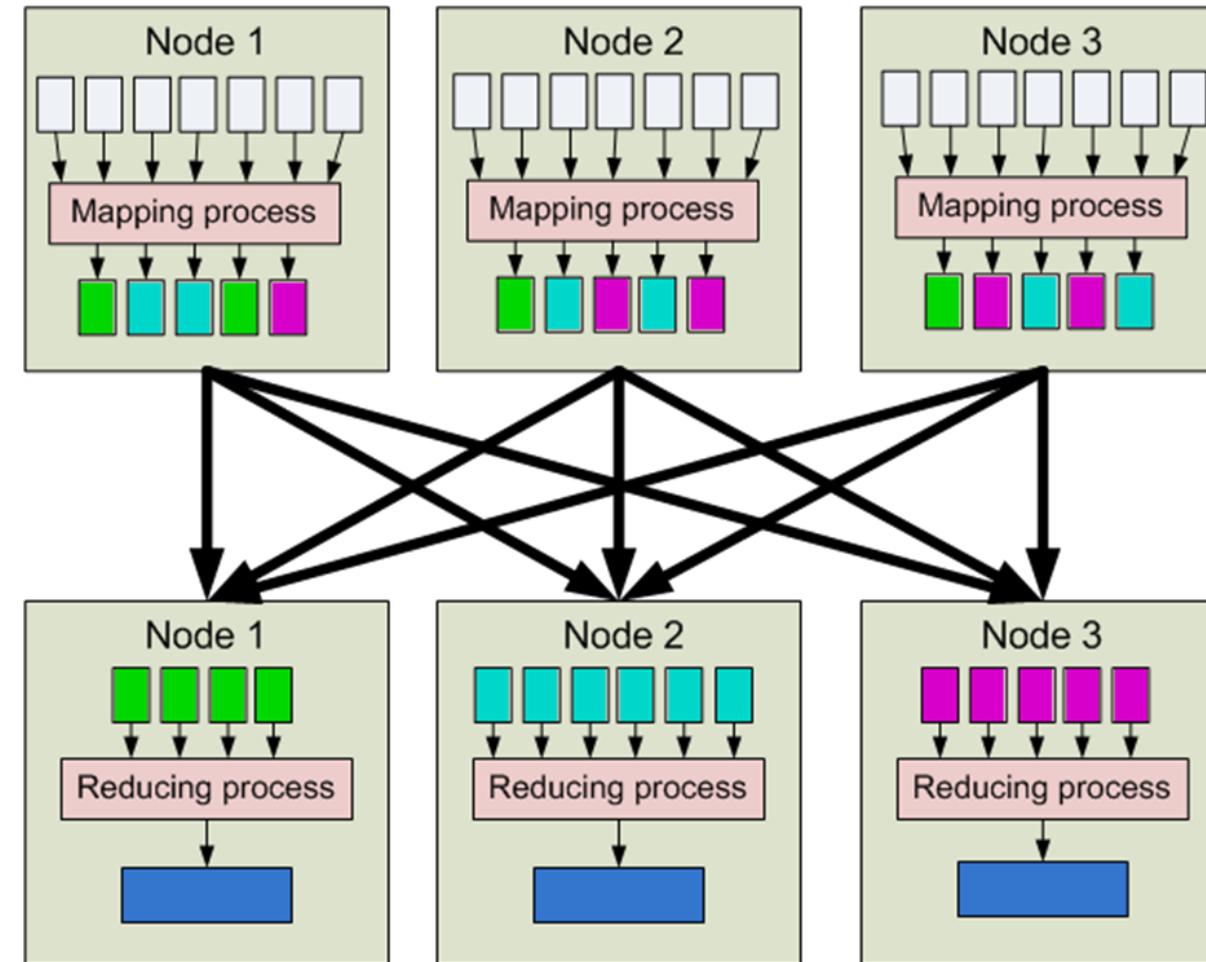
- Based on a richer data structure (dataframe)
- Instead of two operations: map and reduce, it provides more than 20 operators
  - Projection, Selection, Join, Group By... welcome back!
  - Reduces the black box problem in MapReduce
- Models the Spark job execution as a graph defined by the user programmatically
  - MapReduce always follows the same execution scheme: Map, Shuffle, Reduce
  - Spark is more flexible and the user can adapt and configure it
- Avoids writing intermediate results in disk (intensive in-memory processing)

# Putting everything together

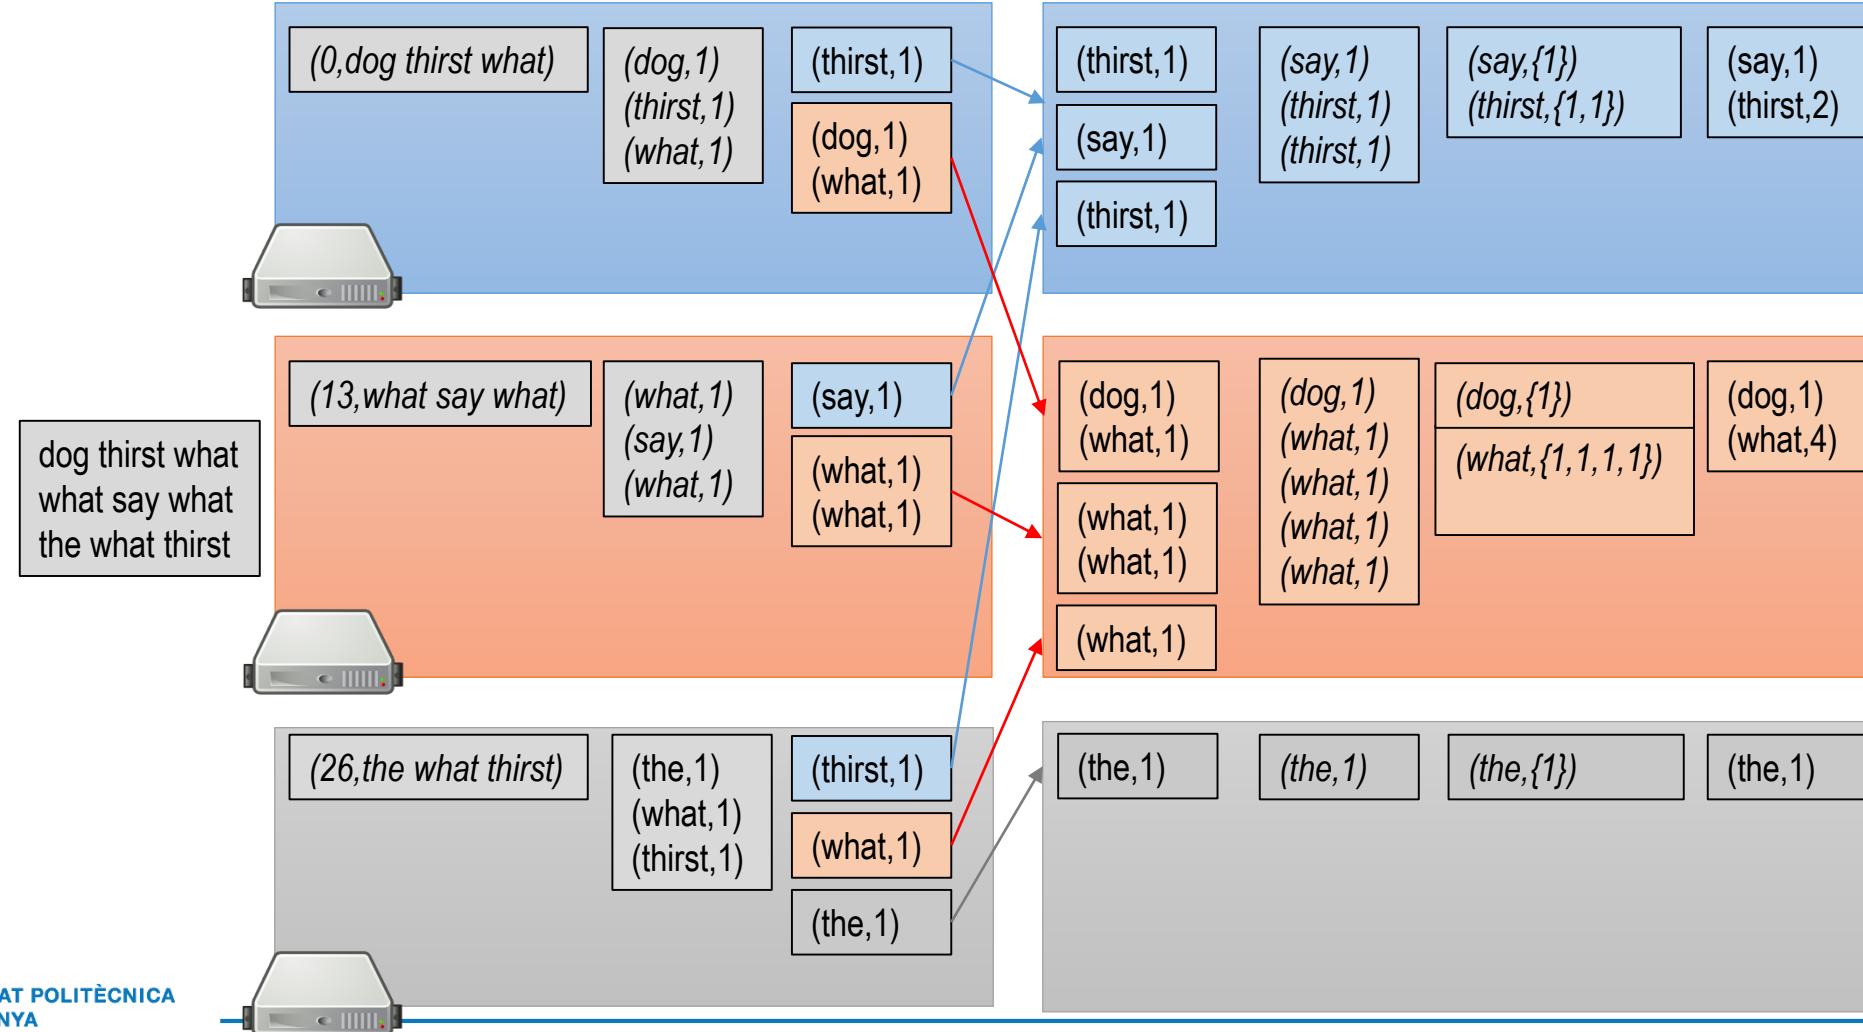
Executing a MapReduce job step by step

# MapReduce

1. Input
2. Map
3. Partition [Sort]
4. Shuffle
5. Merge Sort
6. Reduce
7. Output

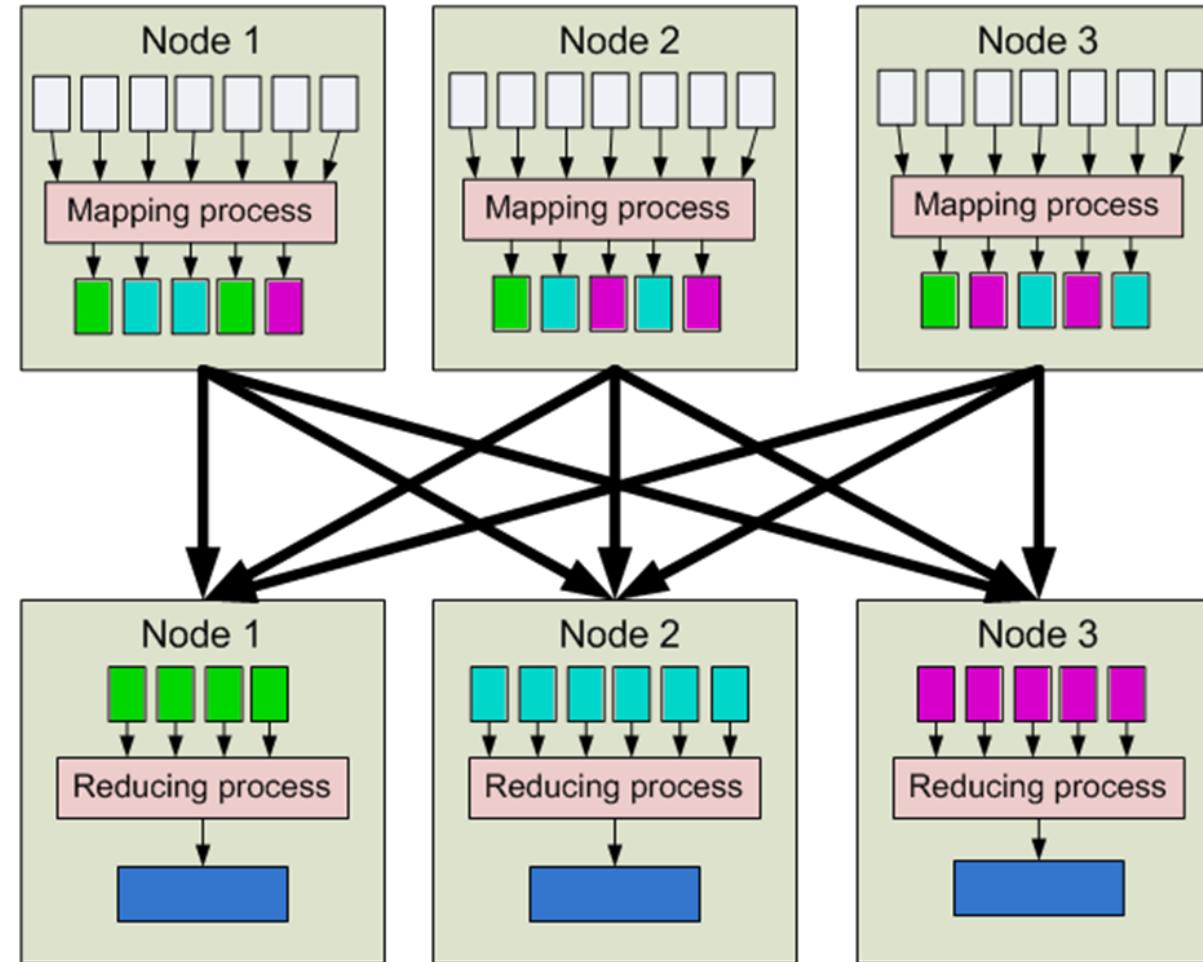


# MapReduce: Counting Words

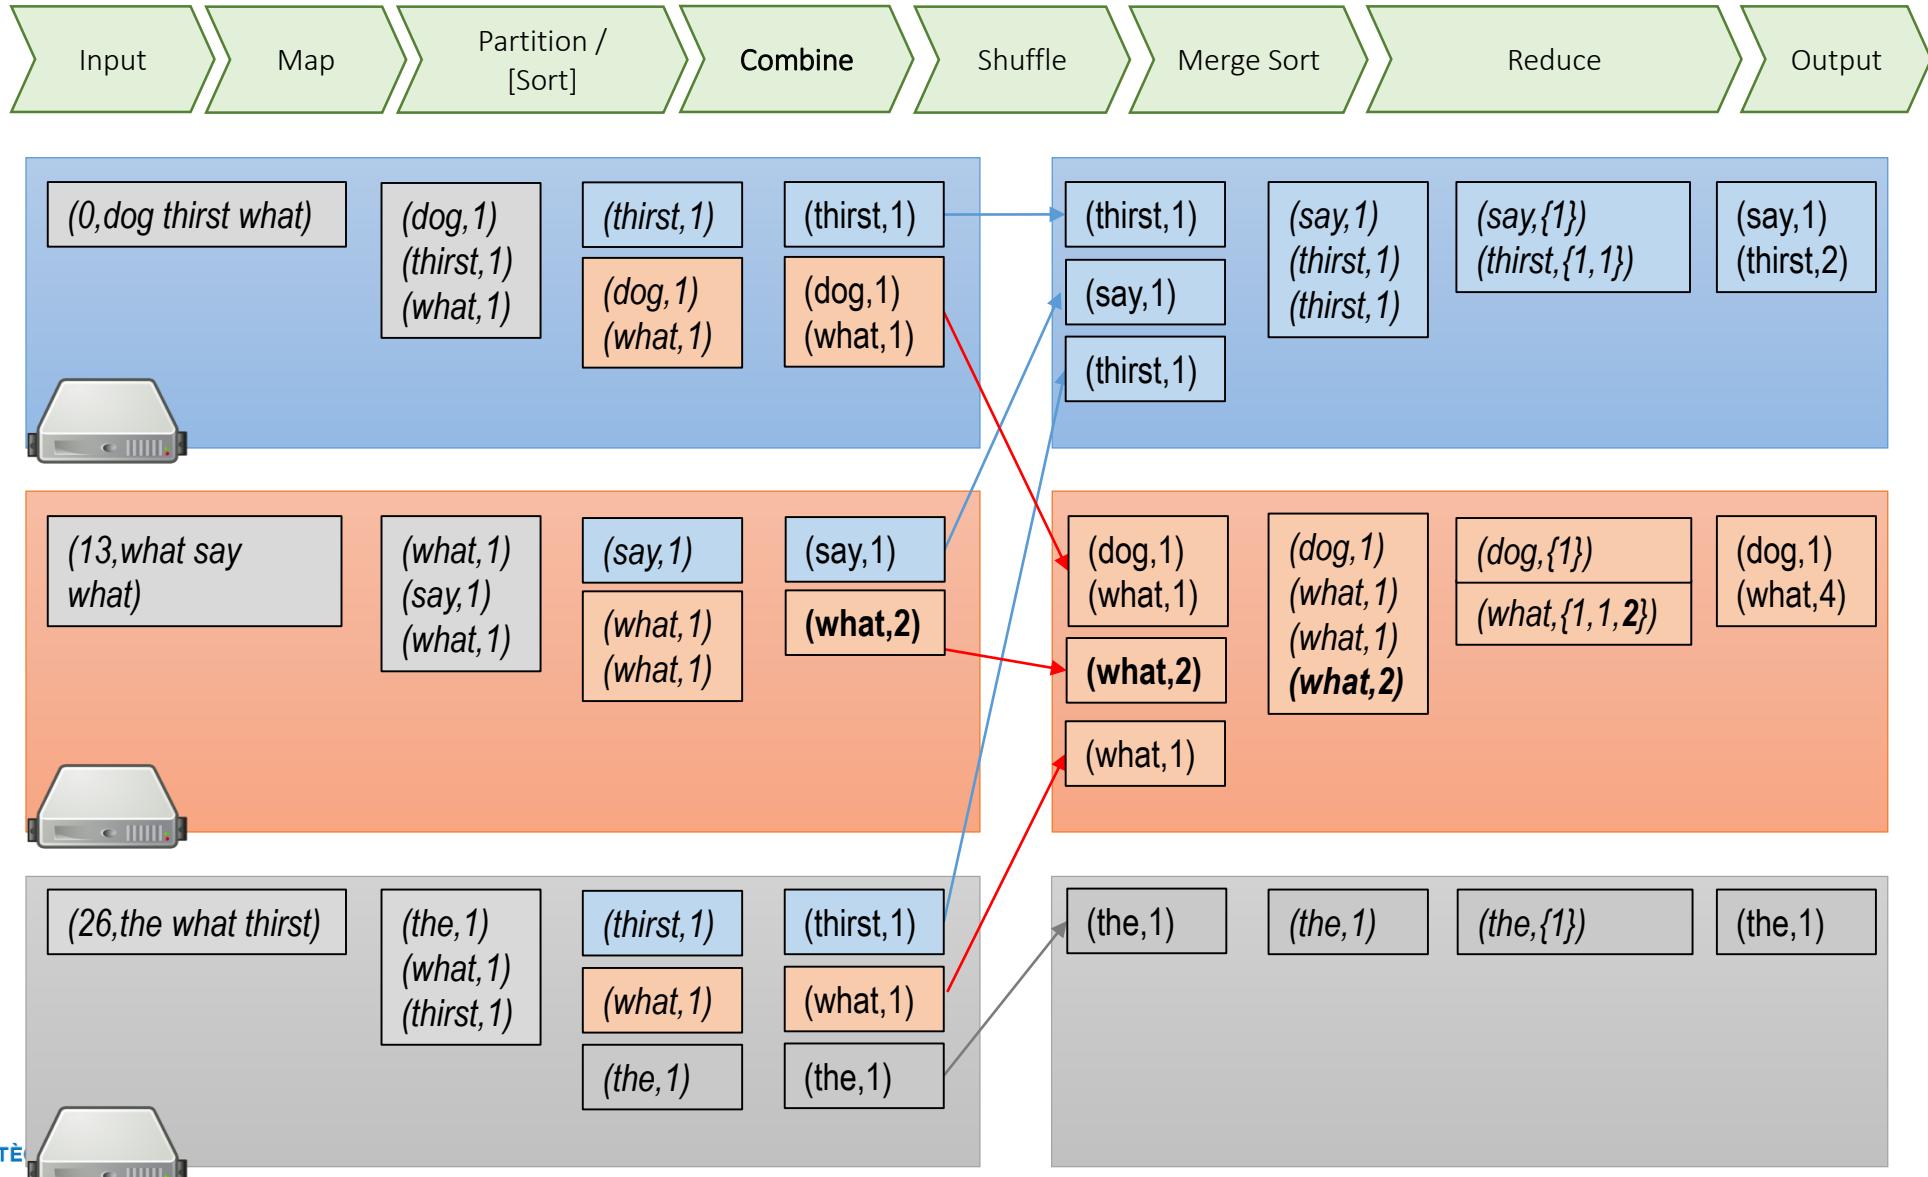


# MapReduce: Combiner

1. Input
2. Map
3. Partition [Sort]  
*("Combine")*
4. Shuffle
5. Merge Sort
6. Reduce
7. Output

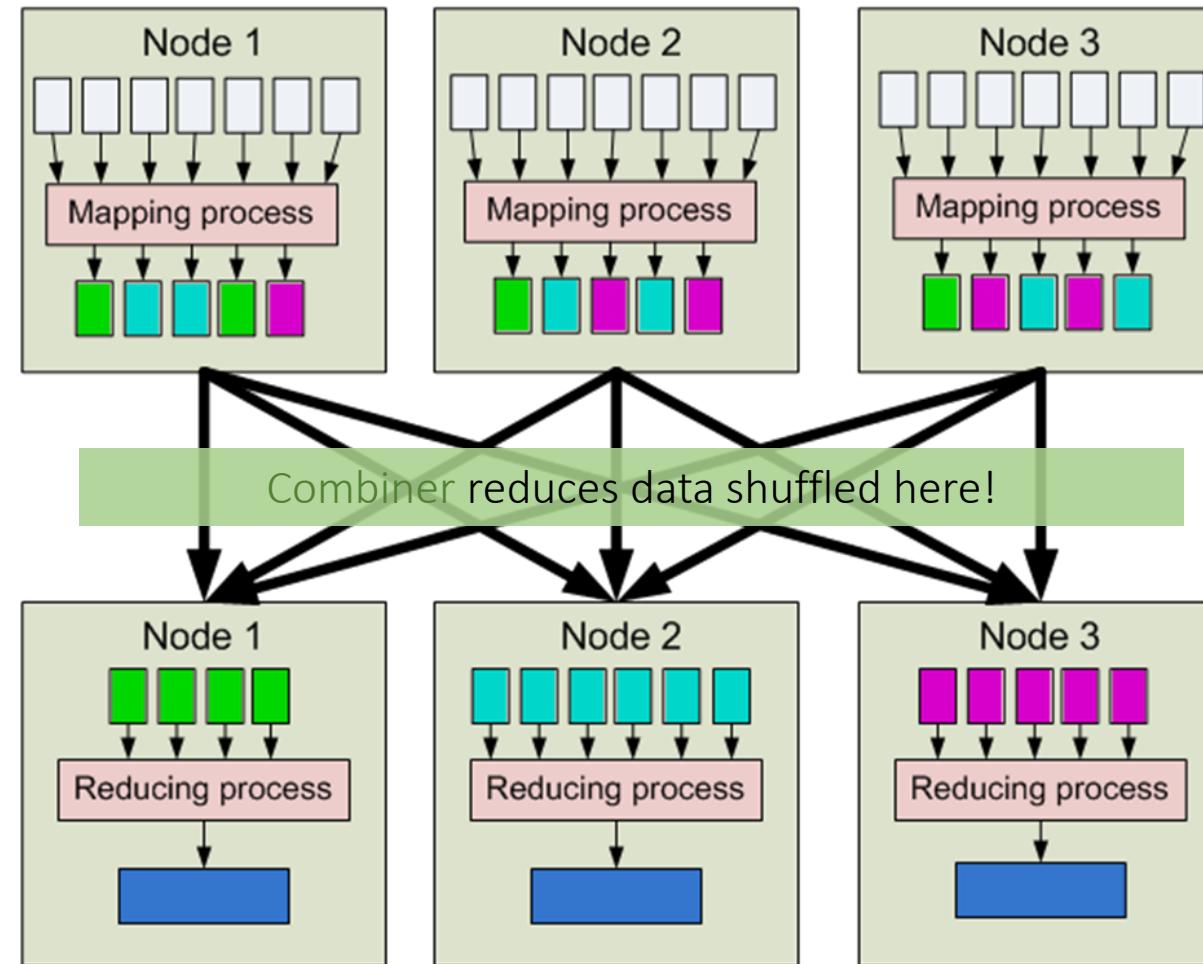


# MapReduce: Combiner



# MapReduce: Combiner

1. Input
2. Map
3. Partition [Sort]  
*(“Combine”)*
4. Shuffle
5. Merge Sort
6. Reduce
7. Output



# Summary

- The Algorithm
  - Map, Shuffle, Reduce
    - Combine
- Drawbacks
- Spark

# Bibliography

- Jeffrey Dean et al. *MapReduce: Simplified Data Processing on Large Clusters*. OSDI'04
- D. Jiang et al. *The performance of MapReduce: An In-depth Study*. VLDB'10
- S. Abiteboul et al. Web Data Management, 2012
- Apache Spark: <https://spark.apache.org/>