

Document Stores

Knowledge Objectives

1. Explain the main difference between key-value and document stores
2. Justify why indexing is a first-class citizen for document-stores and it is not for key-value stores

Application Objectives


1. Given an application layout and a small query workload, design a document-store providing optimal support according to a given set of criteria

Distributed Architectures

Consistent Hashing

Indexes

- Index – associates a key with its (physical) address
 - Trees – logarithmic search complexity (-> distributed trees seen in the key-value lecture)
 - Hash tables – constant search complexity
 - Good for point queries
 - Do not support range search
- And what if the data grows too much?




Direct access to an object
(vs sequential dataset
scan as in HDFS)

bucket != chunk!

A Design Alternative: Distributed Hashing

- Distributed Hashing challenges:
 - Dynamicity: grow and shrink rapidly
 - Distribution: Assign buckets to participating nodes
(all the nodes should share the hash function for it to work)

E.g., $h(x) = x \% \#servers$



- Adding a new server implies modifying h...
 - Communicating the new h' to all servers
 - Re-hashing all the objects!
- Location of the hash directory: any access must go through the hash directory
 - Potential bottleneck

Consistent Hashing: Motivation

- Method initially proposed in the context of distributed caching
 - Results of the most common queries are in *caches* (i.e., in-memory) of several servers
 - A dedicated proxy machine records which server stores which query results
 - Queries are assigned to servers according to a hash function over the query
- Currently applied to distribution of data in distributed data stores
 - Supports high dynamicity of the infrastructure (servers may come and go at a rapid pace)
 - Most current key-value (and document-stores) use distributed hashing
 - Memcached – Key-value
 - Voldemort – Key-value
 - Cassandra – Wide-column
 - SimpleDB / DynamoDB – Key-value + Document
 - CouchDB - Document
 - MongoDB (current release) - Document

Consistent Hashing

- Coping with dynamicity:

- The hash function never changes

- Choose a large domain D (address space) and map server IPs and object keys to such domain

- Organize D as a ring so each node has a successor (clockwise)

- Objects are assigned as follows:

- For an object O , $f(O) = D_o$

- $D_{o'}$ and $D_{o''}$ are two nodes in the ring such that

- $D_{o'} < D_o <= D_{o''}$



- O is assigned to D_o

- Adding a new server is straightforward

- It is placed in the ring and part of its successors objects transferred

- Further refinements:


- Assign to the same server several hash values (virtual servers) to balance load (and deal with heterogeneity)



Heterogeneity: servers with different capacity

Activity: Consistent Hashing

- Objective: Understand how consistent hashing works
- Tasks:
 1. (5') By pairs, solve the following exercise
 - What happens in the structure when we register a new server with IP address "37"? Draw the result
 2. (5') Discussion




Document-Oriented DBs

Key-value enhancements



Integrated Databases vs Application Databases

- SQL and relational databases played a key role as integration mechanism between applications
 - Multiple applications using a common integrated database
 - All applications operate on a consistent set of persistent data
 - More complex structure
 - Changes by different apps need to be coordinated
 - Different apps have different performance needs, thus call for different index structures
 - Complex access policies



A different approach, treat your database as an **application database**

Application Databases

- An application database is only directly accessed by a single application, which makes it much **easier to maintain and evolve**
- Interoperability concerns can now shift to the interfaces of the application
 - During the 2000s we saw a shift to web services, where applications would communicate over HTTP (Service-Oriented Architecture)
- You are able to use **richer data structures** (compared to SQL)
 - Usually represented as documents in XML or, more recently JSON

The application is
responsible for database
integrity


Aggregate data models

- The relational model divides the information that we want to store into **tuples** (rows): this is a very simple structure for data
- **Aggregate orientation** takes a different approach. It recognizes that often you want to operate on data in units that have a **more complex structure**
 - Think of it as a complex record that allows lists and other record structure to be nested inside
- Key-value, document and column-family DBs can all be seen as aggregate-oriented databases
 - They differ in how they structure the aggregate and consequently how they allow for it to be accessed

Aggregate data models

Aggregate: collection of related objects that we wish to treat as a unit

- Example: A purchase order



In RDBMS: FKs and joins

Aggregate data models: benefits

- Dealing with aggregates makes it much easier for the databases to handle **operating on a cluster**, since the aggregate makes a natural unit for replication.
 - Also a natural unit to use for distribution (all the data for an aggregate stored together in one node)
 - And the atomic unit for updates (transactional control)
- They reduce **the impedance mismatch problem**, i.e., the difference between the stored data and the in-memory data structures



From K-V to Documents: Structuring the value

- Essentially, Document stores are Key-Value stores
 - Same design and architectural features
- The value is a document
 - XML (e.g., eXist)
 - JSON (e.g., MongoDB and CouchDB)
- Tightly related to the Web
 - Normally, they provide RESTful HTTP APIs
- So... what is the benefit of having documents?
 - New data model (collections and documents)
 - New atom: from rows to documents
 - Indexing

Includes data + metadata
(structure)


Structure provides more
flexibility in access

Types of Document Stores: XML

- XML is a semistructured data model proposed as the standard for data exchange on the Web
 - Can be represented as a tree
 - Document: the root node of the XML document, denoted by "/"
 - Element: element nodes that correspond to the tagged nodes in the document
 - Attribute: attribute nodes attached to Element nodes
 - Text: text nodes, i.e., untagged leaves of the XML tree
 - Supports XPath, XQuery and XSLT
 - XQuery is a query language for extracting information from collections of XML documents
 - XPath is a language for addressing portions of an XML document
 - Subset of XQuery
 - XSLT is a language for specifying transformations (from XML to XML)
 - XML document stores
 - eXist, MarkLogic
 - Natively supported
 - XML extensions for Oracle, PostgreSQL, etc.
 - Mapped to relational (impedance mismatch!)

XML Example

```
<?xml version="1.0"
      encoding="utf-8"?>
<A>
  <B att1='1'>
    <D>Text 1</D>
    <D>Text 2</D>
  </B>
  <B att1='2'>
    <D>Text 3</D>
  </B>
  <C att2="a"
     att3="b"/>
</A>
```



An XML document is a *labeled, unranked, ordered tree*

Types of Document Stores: JSON

- JSON is a lightweight data interchange format
 - Brackets ([]) represent ordered lists
 - Curly braces ({}) represent key-value dictionaries
 - Keys must be strings, delimited by quotes ("")
 - Values can be strings, numbers, booleans, lists, or key-value dictionaries
- Natively compatible with JavaScript
 - JSON stands for JavaScript Object Notation
 - Web browsers are natural clients for MongoDB / CouchDB
- JSON document stores
 - MongoDB, CouchDB
 - Natively supported
 - JSON extensions for Oracle, PostgreSQL, etc.
 - Mapped to relational (impedance mismatch!)

<http://www.json.org>

JSON Example

- Definition:
 - A document is an object represented with an unbounded nesting of array and object constructs

A document is just a set of key-value pairs (a dictionary)

```
{  
    "title": "The Social network",  
    "year": "2010",  
    "genre": "drama",  
    "summary": "On a fall night in 2003, Harvard undergrad and computer programming genius Mark Zuckerberg sits down at his computer and heatedly begins working on a new idea. In a fury of blogging and programming, what begins in his dorm room soon becomes a global social network and a revolution in communication. A mere six years and 500 million friends later, Mark Zuckerberg is the youngest billionaire in history... but for this entrepreneur, success leads to both personal and legal complications.",  
    "country": "USA",  
    "director": {  
        "last_name": "Fincher",  
        "first_name": "David",  
        "birth_date": "1962"  
    },  
    "actors": [  
        {  
            "first_name": "Jesse",  
            "last_name": "Eisenberg",  
            "birth_date": "1983",  
            "role": "Mark Zuckerberg"  
        },  
        {  
            "first_name": "Rooney",  
            "last_name": "Mara",  
            "birth_date": "1985",  
            "role": "Erica Albright"  
        },  
        {  
            "first_name": "Andrew",  
            "last_name": "Garfield",  
            "birth_date": "1983",  
            "role": "Eduardo Saverin"  
        },  
        {  
            "first_name": "Justin",  
            "last_name": "Timberlake",  
            "birth_date": "1981",  
            "role": "Sean Parker"  
        }  
    ]  
}
```

MongoDB

An example of Document Store



MongoDB Data Model

- Collections
 - Definition: A grouping of MongoDB documents
 - A collection exists within a single database
 - Collections do not enforce a schema
 - MongoDB Namespace: *database.collection*


Pros and Cons of schema-less?

- Documents
 - Definition: JSON documents (serialized as BSON)
 - Basic atom
 - Aggregated view of data
 - Identified by `_id` (user or system generated)
 - May contain
 - References (NOT FKs!)
 - Embedded documents

IDs are globally unique

MongoDB Document Example

- Ordered set of keys with associated values
- Data structure:
 - Map, Hash, Dictionary or Object → JSON (BSON)
e.g., {"greeting" : "Hello, world!", "foo" : 3}
- Keys in a document must be Strings
 - No duplicate keys




MongoDB Document Example

- Plain document

```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact_phone: "123-456-7890",  
  contact_email: "xyz@example.com",  
  access_level: 5  
  access_group: "dev"  
}
```

- Embedded sub-documents

```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



- Array of sub-documents

```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contacts: [  
    { type: "work", phone: "123-456-7890", email: "xyz@example.com" },  
    { type: "home", phone: "098-765-4321", email: "xyz@home.com" },  
  ]  
  access_level: 5  
  access_group: "dev"  
}
```



Designing Documents

- Follow one basic rule: 1 fetch for the whole data object at hand
 - Aggregate data model: check the data needed by your application simultaneously (queries)
 - Do not think relational-wise!
 - Use indexes to identify finer data granularities
- Consequences:
 - Independent documents
 - Avoid pointing at other docs
 - Massive denormalization
 - A change in the application layout might be dramatic
 - It may entail a massive rearrangement of the database documents

From rows to documents ->
New design strategies!

Activity: Modeling in MongoDB


- Objective: Learn how to model documents
- Tasks:
 1. (15') Model the TPC-H database
 2. (5') Discussion



Activity: Modeling in MongoDB

- Objective: Learn how to model documents
- Tasks:
 1. (15') Model the TPC-H database
 - According to the query below
 2. (5') Discussion

```
SELECT l_orderkey,  
sum(l_extendedprice*(1-l_discount)) as  
revenue, o_orderdate, o_shippriority  
FROM customer, orders, lineitem  
WHERE c_mktsegment = '[SEGMENT]' AND  
c_custkey = o_custkey AND l_orderkey =  
o_orderkey AND o_orderdate < '[DATE]'  
AND l_shipdate > '[DATE]'  
GROUP BY l_orderkey, o_orderdate,  
o_shippriority  
ORDER BY revenue desc, o_orderdate;
```



MongoDB API

MongoDB Shell


- show dbs
- show collections
- show users
- use <database>
- coll = db.<collection>
 - coll.insert(*document*)
 - coll.save(*document*) -- updates an existing document or inserts a new document
 - coll.find(*query, projection*) -- coll.find({name:"Joe"}, { name: true })
 - coll.update(*query, update*)
 - coll.deleteOne(*query*) or deleteMany(*query*); -- removes one or many docs from a collection
 - coll.drop(); -- removes a collection from the database
 - coll.createIndex(*keys, options*); -- creates an index on the specified fields
- Notes:
 - *db* refers to the current database
 - *query* is a document (query-by-example)

<https://www.mongodb.com/docs/mongodb-shell/>

MongoDB: Querying

- Find and findOne methods
 - database.collection.find()
 - database.collection.find({ qty: { \$gt: 25 } })
 - database.collection.find({ field: { \$gt: value1, \$lt: value2 } });
- The Aggregation Framework
 - An aggregation pipeline
 - Documents enter a multi-stage pipeline that transforms the documents into an aggregated result
 - Filters that operate like queries
 - Document transformations that modify the form of the output
 - Grouping
 - Sorting
 - Other operations
- MapReduce (deprecated)

MongoDB: Aggregation Framework



<https://docs.mongodb.com/manual/reference/operator/aggregation-pipeline/>


MongoDB Architecture

MongoDB Notation: Shard Clusters

- Shards
 - Nodes containing data
 - A shard may contain several chunks
- Config Servers
 - Nodes containing the global catalog (e.g., hash directory)
- Query routers or Route servers
 - Nodes containing a copy of the hash directory to redirect queries

Documents in a collection are distributed across shards

MongoDB Architecture



<http://horicky.blogspot.com.es/2012/04/mongodb-architecture.html>

Shard Clusters Management


- Query routers are replicas of the config servers
 - Secondary versioning (config servers)
 - Eager replication (to both config servers and query routers)
 - 2PCP (potential distributed deadlocks!)
- Config Servers
 - The hash directory is mandatorily replicated to avoid single-point failures
 - MongoDB asks for 3 config servers
 - Writes happen if:
 - A shard splits
 - A chunk migrates between servers (e.g., adding servers)
- Query routers
 - Read from config servers
 - When they start (or restart)
 - Every time a split / migration happens

Splitting/Migrating Chunks

- Default chunk size: 64MB
- The query router (mongos) asks a shard to split
 - Inserts and updates trigger splits
- Shards rearrange the data (data migration)
 - During the migration, requests to that chunk address the origin shard
 - Changes made during the migration are afterwards applied in the destination shard
 - Finally, changes in the hash directory are made in the config servers
 - Query routers eagerly synchronized
- A *balancer* avoids uneven distributions

Replication

- Each shard (in a shard cluster) is a replica set
 - Maps to a mongod instance (with its config servers)
- Replica Set: Master versioning with lazy replication
 - One master
 - Write / Update / Delete
 - Several replicas
 - Reads
- Replica Set management
 - The master has a recovery system (journaling): WAL
 - Members interconnected by heartbeats
 - If the master fails, voting phase to decide a new master
 - If a replica fails, it catches up with the master once back



Pluggable Storage Engine Architecture

- MongoDB 3.0 introduced the concept of “pluggable storage engine”
 - MMAP V1 – based on consistent hashing (see previous slides)
 - WiredTiger – based on LSM (distributed B+)
 - In-Memory
- WiredTiger
 - Moves to LSM (welcome range queries!)
 - We can choose if store data row-oriented or column-oriented
 - First boosts writes, hurts reads. Just the opposite for the second one
 - First-class citizen compression
- Each node on a MongoDB cluster can use a different storage engine

Query Optimization

- The aggregation framework creates a left-deep tree access plan and applies pipelining
 - Note that the first operation is executed in parallel in all nodes containing data (exploiting data locality)
 - From there on, a node takes care of the query and data is shipped to it to execute the rest of the pipeline
- MongoDB barely applies any optimization technique in its querying flow
 - First versions: Nothing!
 - From version 2.6: Primitive rule-based optimization approach
<https://www.mongodb.com/docs/manual/core/aggregation-pipeline-optimization/>
- Be careful when creating your pipes (you are the most important optimizer!)
 - Push selections and projections to the beginning of the pipeline
 - A cost-based approach badly needed...

<https://www.mongodb.com/docs/manual/core/query-optimization/>

Indexing

- Indexes are (physically) the same as in a relational database. Same rules apply:
 - Selective queries
 - Must fit in memory
- However, indexing management is way poorer
 - No cost-based models
 - For a new query, all indexes are run in parallel and the best plan is chosen from there on (sigh...)
 - The plan is recalculated when massive inserting happens or when the database restarts
- Better you do the job
 - Monitor your queries:
 - <https://www.mongodb.com/docs/manual/tutorial/analyze-query-plan/>
 - Use \$hint to force MongoDB choose an index
 - https://www.mongodb.com/docs/manual/reference/method/cursor_hint/

Limitations (I)

- Architectural Issues
 - Thumb rule: 70% of the database must fit in memory
 - Be careful with updates! (padding)
 - Holes caused by reallocation
 - Compact the database from time to time
 - In WiredTiger this is left for the compactation (the delta memstore smooths it)
 - Limited number of collections per database (10.000)
 - Theoretically, sharding is automatic and transparent. But in practice it is not. Most typical ones:
 - Max. number of elements to migrate (when balancing the workload)
 - LSM + sequential keys will hit only one node (be careful with the key!!)
- Document Issues
 - The resulting document of an aggregation pipeline cannot exceed the maximum document size (16Mb)
 - GridFS for larger documents
 - No more than 100 nesting levels (i.e., embedded documents nesting)
 - Attribute names are kept as they are (no catalog)

<https://www.mongodb.com/docs/manual/reference/limits/>

Limitations (II)

- Querying Issues
 - Limited transaction support
 - ACID guarantees only at document level
 - Strong / loose consistency parametrizable (write-concern) <https://www.mongodb.com/docs/manual/reference/write-concern/>
 - Thumb rule: A query must attack a single collection
 - Arrays are a mess! → \$unwind
 - No optimizer!
 - Be careful with distributed queries (may have to be sent to ALL shards): <https://www.mongodb.com/docs/manual/core/distributed-queries/>
- The aggregation framework
 - Pipe stages are limited to 100MB of RAM: <https://www.mongodb.com/docs/manual/core/aggregation-pipeline-limits/>
 - Disk usage for bigger pipeline operations must be specified
 - Performance deteriorates hugely!
 - Does not solve the problem in many cases
- Indexing
 - MongoDB runs all the queries with all indexing possibilities in parallel: <https://www.mongodb.com/docs/manual/core/query-plans/>
 - From there on, that index will be always chosen (!!!)
 - Check its performance with .explain() to see the access plan

MongoDB: Conclusions

- MongoDB has its limitations, but it is one of the most supported and mature NOSQL tools
 - Still, its robustness is far away from a relational database
- Many tools and functionalities available
 - Managing and Monitoring
 - OpsManager: Be careful! The terms of use say your data is periodically sent to MongoDB (the Company)
 - db.collection.explain()
 - GUI: MongoDB Compass
 - Supporting GeoSpatial data and queries
 - Support from 3rd parties
 - Tableaux
 - Pentaho BI Suite
 - Cubes: OLAP-lightweight Engine (<http://cubes.databrewery.org/>)
 - Good pluggability with almost any language (Python, Ruby, Perl, Java, Scala, PHP, etc.)
 - Most Cloud providers offer MongoDB as a service
 - Amazon, DigitalOcean, Rackspace, Openshift, Azure, etc.
 - Compose (MongoDB as a Service) + Heroku (great combo!)

Summary

- Document stores
 - Semi-structured value
 - Indexing
- Designing document stores

Bibliography

- S. Abiteboul et al. Web Data Management, 2012
- J. Dittrich et al. Hadoop++: Making a yellow elefant run like a cheetah (without it even noticing)
- D. Jiang et al. The performance of MapReduce: An In-depth Study. VLDB'10
- E. Brewer, "Towards Robust Distributed Systems," Proc. 19th Ann. ACM Symp.Principles of Distributed Computing (PODC 00), ACM, 2000, pp. 7-10.
- F. Chang et all. Bigtable: A Distributed Storage System for Structured Data. OSDI'06
- Sanjay Ghemawat et al. The Google File System. OSDI'03
- Jeffrey Dean et al. MapReduce: Simplified Data Processing on Large Clusters. OSDI'04
- D. Battre et al. Nephele/PACTs: A Programming Model andExecution Framework forWebScale Analytical Processing. SoCC'10
- L. Liu and M.T. Özsü (Eds.). Encyclopedia of Database Systems. Springer, 2009
- P. Sadalge, M. Fowler. NoSQL Distilled. Addison Wesley, 2012.