

# Data Integration

Past, Present and Future



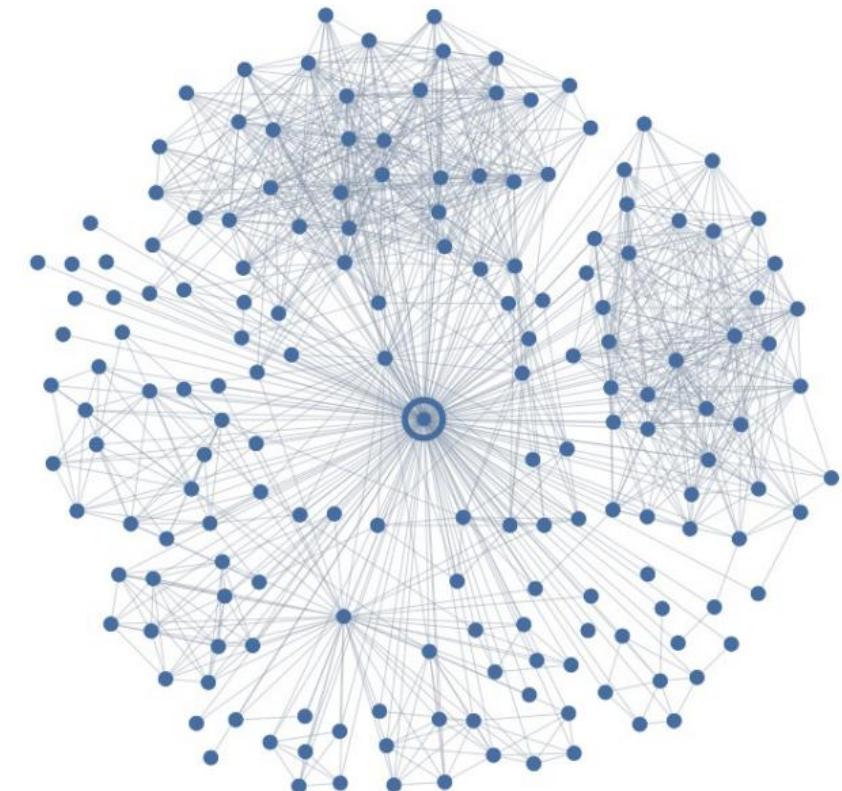
**“WITHOUT DATA,  
YOU’RE JUST  
ANOTHER PERSON  
WITH AN OPINION”**

W. Edwards Deming, American Statistician



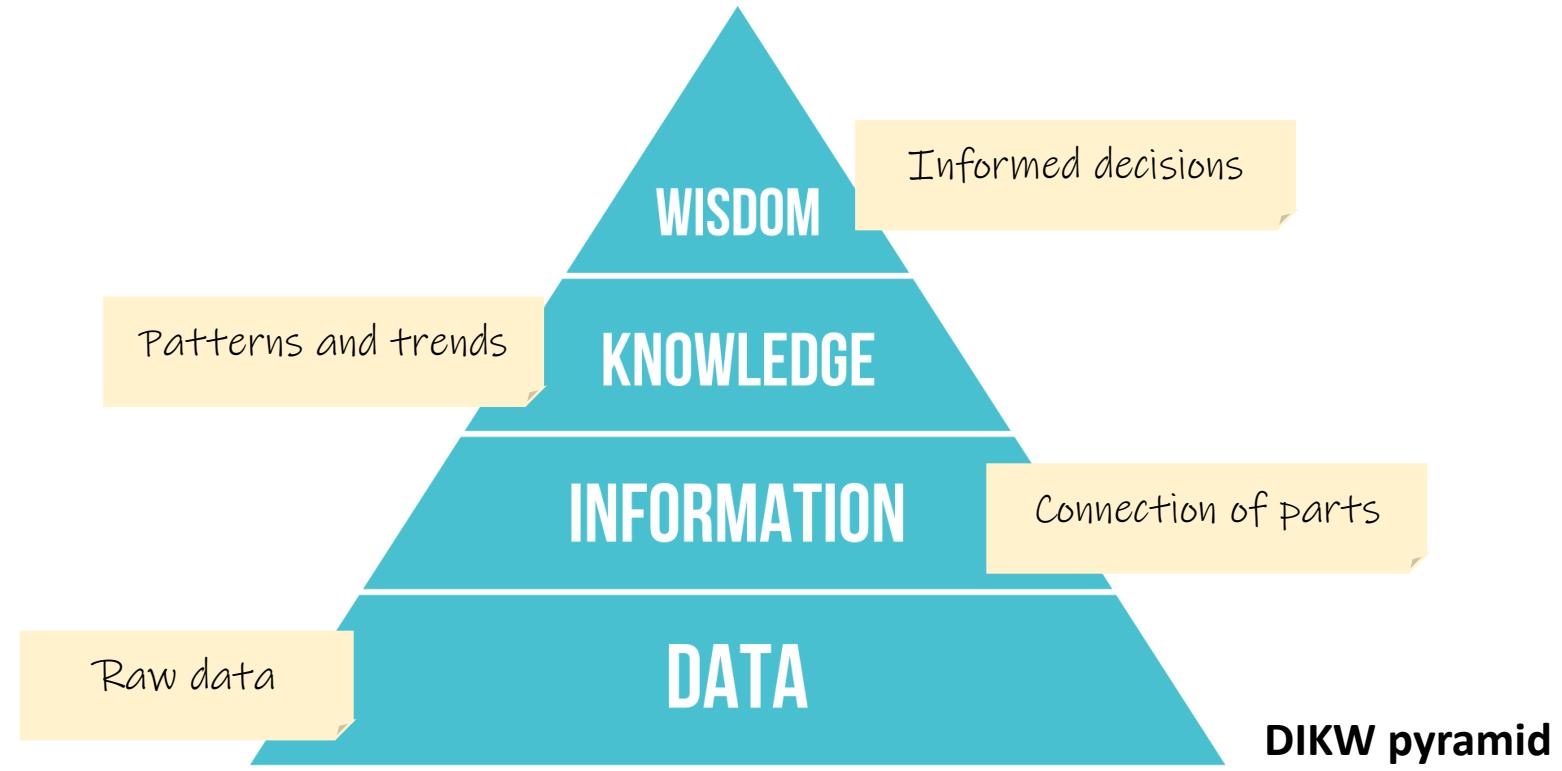
# An Example: Will They Split Up?

- Crossing data from social networks it is possible to identify a graph like the one that follows:
  - In the centre there is a specific person  $P$
  - The rest are  $P$  connections and connections among them
- Using sociology techniques...
  - We can identify  $P$  social foci:
    - Dense clusters of friends corresponding to long periods of interaction
    - Typically, college friends, coworkers, relatives, etc.
  - The *significant other* can be identified by a high *dispersion* rate
    - Highly connected with  $P$  connections,
    - But with a high dispersion degree wrt  $P$  social foci
- **Hypothesis:** when the node with higher dispersion degree Identified is not the partner, this couple is likely to split up in a period of 60 days
- L. Backstrom, J. Kleinberg. Romantic Partnerships and the Dispersion of Social Ties: A Network Analysis of Relationship Status on Facebook  
<https://arxiv.org/pdf/1310.6753v1.pdf>



# Decision Support Systems

- IT systems aimed at exploiting data and transform it into information and knowledge

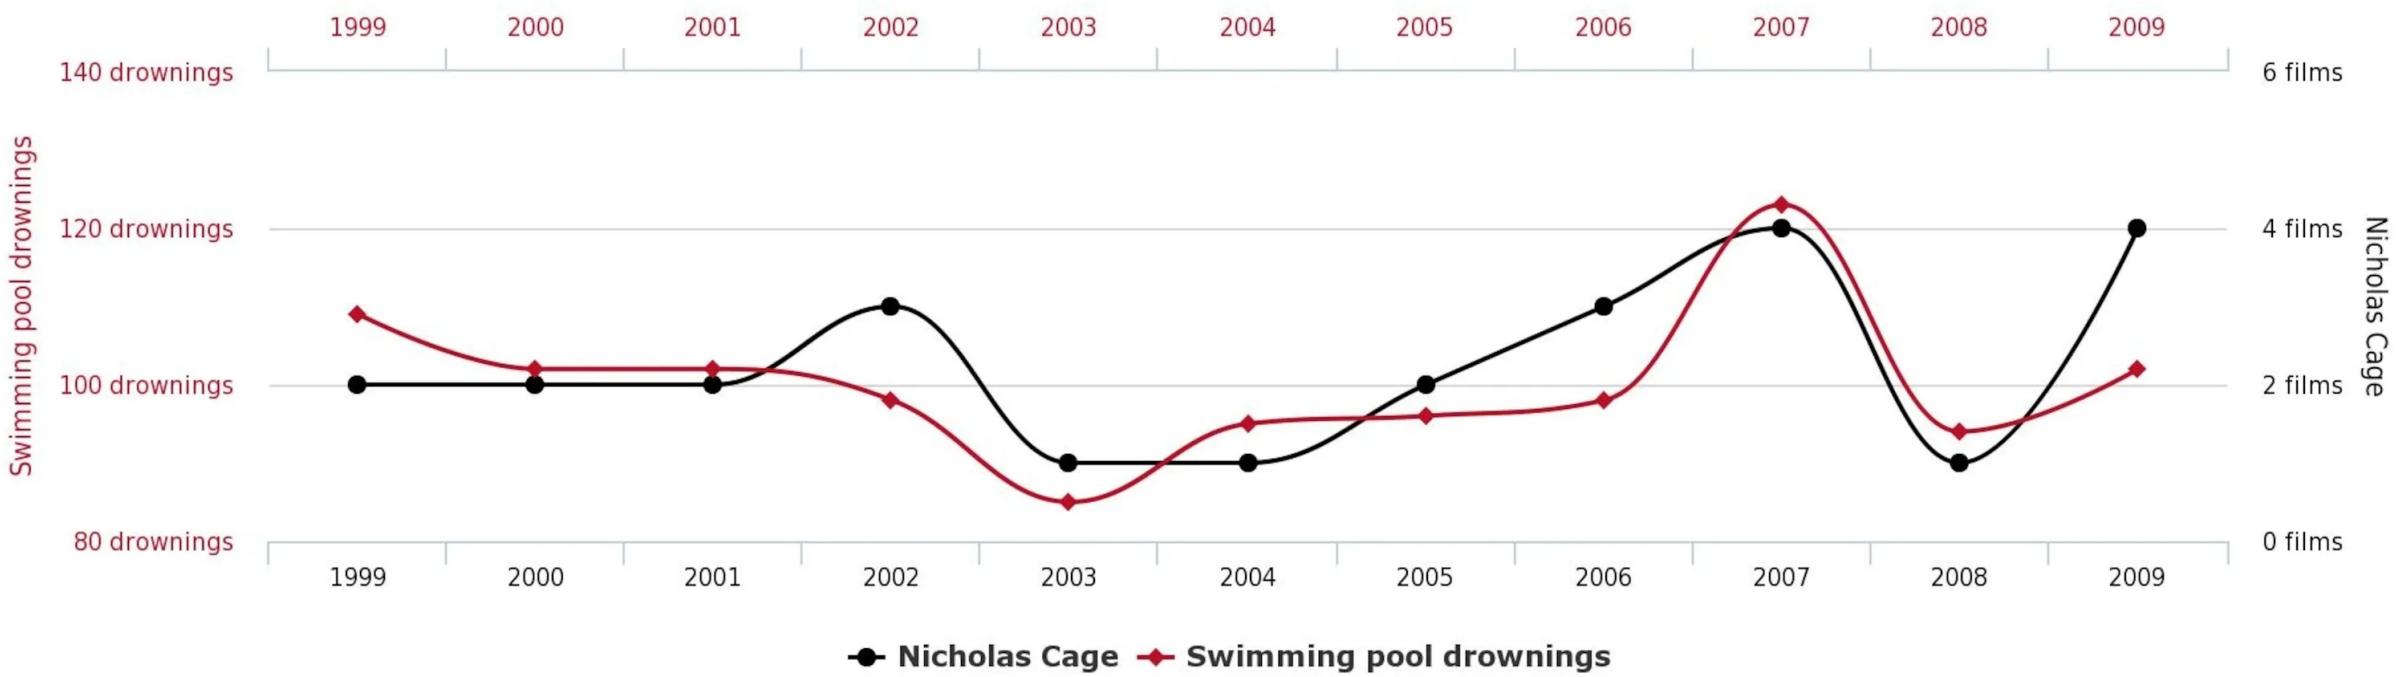


# From Information to Knowledge: a Counterexample

**Number of people who drowned by falling into a pool**

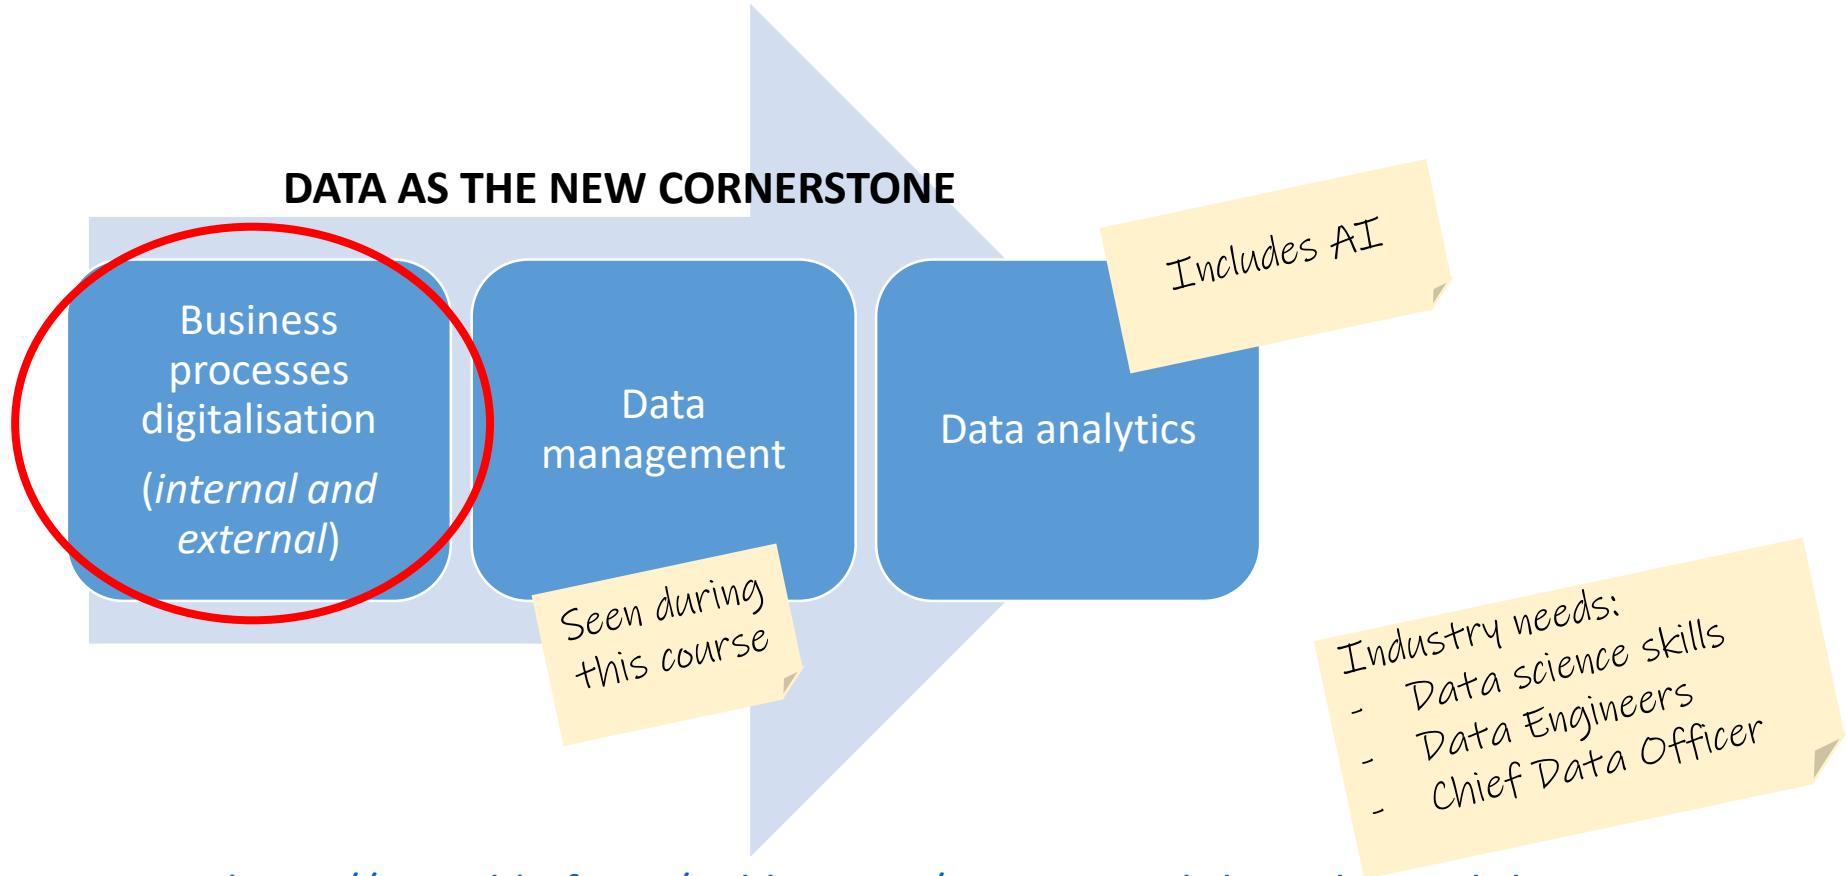
correlates with

**Films Nicolas Cage appeared in**



[tylervigen.com](http://tylervigen.com)

# Decision Support Systems



The IBM Quant Crunch report, 2017: <https://www.bhef.com/publications/quant-crunch-how-demand-data-science-skills-disrupting-job-market>

# Linking Data, Analysis and Business

## This is why AI has yet to reshape most businesses

For many companies, deploying AI is slower and more expensive than it might seem.

10% AI  
90% Data

by Brian Berastein

February 13, 2019

Biggest obstacles:

- Get and prepare all required data
- Lack of qualified professionals
- Reluctance of end users (trust)

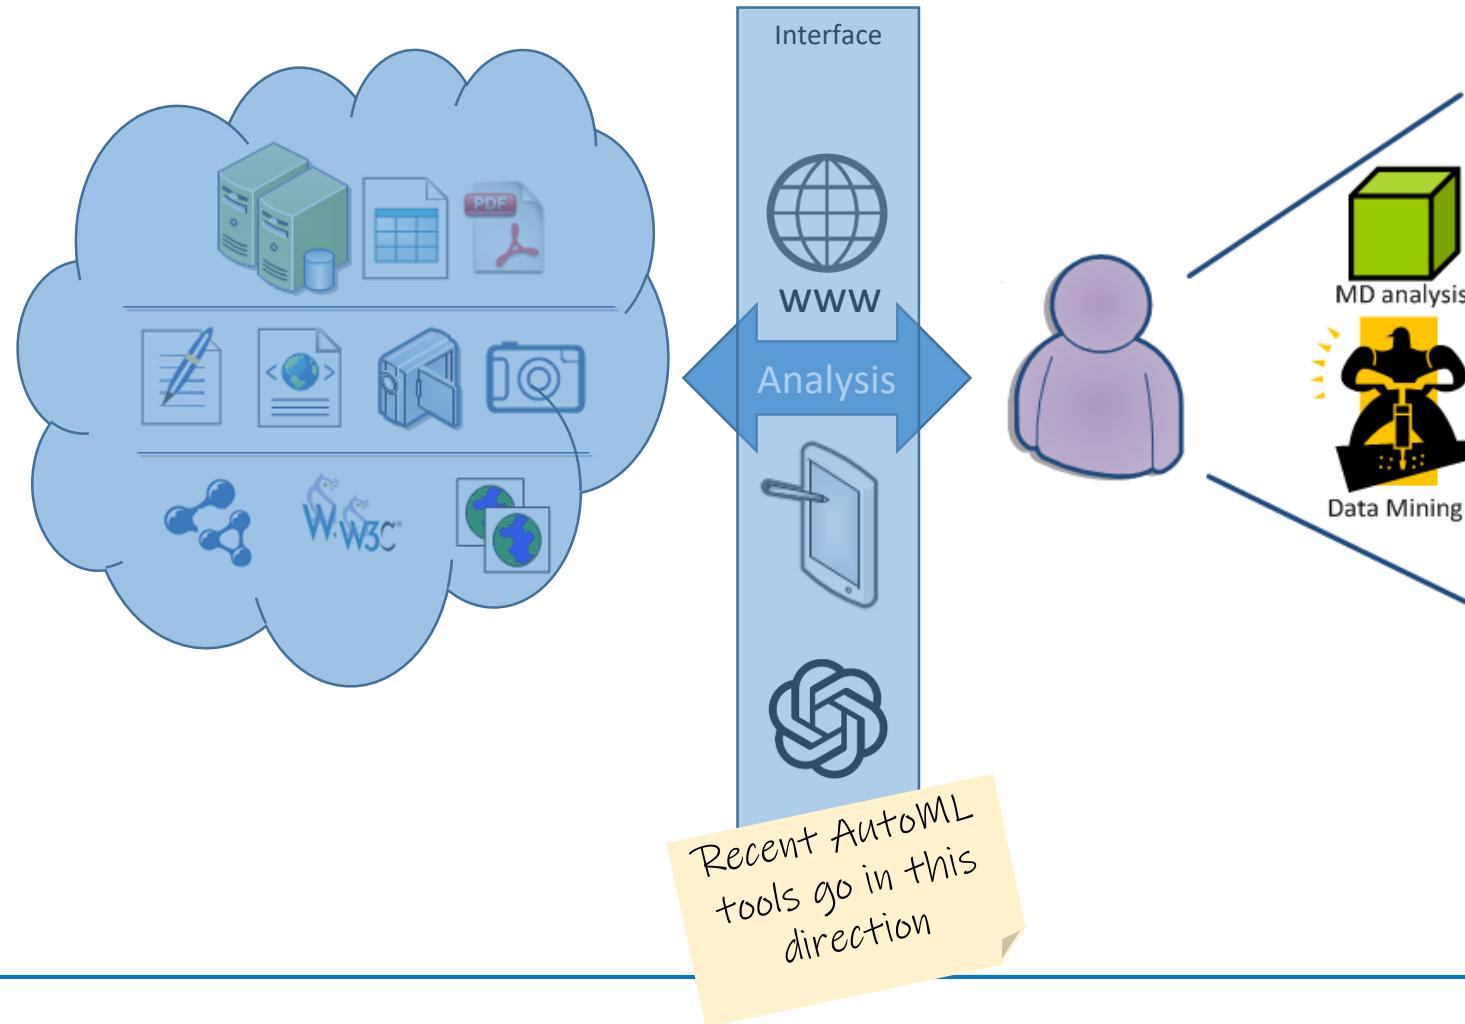
...link. <https://www.technologyreview.com/2019/02/13/137047/this-is-why-ai-has-yet-to-reshape-most-businesses/>

Solutions:

- Automation of data-related tasks
- Data literacy (citizen data scientist)

# Data Analysis Democratisation

- From the point of view of the end-user:



# Examples

- A company wants to enrich their DW data (*products, customers, shops*) with external data coming from Twitter (*opinion about their products*), logs generated by their web and feedback received by phone, web or the company app. The constraint is that external data must be loosely coupled but be completely integrated from an analysis point of view
- A journalist wants to analyze the evolution of asylum demands in Europe in the last 10 years. What data are available in the EU open data portals? How data should be fetched and crossed in order to answer questions?
- A company with a large dataset of information from their activities wants to deploy a flexible Machine Learning pipeline. They want to avoid data analysts spend 80% of their time pre-processing data. There are many reasons: (i) data analysts should not repeat the same transformations once and again, (ii) they should collaboratively share their knowledge in pre-processing data and (iii) they want to avoid lock-in knowledge (i.e., analysts keep their knowledge in personal scripts not shared with the company)

# Exercise 1: Looking for clusters visually

From the course *Transition to Data Science*. [Buy the entire course for just \\$10](#) for many more exercises and helpful video lectures.

You are given an array `points` of size  $300 \times 2$ , where each row gives the  $(x, y)$  co-ordinates of a point on a map. Make a scatter plot of these points, and use the scatter plot to guess how many clusters there are.

**Step 1:** Load the dataset (written for you).

```
In [1]: import pandas as pd  
  
df = pd.read_csv('../datasets/chlex1.csv')  
points = df.values
```

Read data from temporal files

**Step 2:** Import PyPlot

```
In [2]: import matplotlib.pyplot as plt
```

Use off-the-shelf ML libraries

**Step 3:** Create an array called `xs` that contains the values of `points[:, 0]` - that is, column 0 of `points`.

```
In [3]: xs = points[:, 0]
```

**Step 3:** Create an array called `ys` that contains the values of `points[:, 1]` - that is, column 1 of `points`.

```
In [4]: ys = points[:, 1]
```

Pre-process the data to fit the needs of the algorithm input parameters

**Step 4:** Make a scatter plot by passing `xs` and `ys` to the `plt.scatter()` function.

```
In [5]: plt.scatter(xs, ys)
```

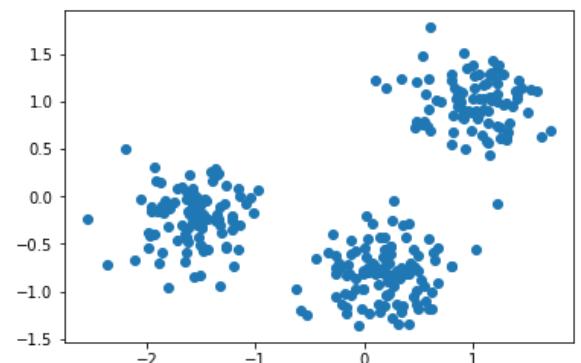
Run the algorithm

```
Out[5]: <matplotlib.collections.PathCollection at 0x110fc45c0>
```

**Step 5:** Call the `plt.show()` function to show your plot.

```
In [6]: plt.show()
```

Visualise the result to allow its interpretation



Extracted from: <https://github.com/benjaminwilson>

# Challenges in Data Management

From the IT point of view



# What is Big Data?

VOLUME

Veracity

*Velocity*

Value

VariaBiLiTy

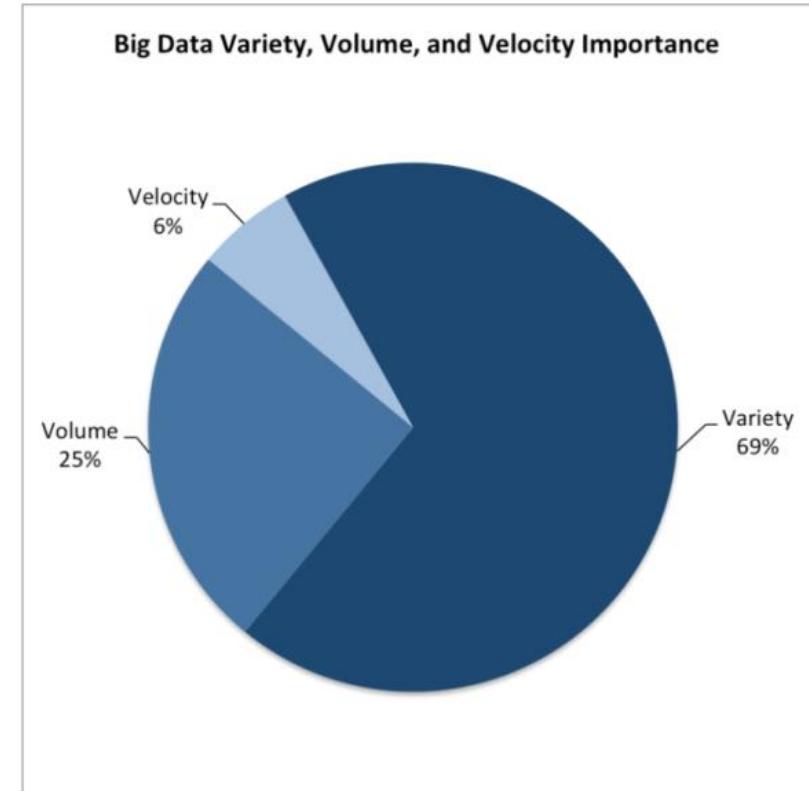
Variety

# Today, the Focus is on Variety

That Big Data is synonymous with large volumes of data is a **myth**

*“Rather, it is the ability to **integrate** more sources of data than ever before — new data, old data, big data, small data, structured data, unstructured data, social media data, behavioral data, and legacy data”*

The Variety Challenge



MIT Sloan Management Review (2016): <http://sloanreview.mit.edu/article/variety-not-volume-is-driving-big-data-initiatives/>

# The Problem: Motivation

Currently, most organizations use different databases for different applications

- Example:



Polyglot persistence, Martin Fowler

# The Problem: Motivation

When performing data analytics (e.g., reporting, recommendations, machine learning) we may need to extract / query data from different databases.

For example, to perform recommendations we may need to cross information from the product catalog, user sessions and user activity logs:

- Collaborative recommendations: find *similar* users and suggest to them what their similars bought
- Similarity is built based on their session and activity information, and also based on the characteristics of the product bought

# The Solution

*Data integration involves combining data residing in different sources and providing users with a unified view of them*

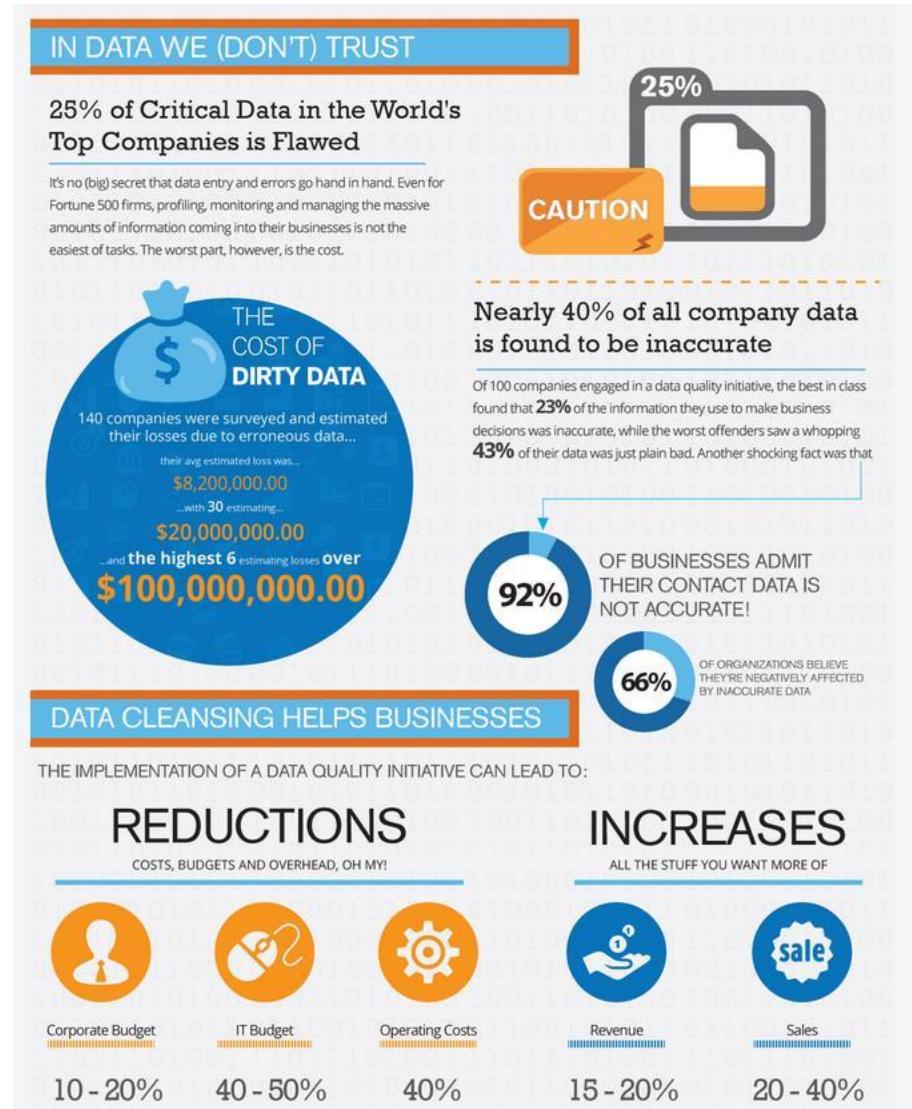
Maurizio Lenzerini (2002). "Data Integration: A Theoretical Perspective". PODS 2002. pp. 233–246.

Today: Data sources  
can be within or  
outside the  
organization

# Data Quality

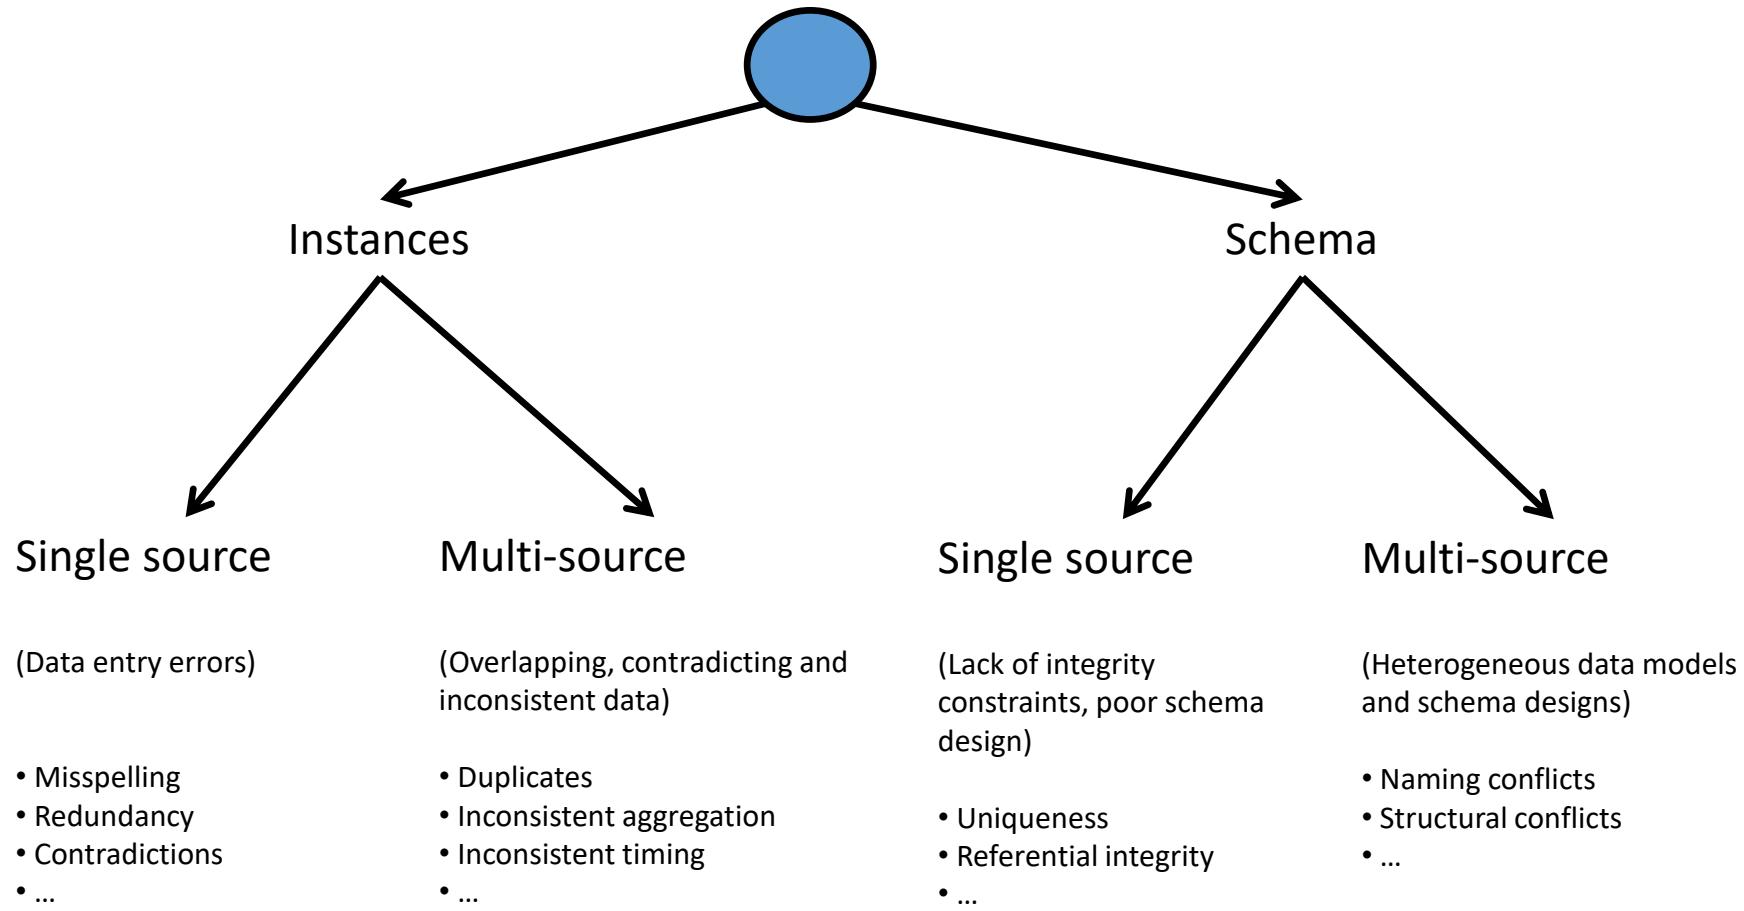
The Relevance of Data Cleaning in Data Integration

# Motivation



Can you think of some examples of data quality problems?

# Data Conflicts



*They may appear at every possible stage of a data pipeline!*

# Typical Quality Measures

- Completeness
- Accuracy
- Consistency
- Timeliness (Freshness)

# Completeness

“The degree to which a given collection of data describes the corresponding set of real-world objects”

- Issues:
  - Missing values (instance level)
  - Missing entities (instance level)
  - Missing attributes (schema level)
- Solution:
  - Imputation of values
  - Additional data sources

Not all NULL values have  
the same meaning or  
importance

# Accuracy

"The extent to which data are correct, reliable and certified error free"

- Issues:
  - Typing errors
  - Inappropriate precision
- Solution:
  - Using dictionaries or ontologies
  - Removing inexact values
  - Applying data cleaning techniques

Distance between the  
data and the value in the  
real world

# Consistency

“The degree of violation of semantic rules defined over a set of data items”

- Issues:
  - Integrity constraint violations
    - Entity, Domain, Referential, User-defined
  - Inconsistency between copies
    - Temporal or permanent
- Solution:
  - Removing or replacing inconsistent data

Rules can be specified by  
the user, or discovered  
from the data

# Timeliness (Freshness)

“How old the stored value of an attribute is with regard to the current value in the real world”

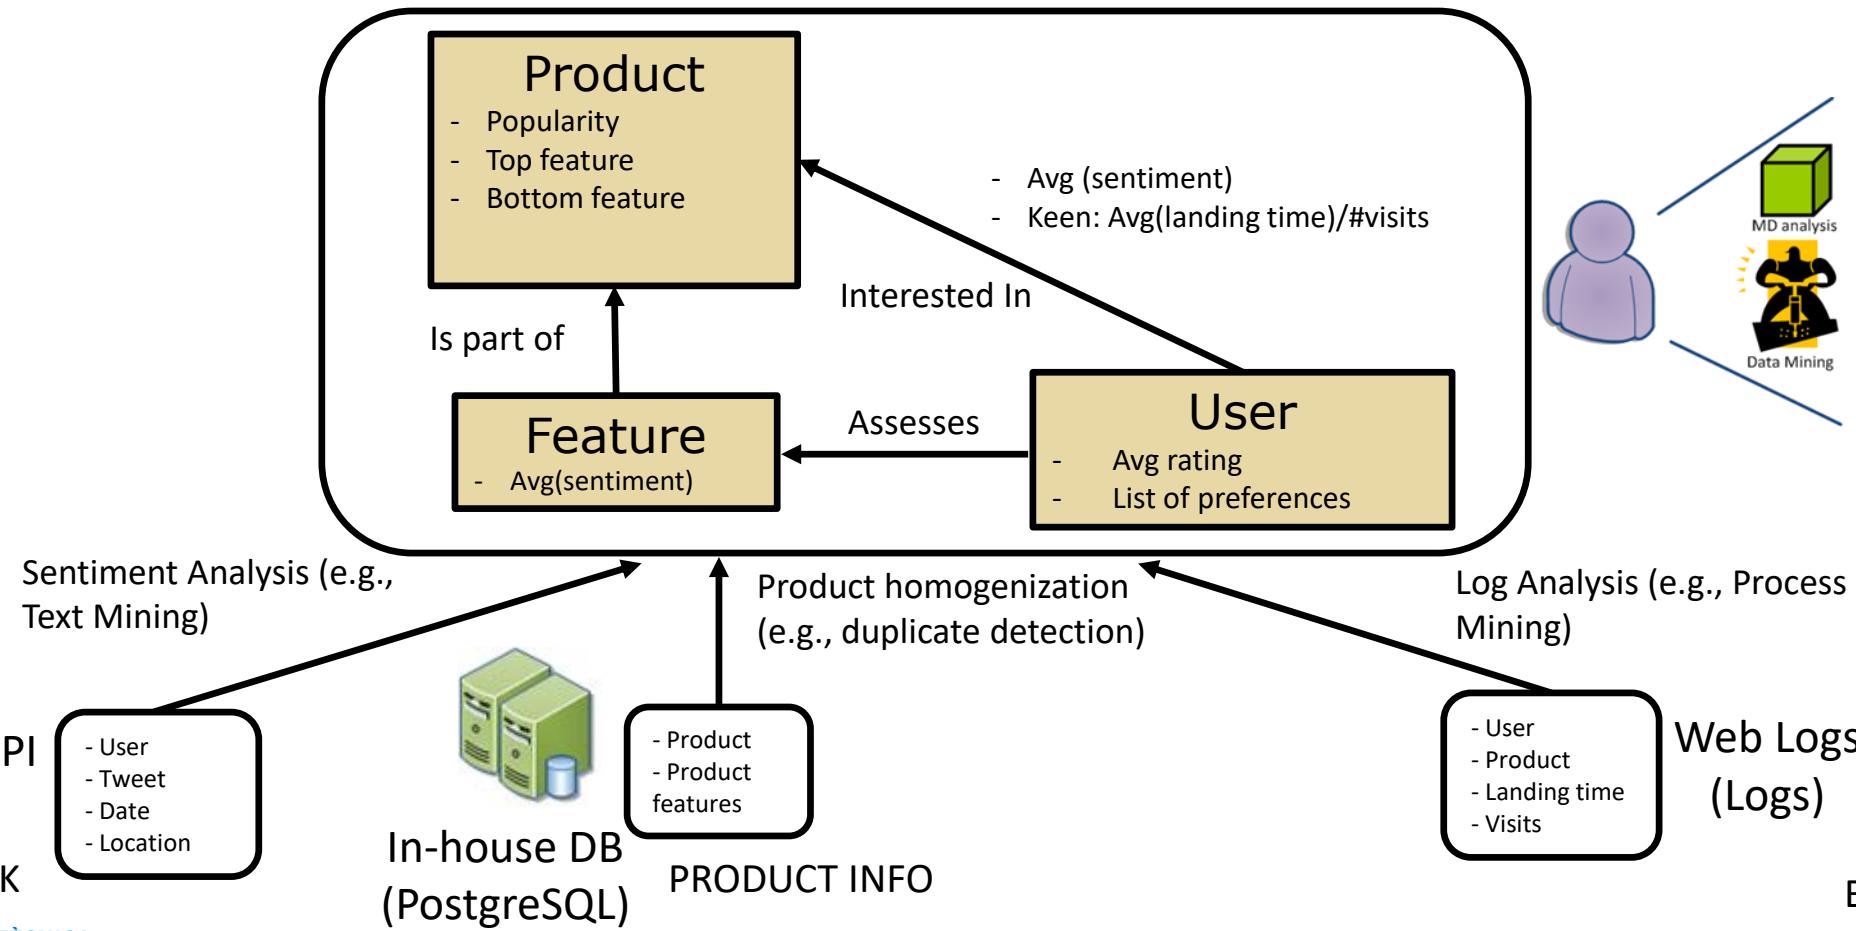
- Issues:
  - Outdated data
    - Depends on the update frequency
    - ...And of the goals of the user
- Solution:
  - Additional data sources

# The Future of Data Integration

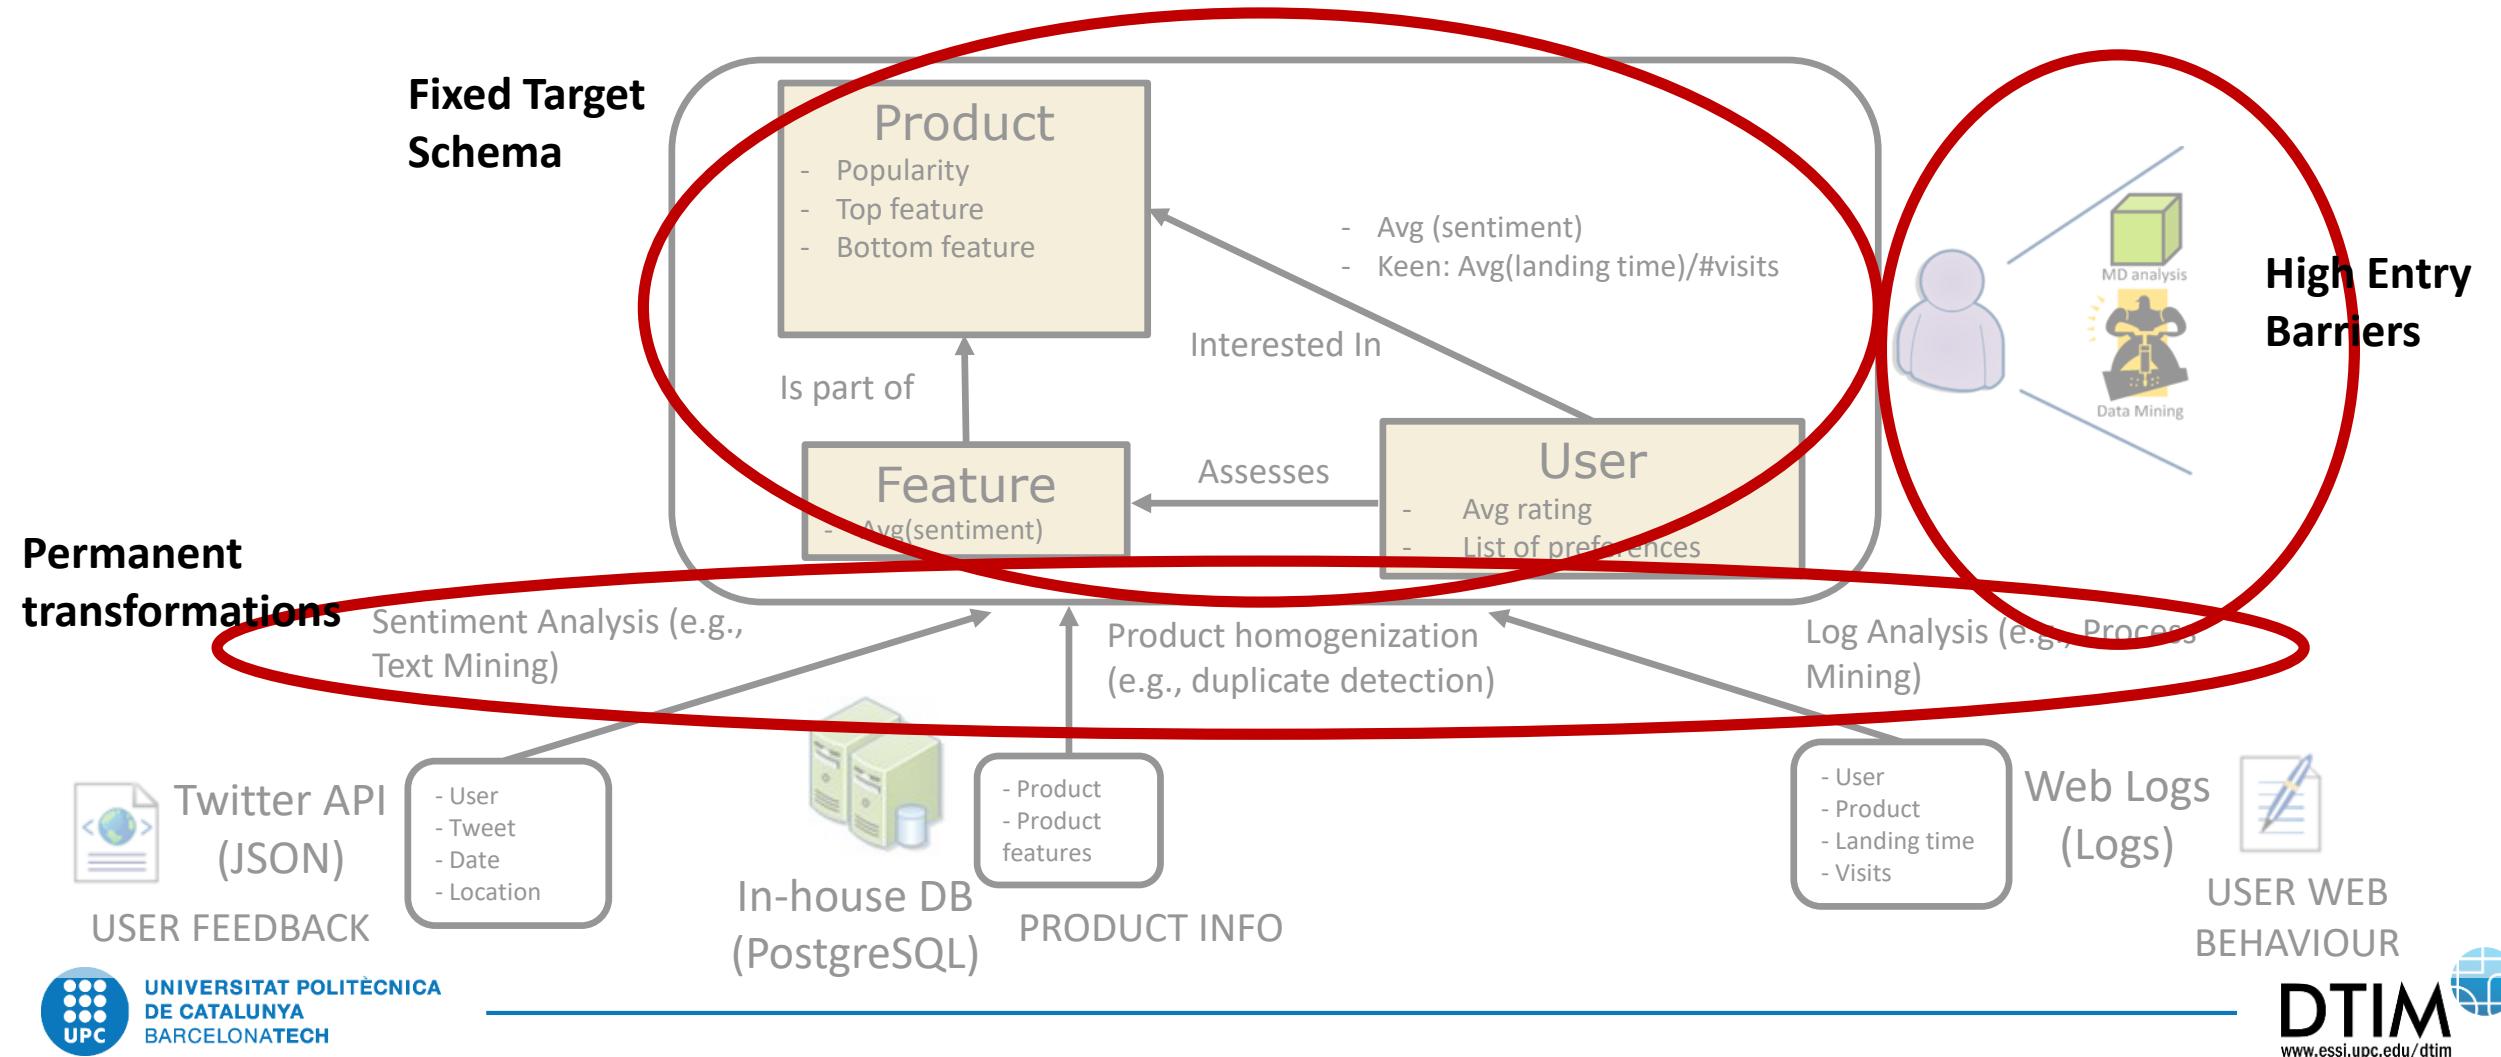
Data Lakes and DataOps

# Traditionally: Model-First (Load-Later)

Typical of Data Warehousing (ETL)



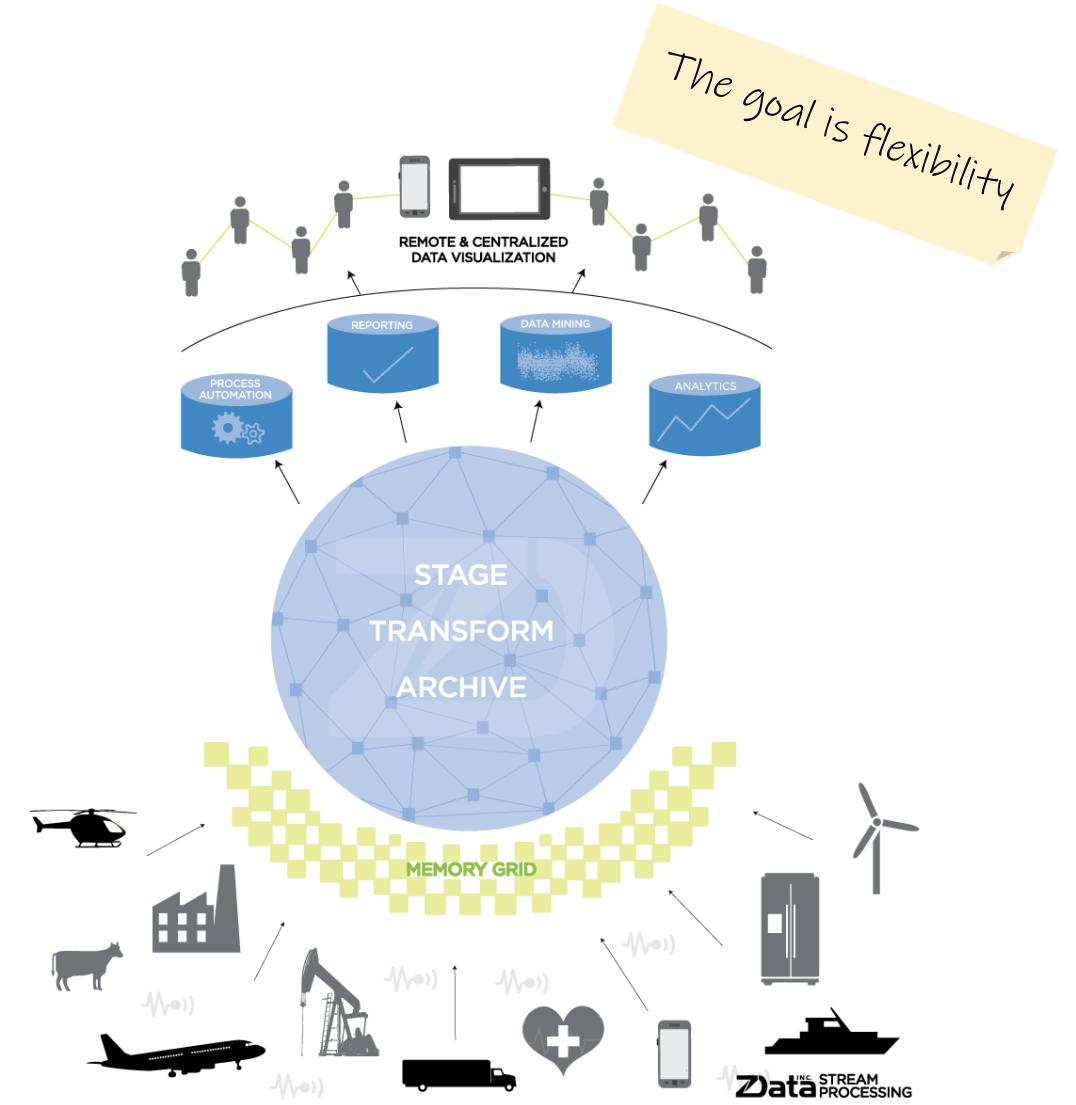
# Drawbacks



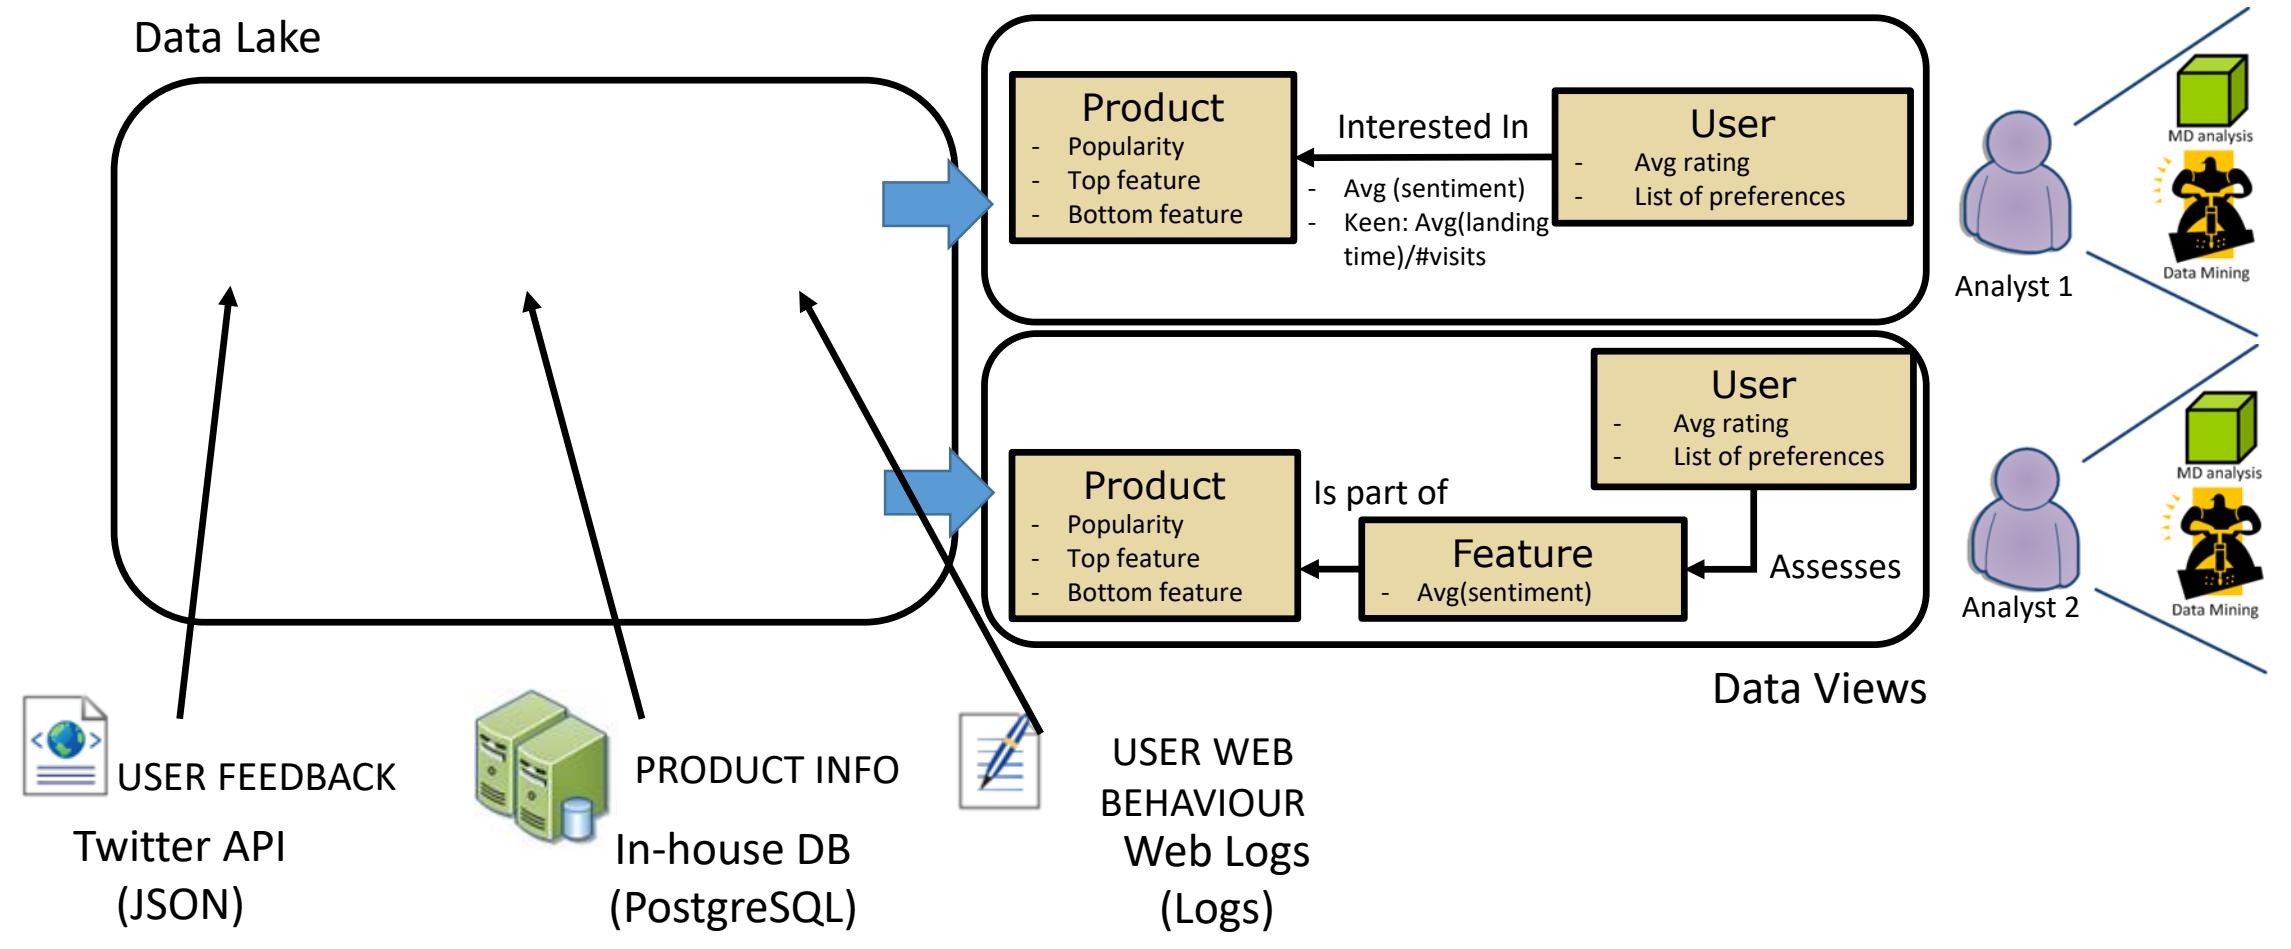
# The Data Lake

- IDEA: Load-first, Model-Later
- Modeling at load time restricts the potential analysis that can be done later (Big Analytics)
- Store raw data and create on-demand views to handle with precise analysis needs

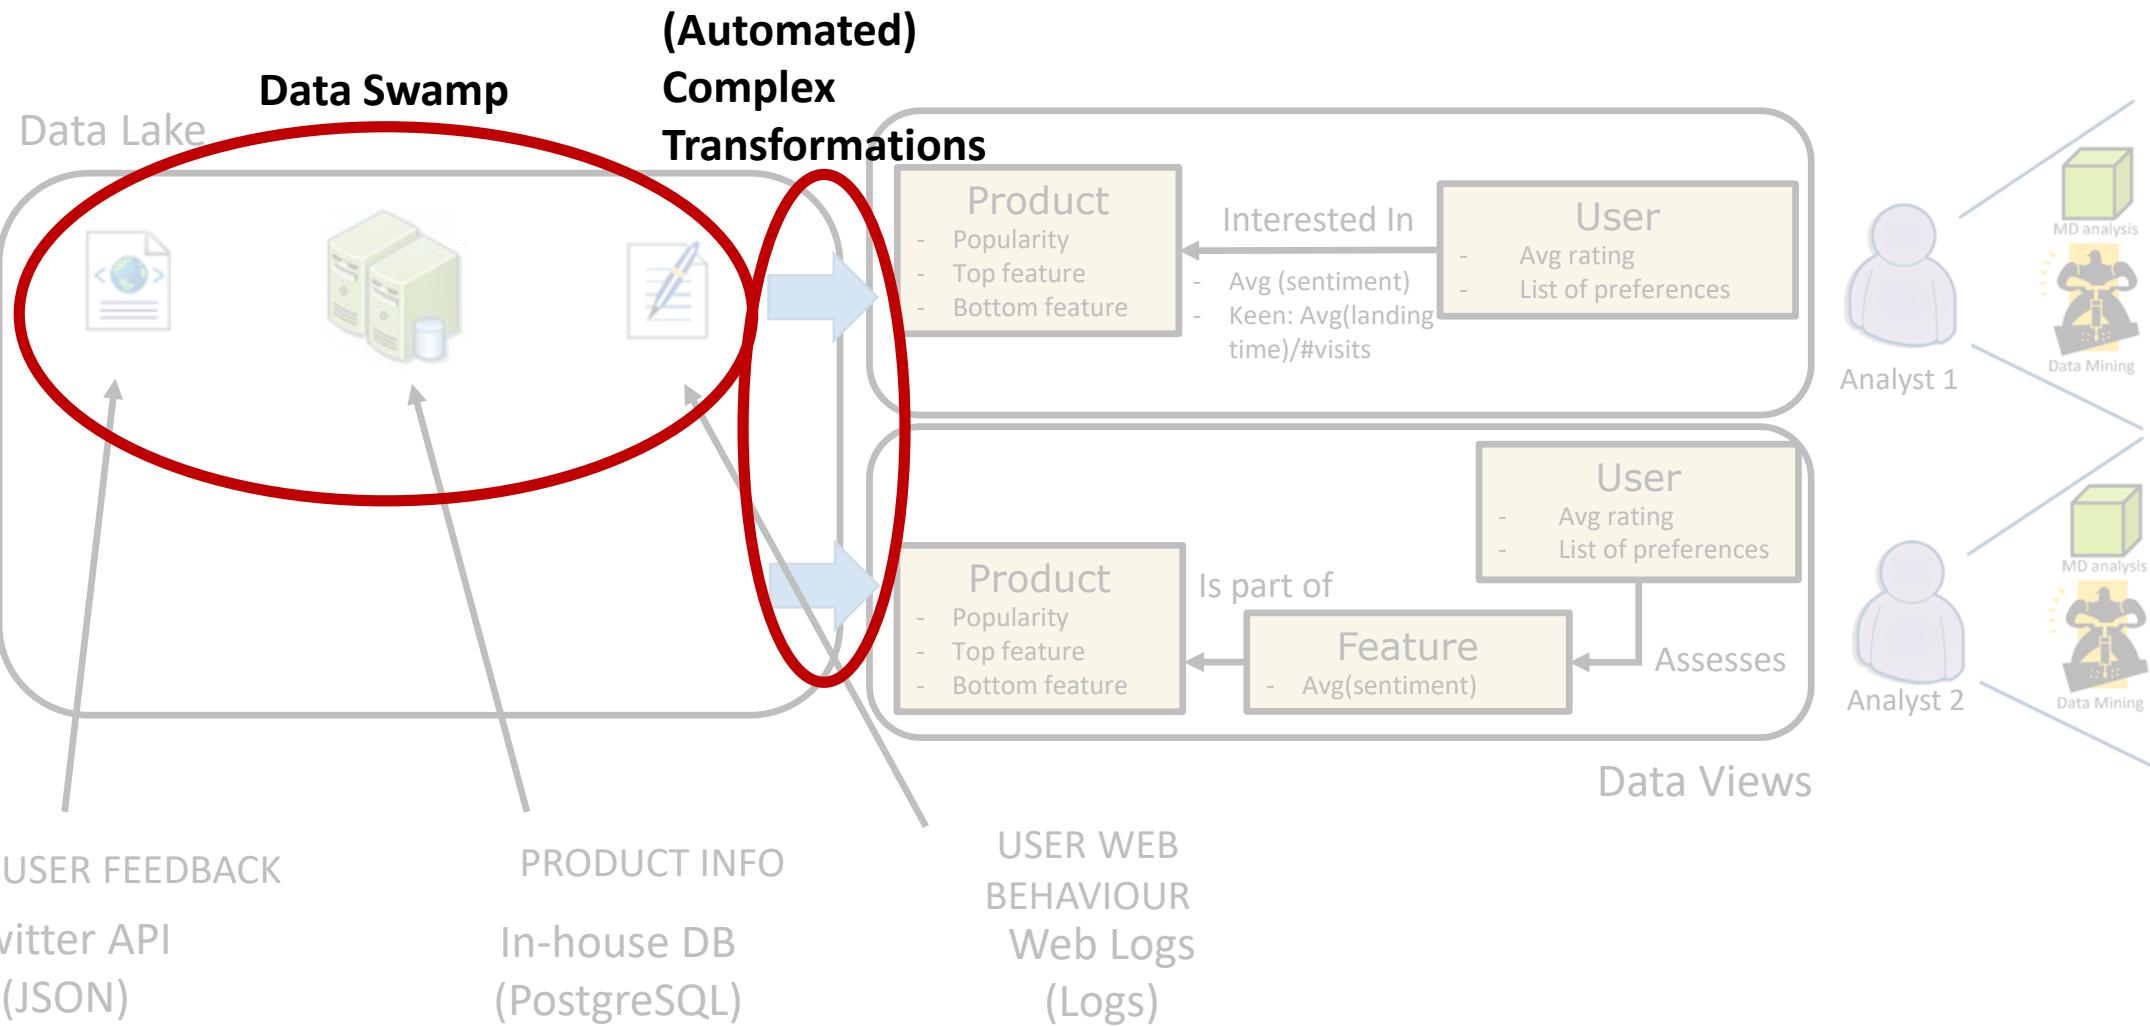
It is not a technology or architecture, it is a philosophy



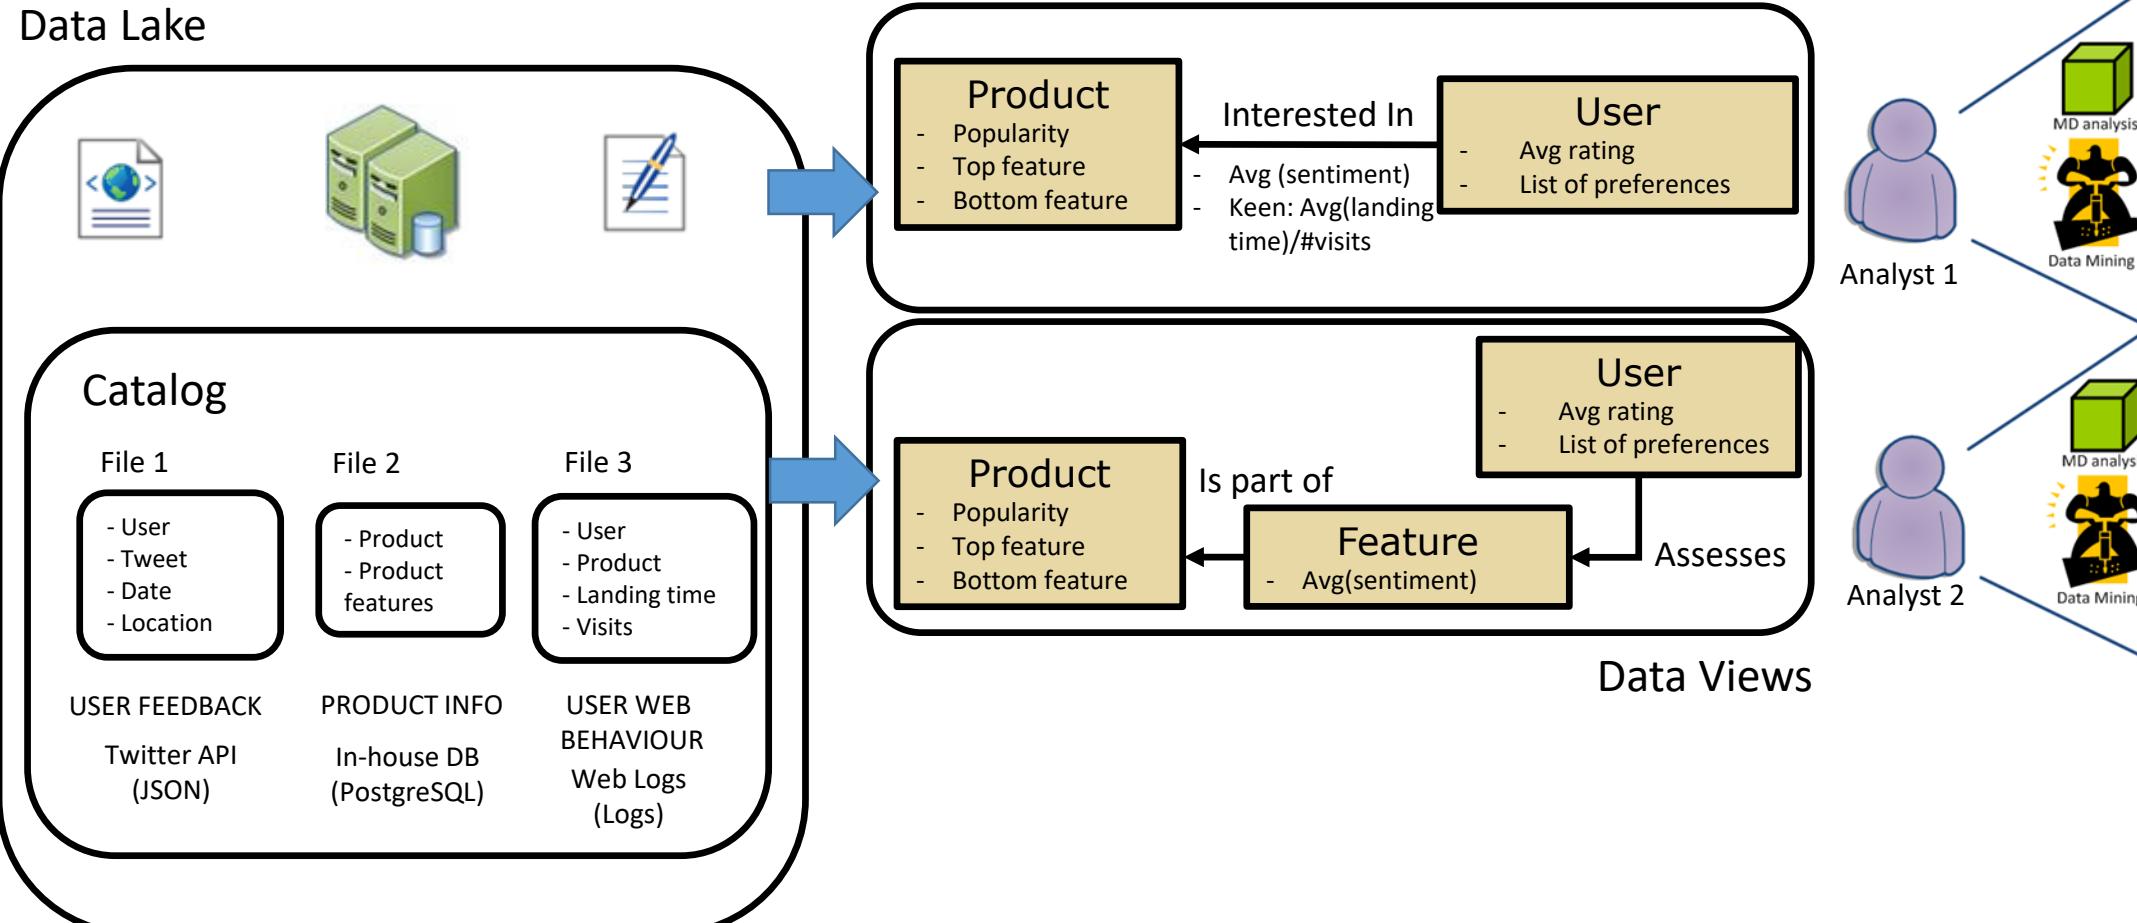
# Load-First Model-Later



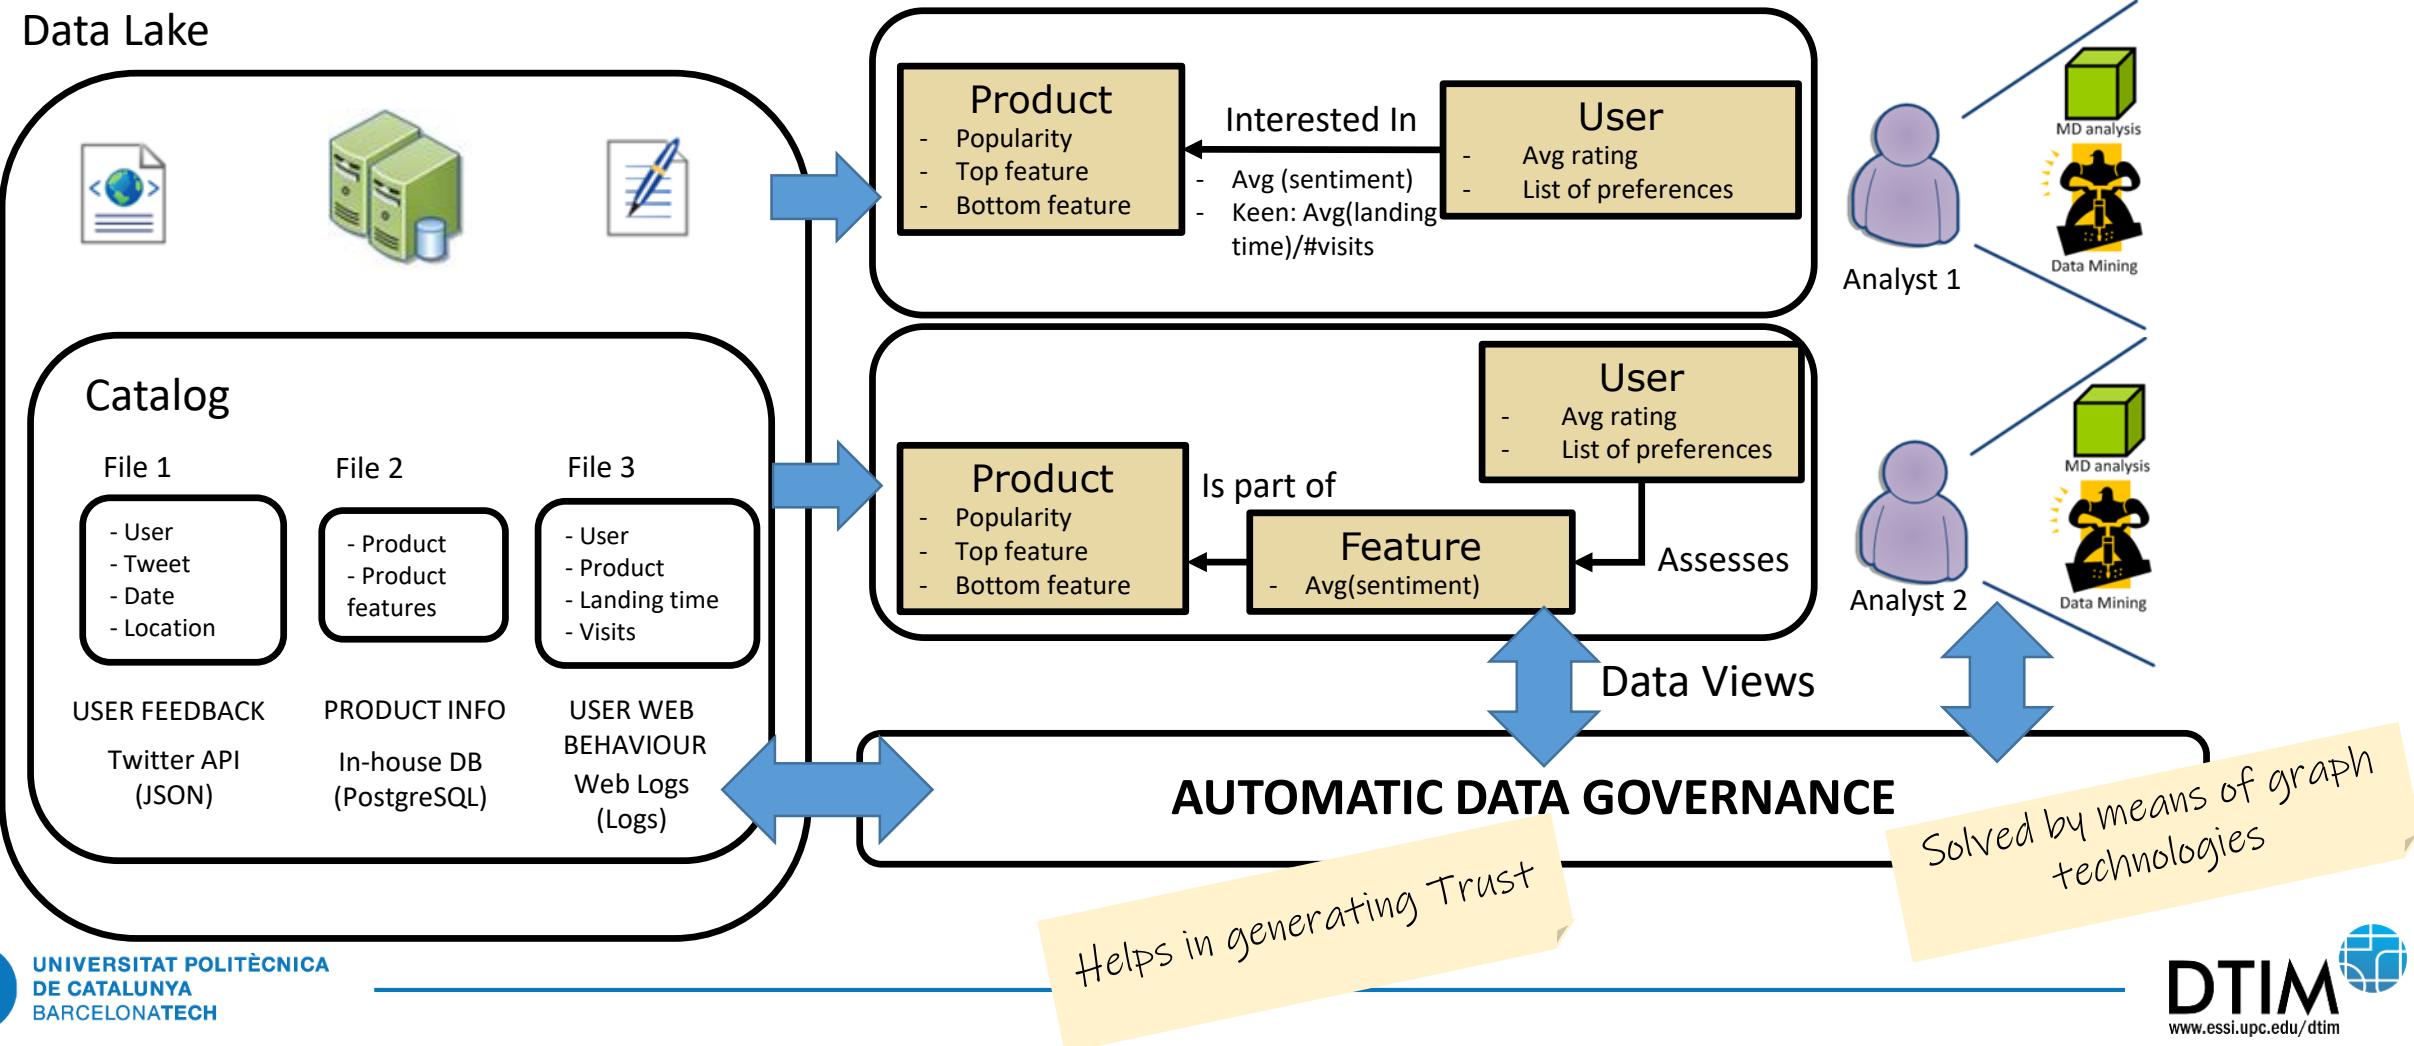
# Drawbacks



# From Data Swamps to Semantic Data Lakes



# From IT-Centered to User-Centered

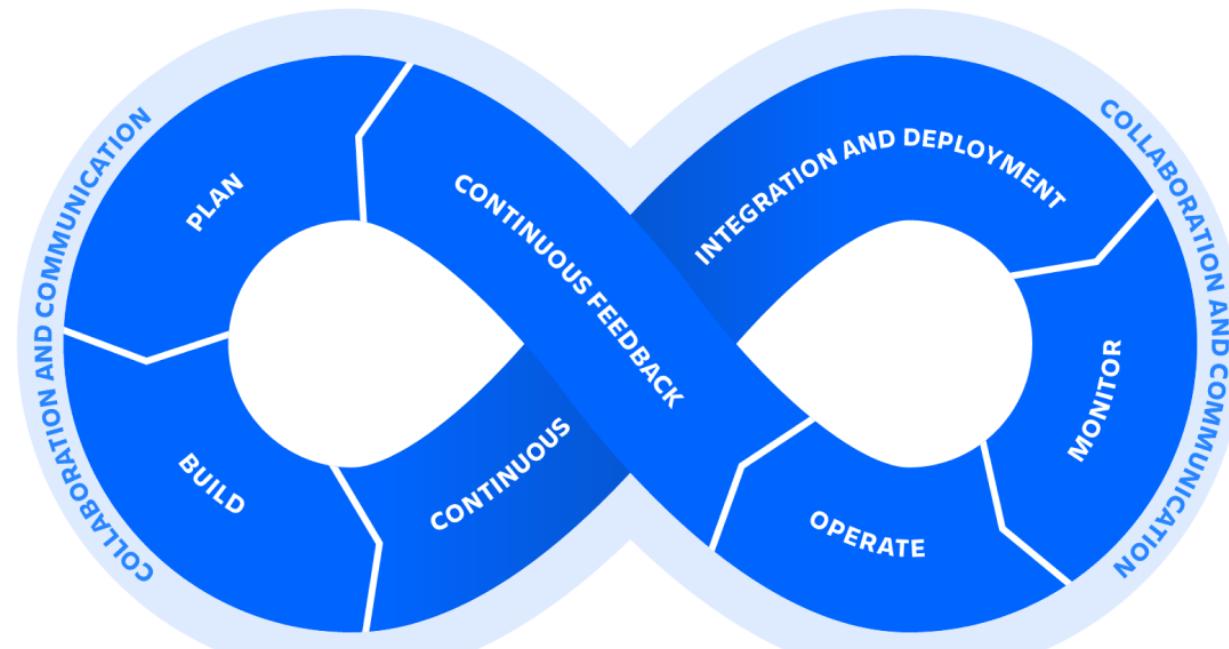


# DevOps

DevOps is about fast, flexible development and provisioning processes.

*It integrates the two worlds of development and operation, using automated development, deployment, and infrastructure monitoring. Its an organizational shift in which, instead of distributed isolated groups performing functions separately, cross-functional teams work on continuous operational feature deliveries.*

Ebert, C., Gallardo, G., Hernantes, J., Serrano, N. (2016). DevOps. *IEEE Software* 33(3): 94-100.



# From DevOps to DataOps / MLOps

DevOps ignores a key aspect in Data Science: data and its life cycle

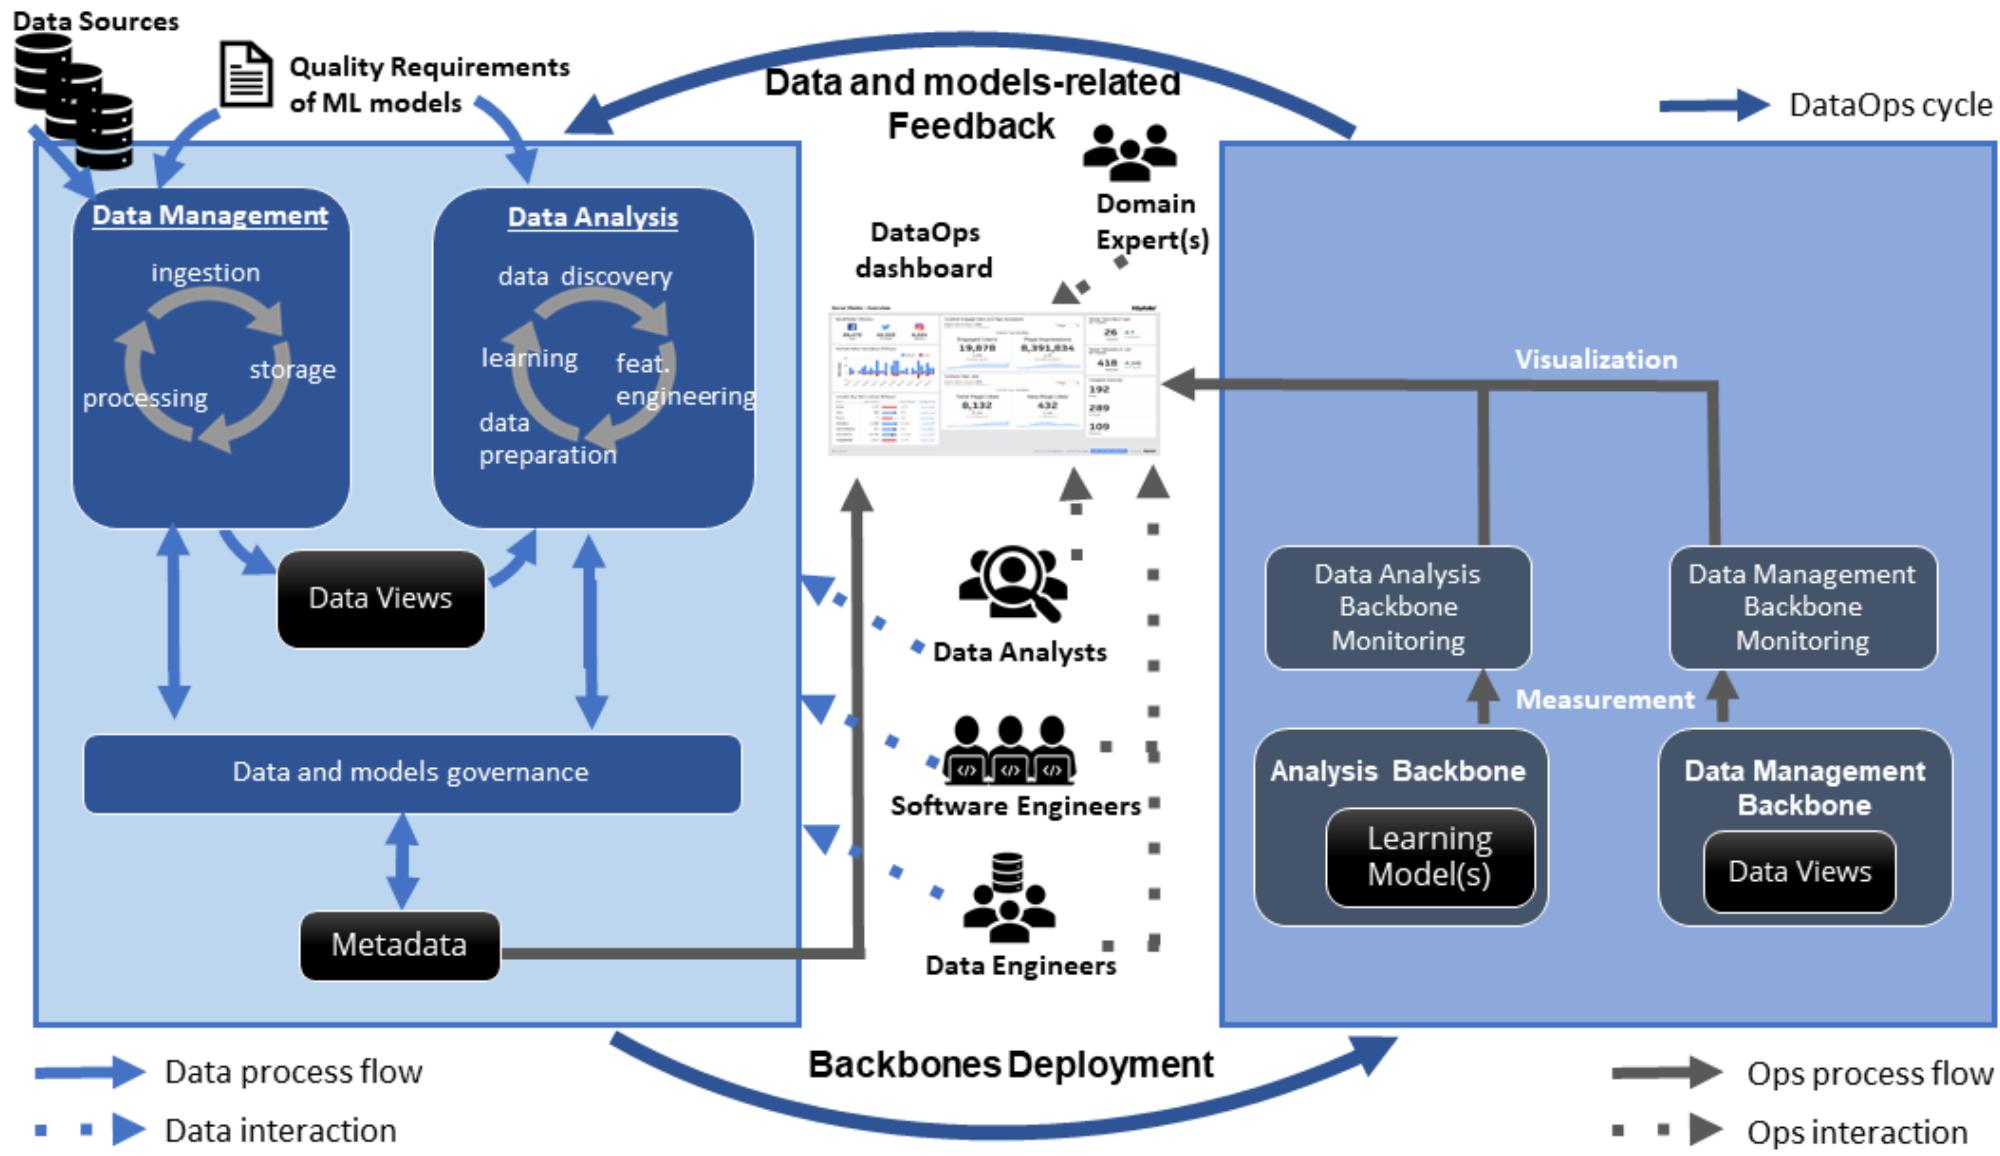
*DataOps* was introduced to cover this gap by *combining an integrated and process-oriented perspective on data with automation and methods from agile software engineering, like DevOps, to improve quality, speed, and collaboration and promote a culture of continuous improvement.*

Ereth, J. (2018). DataOps - Towards a Definition. LWDA 2018: 104-112.

*In short, the DataOps life cycle is needed to manage the complexity of the data life cycle in data science projects (*data engineering*)*

*Disclaimer:* you may read about MLOps too. However, in essence, DataOps / MLOps talk about the same problem. Simply, the latter focuses more on the ML part while the former covers the whole data lifecycle.

# DataOps in a Nutshell



# DataOps in a Nutshell

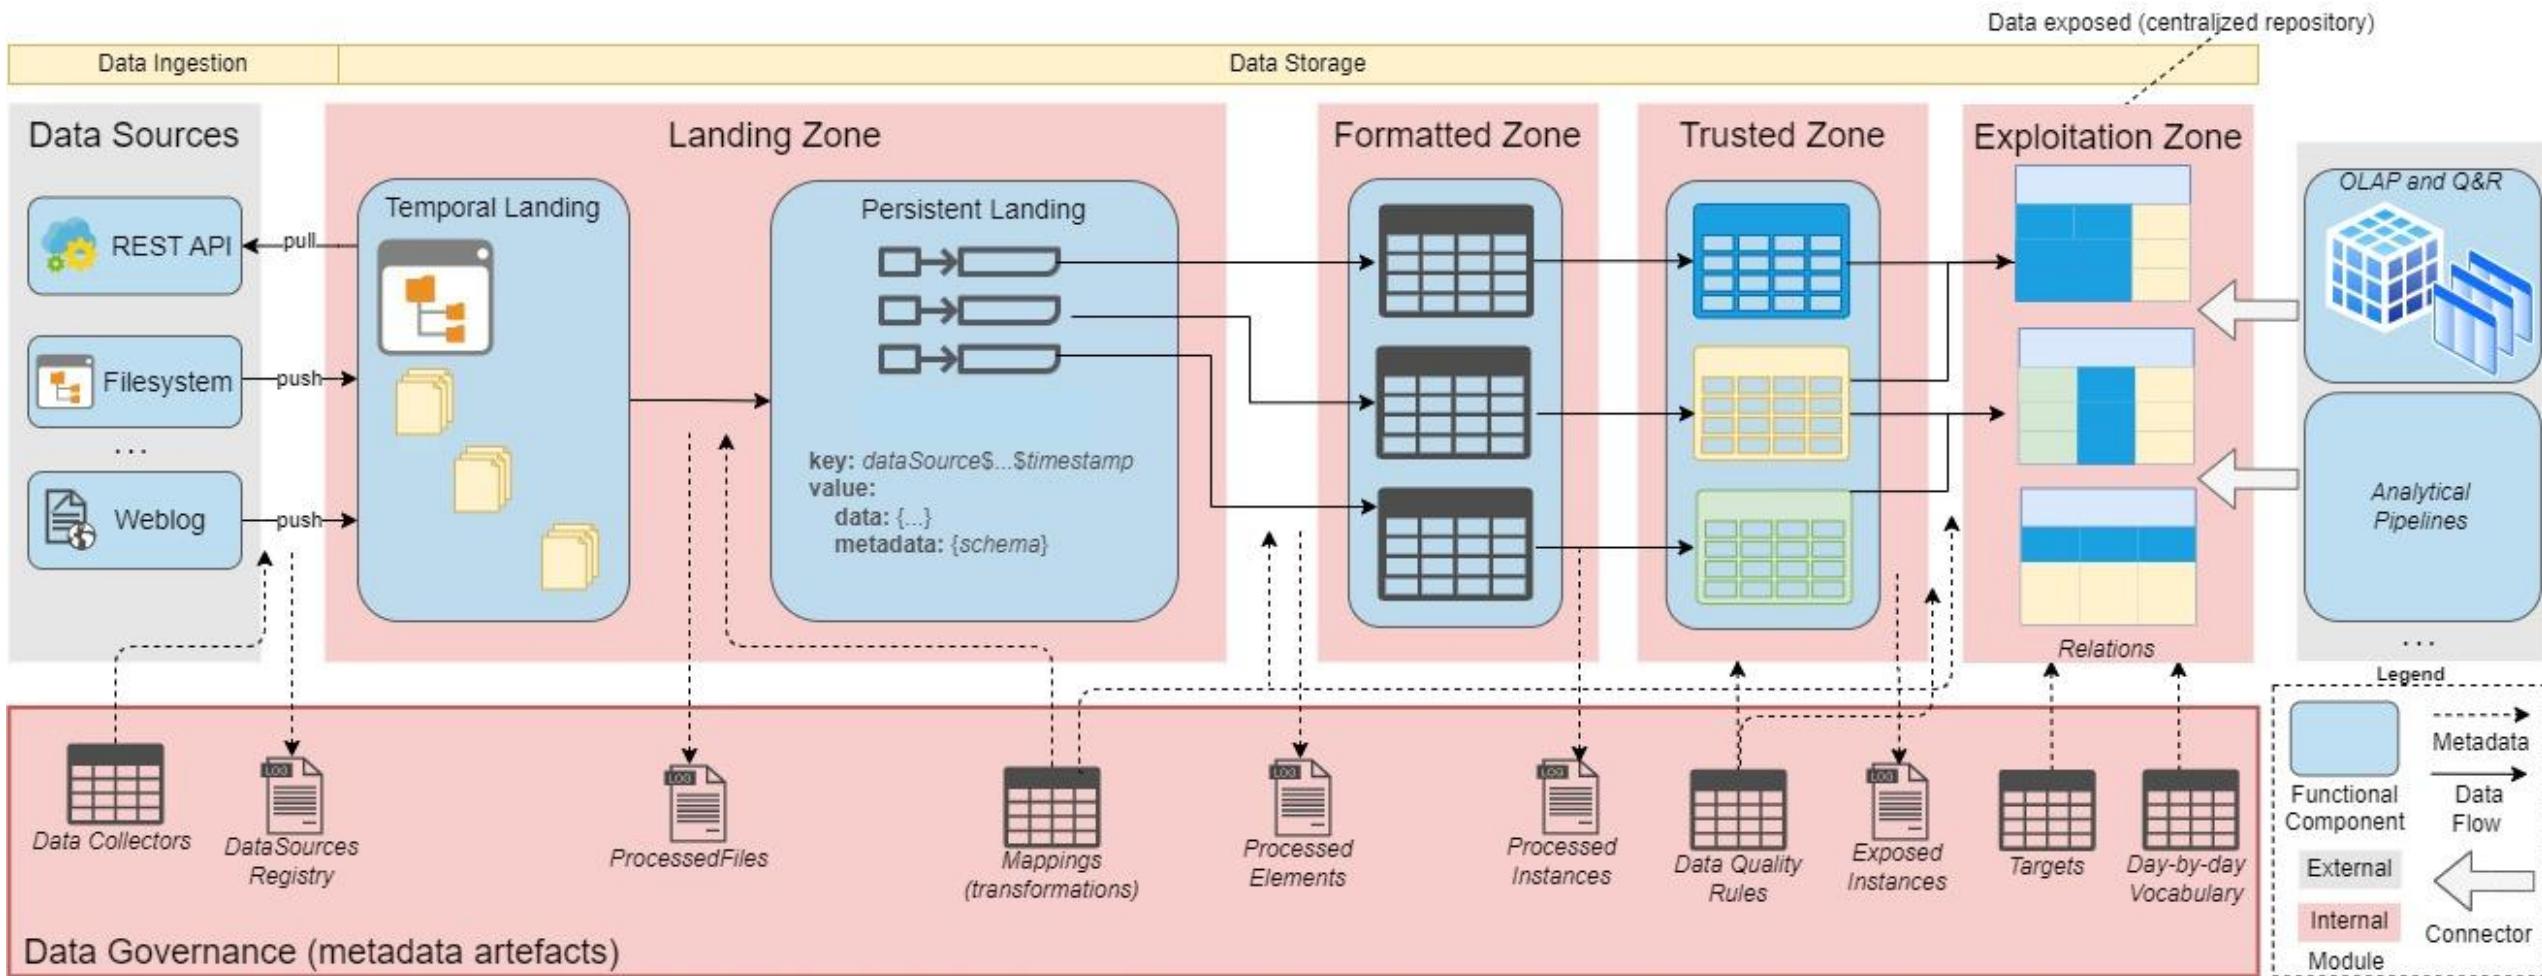
## The data management backbone (common for the whole organisation)

- Ingest and store external data into the system
  - Data Integration (standardization and data crossing)
    - Syntactic homogenization of data
    - Semantic homogenization of data
  - Clean data / eliminate duplicates / ...
- Expose a cleaned and centralized repository
  - Data profiling

## The data analysis backbone (repeats for every analytical pipeline)

- Extract a data view from the centralized repository
- Feature engineering
- Specific pre-processing for the given analytical task at hand
  - Labeling
  - Data preparation specific for the algorithm chosen
- Create test and validation datasets
- Learn models (either descriptive statistical analysis or advanced predictive models)
- Validate and interpret the model

# The Data Management Backbone



# The Data Management Backbone

- Data is ingested in the **landing zone** as it is produced (raw data)
  - The temporal landing stores temporary the files, until they are processed
  - The persistent landing tracks ingested files by data source and timestamp (i.e., versions)
- Data is then homogenized, according to a canonical data model in the **formatted zone** (syntactic homogenization)
- The **trusted zone** is where data cleaning happens.
  - Data quality processes: per dataset
- The **exploitation zone** exposes data ready to be consumed / analysed either by advanced analytical pipelines or external tools. Two main kind of tasks are conducted to generate this zone (semantic homogenization):
  - Data integration: new data views are generated by combining the instances from the trusted zone. A view may serve several data analysis.
  - Data quality processes: across datasets

# Governing the Data Management Backbone

The **data governance** layer must fulfill the following tasks:

- **Monitor** all the data flow executions between zones, as well as the required metadata to allow their incremental execution (logs)
- Metadata artifacts capturing the **schema** of the data views created in the exploitation zone, but also in any of the intermediate zones
- Metadata required to **automate processes**: templates to process different types of data, database connections, etc.
- **Mappings** between zones: transformations conducted between elements of two consecutive zones to enable data lineage / traceability
- **Profiling** information of every zone: for maintenance purposes for the landing, trusted and formatted zones and analytical purpose in the exploitation zone

# Tools for DataOps



# Tools for the Data Management Backbone

- **Ingestion**: it can be performed by means of off-the-shelf tools or ad-hoc scripting. Further, it can be pull or push
  - Scripting: Java, Python, bash scripts, etc.
  - Tools: Kafka, Spark, etc.
- **Storage**: when facing Big Data, distributed storage is a must. Anyway, specially in the exploitation zone, a vast range of tools might be used to store a given data view
  - Distributed storage: distributed file Systems (e.g., HDFS) and distributed databases (e.g., HBase, SimpleDB, Cassandra).
  - Centralized storage: relational databases (e.g., PostgreSQL, DuckDB), etc.
- **Data quality and profiling**: in the past, they were tightly related to the analytical tools (e.g., SAP, IBM, etc.). But more and more there is an extensive offer of stand-alone tools that fit better the required *openness* view of data science: Trifecta, HoloClean, Attacama, Informatica, Dataprep.ai, etc.
- **Transformations between zones**: Spark and Flink the most used ones. There is the option of going to a lower level of abstraction and use frameworks that allow to tailor the execution to your needs and gain efficiency: e.g., Akka.

# Tools for Governance

Knowledge graphs are nowadays the standard the facto to govern Data Science projects. We differentiate repositories from generic frameworks supporting all data governance steps

- **Repositories**: GraphDB, Stardog, Amazon Neptune, etc.
- **Governance frameworks**: typically embedded as part of the whole data management backbone (specially in Cloud services). However, there are quite a few stand-alone solutions:
  - Cloud services: Delta Lake (Databricks), Azure Purview (Microsoft), IBM Data Governance, Data Catalog (Google Cloud), etc.
  - Standalone: Apache Atlas, Collibra, Semantic Web Company, Salesforce Einstein, Alteryx, DataRobot, H2O Driverless AI, etc.

# Main Programming Languages

- Development environment

- An exploratory phase where data scientists aim to flexibly explore the data sources at hand, select the most relevant variables or evaluate different models.
- Desired features
  - Dynamically typed languages
  - Interpreted code (no compilation required)
  - Unnecessary verbose code
- Exemplary languages: Python, R, Perl

- Operations environment

- A consolidation phase where the goal is to deploy a data infrastructure with performance guarantees and robust to external changes.
- Desired features
  - Statically typed languages
  - Compiled code (low-level optimization and memory management)
  - Built-in parallelism for concurrent applications
  - Modularity
  - Advanced error and exception management
- Exemplary languages: Java, Scala, C++

# Summary

- Variety in Big Data requires Data Integration to obtain valuable insights
- Data Quality must be guaranteed
  - Accuracy
  - Completeness
  - Consistency
  - Timeliness
- DataOps to operationalize data management and analysis

# Bibliography

- M. Golfarelli and S. Rizzi. *Data Warehouse Design*. McGrau-Hill, 2009
- R. Kimball, J. Caserta. *The Data Warehouse ETL Toolkit*. Wiley Publishing, 2004
- J. Han and M. Kamber. "*Data Mining: Concepts and Techniques*". Morgan Kauffman Publishers, 2000
- Carlo Batini, et al. *Methodologies for data quality assessment and improvement*. ACM Computing Surveys, 41(3), 2009
- Arkady Maydanchik. *Data Quality Assessment*. Technics Publications, 2007
- H. Garcia-Molina et al. *Database Systems*. Prentice Hall, 2009
- J. Bleiholder and F. Naumann. *Data Fusion*. ACM Computing Surveys, 41(1), 2008